

Bills That Never Reached the Floor: Recovering Negative Agenda Control from Text

Kisoo Kim*

July 2, 2026

Abstract

The majority party rarely loses on the House floor, and the standard explanation is that it blocks the bills a cross-party majority would pass over its objection. How often that would have happened has never been measured, because those bills never came to a vote. A TF-IDF ridge model trained on three decades of House roll calls reads each never-voted bill's would-be floor position and computes the probability it would have rolled the majority party. The bills the party never let reach a vote carried several times the roll risk of the bills it did, and that risk falls gate by gate from committee to the floor: the screening the cartel account predicts and majoritarian accounts do not. Blocking concentrates on broadly cosponsored coalition threats and on the social, labor, and distributive domains that split the party: a policy map of party power that no floor-based measure can draw.

Keywords: negative agenda control; agenda setting; party cartel; roll rates; text as data; U.S. House of Representatives

Word count: 10,257 (including tables and notes; excluding references and the Online Appendix)

*Email: kisoo.kim@vanderbilt.edu

1. Introduction

In October 2013, House Democrats introduced H.R. 15, a comprehensive immigration-reform bill built on legislation the Senate had passed months earlier with sixty-eight votes. It drew two hundred cosponsors, Republicans among them, and contemporaneous vote counts suggested a floor majority stood ready to pass it. It never received a vote. Speaker John Boehner declined to schedule it, holding to the practice of refusing floor time to any bill opposed by most of the majority party, and a discharge petition to force the question stalled well short of the signatures it needed. The death of the bill itself was ordinary. Most House bills die without ever receiving a vote. What makes H.R. 15 unusual is that its death left a record: the contemporaneous vote counts preserve what the House stood ready to do with a bill it never voted on. For the thousands of other bills that die the same way in every Congress, no one is counting, and what the floor would have done with them goes unrecorded. The measure developed in this paper recovers that answer. Reading nothing but the bill's text against the 113th House's revealed voting geometry, it puts the probability that H.R. 15 would have rolled the majority party (passed over the opposition of most of its members) at 87 percent, higher than nine in ten of the bills that shared its fate.

The majority party in the U.S. House of Representatives almost never loses on the floor. Across recent decades, the majority party has been on the losing side of a final-passage vote, or "rolled," on only about one percent of the bills that reached a vote. Cox and McCubbins (2005) read this near-absence of defeats as the signature of negative agenda control: the party uses its command of the floor schedule to keep bills that would divide or defeat it from ever coming up. The low roll rate is the evidence. But it is a peculiar kind of evidence, because it is measured only on the bills the majority party chose to let through. The bills it actually blocked, the ones on which it would have been rolled, never reached a vote, and so never enter the measure. The power is inferred from the contests the majority party wins, while its exercise lies in the contests it never allowed

to happen.

This is a selection problem at the center of how political scientists measure party power in legislatures. The roll rate, and the larger family of measures built on recorded votes, from spatial scalings of legislator ideal points to pivot models of gridlock, can only see bills that reached the floor (Poole and Rosenthal, 1997; Krehbiel, 1998). What a blocked bill would have done had it come up for a vote is, by construction, unobserved. The scale of the omission is not small: for every House bill that received a passage vote over the past three decades, roughly fifteen died without one. So the dominant measurement tradition is silent on two questions that its own theory makes central. How much agenda control is hidden by this selection: were the blocked bills meaningfully more threatening to the majority party than the ones it let through, or only marginally so? And where does that control concentrate: on which bills, in which areas of policy, does the party's blocking power actually operate? Existing measures cannot answer either, because they never observe the object of interest.

I ask how roll-prone the bills the majority party kept off the floor would have been, relative to the bills it let through, and on what kinds of bills the party's blocking power concentrates.

My answer begins from a simple observation: a bill's text carries its ideological content, and that content is what determines whether it would split the majority party on the floor. Within a single Congress, the words of a bill predict the ideological cutpoint it would occupy in the chamber's revealed voting space well enough to place it on the floor's left-right map. I build a recovery model that learns this text-to-position mapping on the bills that did reach a vote, where the true floor position is observed, and then deploys it on the universe of bills that never did. For each blocked bill I recover the position it would have taken and, given the chamber's estimated ideal points, the probability it would have rolled the majority party. The gap between the roll-proneness of the blocked bills and the same model's roll-proneness on the bills that reached the floor — what I call

the agenda-control wedge — is the magnitude of negative agenda control, measured on the bills the theory is actually about.

I answer these questions using the full text of House legislation. Bill text from the BICAM corpus (Delano et al., 2025), linked to Voteview roll-call records, covers the House across Congresses 103 through 118, with more than eight thousand eligible final-passage votes and a never-voted universe of roughly 122,000 bills. A per-Congress item-response model of the recorded votes supplies the floor’s ideal points and each passed bill’s true position; these serve both as the target the recovery model learns and as the ground on which I validate it. The recovery is checkable precisely where it matters: on the floor bills, I can compare the text-recovered position against the observed one, and I can confirm that feeding the model the true parameters reproduces the observed roll rate almost exactly. Where no vote exists to check against, the deployment is validated against an independent signal the recovery model never sees: the revealed ideology of each bill’s cosponsoring coalition.

I find that the bills the majority party never let reach a vote carried several times the majority-roll risk of the bills it did: the hidden exercise of agenda control is large, not marginal. The screening is visible at every stage of the legislative process — roll risk falls step by step as bills clear committee and approach the floor, the profile of a filter working gate by gate on precisely the dimension the cartel account says the majority party guards. And the control is organized by what bills are about. Blocking power concentrates on social, labor, and constituency-facing distributive legislation and is weakest on fiscal, procedural, and trade measures, a policy-domain map of agenda control that no floor-only measure can draw; the most roll-prone blocked bills are broadly cosponsored coalition threats, not obscure lone-sponsor proposals. At the same time, who sponsors a bill remains the larger single driver of whether it would roll the majority party. Text earns its place not by beating party but by reading a bill’s content beyond its sponsor’s identity, recalibrating the positions sponsor identity alone would overstate, and mapping agenda

control onto the substance of policy.

My argument speaks to a literature that has measured party agenda power almost entirely from what happens on the floor. The roll-rate tradition established that the majority party is rarely defeated and interpreted that fact as negative agenda control (Cox and McCubbins, 2005; Den Hartog and Monroe, 2011); the pivotal-politics account offered a majoritarian alternative in which gridlock, not party scheduling, produces the same floor patterns (Krehbiel, 1998). Both sides of that debate reason from the bills that reached a vote. I move the analysis to the bills that did not — and the move has adjudicative bite, because a majoritarian filter should be indifferent to which party a passable bill endangers, so a blocked pool sorted by majority-party roll risk, gate by gate and within sponsor party, is evidence on the central question that floor-based tests could not supply in principle.

The recovery model belongs to a text-as-data lineage, and the difference from that lineage is where the method’s contribution lies. Models that read congressional text to predict voting have been trained and evaluated on bills that reached the floor (Gerrish and Blei, 2011), and models of bill progression have used text to forecast which bills survive committee (Yano et al., 2012). Both stop at the boundary this study crosses: the first never leaves the voted sample, and the second predicts whether a bill advances rather than what its passage would have revealed about the majority party. Likewise, where roll-rate measures know a bill’s sponsor and party but not its would-be floor position, and so can characterize agenda control only in terms of who is blocked, a text-based recovery knows what each bill is about and can locate agenda control in the substance of policy — a distinction that turns out to be the empirical heart of the paper.

The contribution is thus twofold. Methodologically, I introduce a recovery model that reads a blocked bill’s would-be floor position from its text, deployed on the never-voted universe, validated where the counterfactual is checkable, and hardened against the concern that it merely relabels sponsor identity. Substantively, I quantify how much agenda

control the floor-only view hides and, more consequentially, show that it is organized around the policy domains that split the majority party — a substantive map of party power that the field's standard measures are structurally unable to produce.

The model cannot say who killed H.R. 15. It can say, for that bill and for each of the 122,000 like it, how much danger it posed to the majority party, and that is the quantity the study of party power has been missing.

2. Theory: negative agenda control and the counterfactual roll

Negative agenda control is the power to keep bills off the floor. In the cartel account of legislative organization, the majority party functions as a procedural coalition that monopolizes the scheduling prerogatives of the chamber, from control of the Rules Committee to the calendar to what is recognized for a vote, and uses them to protect itself from losing (Cox and McCubbins, 1993, 2005). This account belongs to a longer tradition treating the party as an organized team whose cohesion rises with the homogeneity of its members' preferences (Rohde, 1991; Aldrich, 1995). The protection has a specific target. The majority party is endangered not by bills the minority party favors and it opposes, which it can defeat on the floor by simple numbers, but by bills that split it: measures a floor majority would pass but that a majority of the majority party opposes. Letting such a bill reach a vote means either losing on the floor or forcing the party's own members into a damaging choice. Keeping it off the floor avoids both. The observable trace of this power is the roll, an instance in which a bill passes over the opposition of most of the majority party. If negative agenda control operates, rolls should be rare, and empirically they are: the majority party is rolled on only a small fraction of floor votes, the fact from which the theory draws its principal support.

The difficulty is that the roll rate is computed only over bills that reached the floor,

which are precisely the bills the majority party's scheduling power selected in. A low floor roll rate is consistent with strong agenda control, because the dangerous bills were screened out, but the screening itself is never measured. The bills that were kept off the floor are the ones on which the majority party anticipated a roll, and those bills, not the ones that passed, are where the theory locates the exercise of party power. To measure agenda control on its own terms, I need the quantity the floor never reveals: for a bill that did not reach a vote, the probability that it would have rolled the majority party had it come up. I call this the bill's counterfactual roll. It is defined by the same logic as an observed roll (a majority of the chamber in favor, a majority of the majority party opposed) but evaluated against the chamber's revealed ideal points rather than an actual tally, because no tally exists. The counterfactual is deliberately spare: the bill as written, put to an up-or-down vote, with members voting their revealed positions — no amendments, no floor discipline. That spareness is the point. An as-written, whipping-free floor test is precisely the event negative agenda control exists to forestall, and it is the event the chamber's estimated ideal points make calculable.

The counterfactual roll makes agenda control measurable as a single quantity. Let the roll-proneness of a set of bills be the average counterfactual roll probability over that set. The agenda-control wedge is the difference between the roll-proneness of the bills that never reached a passage vote and the roll-proneness of the bills that did, with both sides evaluated by the same recovery model, so that systematic reading error common to the two pools appears on both sides of the comparison and cancels rather than masquerading as a wedge. Error that differed between the pools would not cancel; Section 7 tests for exactly that. A wedge of zero would mean the chamber's filter passes bills to the floor without regard to the majority party's exposure: scheduling unrelated to the threat of a roll, and no negative agenda control worth the name. A positive wedge means the bills held back were systematically more roll-prone than the bills let through (the filter screens on exactly the dimension the cartel account says the majority party guards), and the size

of the wedge is the amount of party power the floor-only roll rate leaves out.

One definitional discipline matters here. A bill that never reached a passage vote is not, bill by bill, proof of a majority-party veto: bills die in committee for many reasons, expire on crowded calendars, or are introduced with no expectation of movement. What never-voted status identifies is the output of the chamber's filter — the sequence of majority-controlled gates, from committee through the Rules Committee and the scheduling power, that every bill must clear. The wedge is accordingly the footprint of that filter: a measure of what the process that selects bills onto the floor screens on, taken over the universe of bills introduced, not a causal effect of a designated act of blocking. The footprint reading also makes the quantity testable at finer grain. A filter with several gates should leave its mark at each of them: if the screen operates on roll risk, the bills that advance at every stage should carry less of it than the bills that stall there.

Framed this way, the wedge separates the two leading accounts of legislative organization on ground neither has been able to reach. In the majoritarian account, floor access is governed by chamber majorities rather than party gatekeepers (Krehbiel, 1998): conditional on a bill's capacity to command a floor majority, the danger it poses to the majority party should not predict whether it reaches a vote. In the cartel account, that danger is precisely what the gates are for. A positive wedge — a never-voted pool weighted toward bills that would pass while splitting the majority party — is what the cartel predicts and a majoritarian filter should not generate; a wedge near zero would say the floor's low roll rate reflects gridlock and consensus rather than partisan screening. The counterfactual roll thus turns the central disagreement about party power into an estimable comparison.

My theoretical expectations concern not only whether the wedge is positive but where it concentrates. Two predictions follow from the logic of the cartel. First, agenda control should concentrate on organized threats. A bill with broad cosponsorship is a credible candidate to reach the floor, since it has a coalition behind it, and blocking it is therefore a meaningful act of control, whereas a lone-sponsor bill that would die on its own requires

no protection. Roll-proneness among blocked bills should thus rise with the breadth of the coalition supporting them, not fall. Second, and more distinctively, agenda control should concentrate in the policy domains that split the majority party. The bills that cross-pressure a majority party are those touching the social, labor, and redistributive questions on which its members' districts diverge; fiscal housekeeping and procedural measures rarely divide it. If negative agenda control is organized by the substance of policy, its footprint should be visibly heavier in the former domains than the latter.

This second prediction is what separates my account from a simpler alternative. One might explain a positive wedge without any appeal to policy content: the minority party sponsors bills far from the majority party's position, everyone can see they would fail or embarrass it, and so they are kept off the floor. On that story the wedge is entirely a matter of who sponsors a bill, a party-composition effect that would leave no structure once sponsor identity is held fixed. The content account makes a different prediction. If agenda control is organized by what bills are about, then the concentration of blocking power across policy domains should survive holding sponsor party constant: among bills sponsored by the same party, the socially and redistributively charged domains should still be the most roll-prone. The two accounts agree that blocked bills lean toward the minority; they disagree about whether, beyond that, agenda control has a policy structure. That disagreement yields a clean test, and settling it requires reading a bill's content, not just its sponsor, which is exactly what recovering position from text allows and what a measure built on sponsor identity cannot do.

Recovering the counterfactual roll therefore rests on one measurement premise: that a bill's text reveals the floor position it would have occupied. The premise is not obvious, and much of what follows is devoted to establishing and stress-testing it: showing that within a Congress text recovers a bill's revealed cutpoint, that it does so beyond what the bill's sponsor already reveals, and that the recovery survives deployment on bills that never reached a vote. Establishing that premise, and carrying it from the voted bills onto

the never-voted ones, is the work of the recovery model I develop next.

3. Data and the recovery model

The empirical strategy has two moving parts: a floor on which the counterfactual can be validated, and a model that reads a bill’s position from its words. Both are built on the House of Representatives from the 103rd through the 118th Congress, a span of thirty-two years that brackets the rise of partisan polarization and places majority control in both parties’ hands at different times.

The recorded votes come from Voteview; the bill texts come from BICAM (Bulk Ingestion of Congressional Actions and Materials), a corpus that pairs the full legislative text of House and Senate bills, version by version, with a crosswalk to the Voteview roll-call record (Delano et al., 2025). Linking the two is where a subtle but consequential measurement decision enters. A bill exists in several textual versions as it moves through the chamber, and the version that carries information about the vote (the “Engrossed” text produced *after* passage) leaks the outcome the model is supposed to predict. I therefore use only pre-decision versions: the latest text available before the vote, drawn from the introduced, reported, referred, or calendared stages. A fifty-vote manual audit confirms that this rule selects a genuinely pre-decision version in forty-six of fifty cases, against a naive “latest available” rule that pulls post-passage text in most rows. This is the leakage linchpin of the design, and everything downstream depends on it.

The linked corpus is large. Of the House roll calls in the period, 8,324 are eligible final-passage votes (votes on the text of a bill or resolution itself, not on amendments or procedure) on 8,021 distinct bills, 93 percent of them House-originated, and text is recovered for 99 percent. Beyond the voted bills lies the object of the study: roughly 122,000 House bills that were introduced with text but never reached a final-passage vote, drawn from the same period under the same pre-decision text rule as the voted sample

(Appendix A).

The floor’s geometry, where each legislator sits and where each bill’s cutpoint falls, is estimated separately within each Congress by an expectation-maximization item-response model of the recorded votes (Clinton et al., 2004; Imai et al., 2016). This yields, for every passed bill, a discrimination parameter and a cutpoint: the point on the left-right scale at which a legislator is indifferent between yea and nay. These estimated positions are the target the recovery model learns and the ground truth against which it is checked, so their credibility matters. Two validations support them. The estimated ideal points correlate with first-dimension DW-NOMINATE scores at 0.978 across the sixteen Congresses, essentially the same map by a different method; the estimated cutpoints correlate with the NOMINATE cutting-line midpoints at 0.75, a strong agreement between two quantities constructed by different methods.

The recovery model itself is deliberately simple: a bag-of-words representation of each bill’s text, weighted by term frequency and inverse document frequency, mapped to the bill’s cutpoint by cross-validated ridge regression (Grimmer and Stewart, 2013). Simplicity is a choice with a rationale, established below and in the robustness section: a far heavier full-text neural model buys almost nothing here, so the interpretable linear model is the workhorse. I evaluate it on two held-out splits. A random split holds out a fifth of bills drawn from across all Congresses; a temporal split trains on the earlier Congresses and predicts the latest three (116–118). The random split leaves bills from each test bill’s own Congress in training, so Congress-specific vocabulary and that Congress’s cutpoint scale reach the test set; holding out whole Congresses removes that leak, which makes the temporal split the honest test of whether the mapping travels across time.

Text recovers a bill’s cutpoint well on held-out bills. On the random split the text-predicted cutpoint correlates with the true cutpoint at 0.69 (Table 1); the permutation null in the Online Appendix puts the 97.5th percentile of chance recovery at 0.07. The recovery is not an artifact of the winsorization in the target’s construction. Cutpoints are

winsorized to the interval $[-2, 2]$ before evaluation (a quarter of well-identified bills have raw cutpoints beyond that range and are pinned at the bound), yet the rank correlation between predicted and true position, which is insensitive to that capping, is 0.70, indistinguishable from the Pearson figure. The bill’s margin of passage, its yea-share, is recovered even better and by text alone, at 0.78.

Table 1. Recovery of a bill’s floor position from its text (bag-of-words TF-IDF, cross-validated ridge).

quantity	across periods (temporal	
	within period (random split)	split)
cutpoint — Pearson r	0.69	0.36
cutpoint — Spearman ρ	0.70	0.37
margin of passage (yea-share) — r	0.78	0.67
discrimination (d_1) — r	0.89	0.23

Note: Well-identified test bills ($|d_1| > 0.3$, non-unanimous): $n = 803$ (random), 1,004 (temporal). Cutpoints winsorized to $[-2, 2]$; the Spearman ρ is cap-insensitive. Within-era recovery restricting train and test to Congresses 116–118 is 0.60 ($n = 209$). Temporal split trains on Congresses ≤ 115 and predicts 116–118; its Pearson r uses a train-only vocabulary fit (0.3554), while the Spearman ρ comes from the full-vocabulary fit; the two fits differ by less than 0.01.

The pooled 0.69 should not be read as the recoverability of any single Congress, however. The pooled random split lets the model exploit the spread of cutpoints *between* Congresses, which inflates a within-Congress reading. Estimated one Congress at a time, recovery ranges from 0.21 to 0.76, and the spread is wider than sampling error alone can produce: with 23 to 103 test bills per Congress the individual estimates are imprecise, yet

a homogeneity test across the sixteen Congresses still rejects a common value (Appendix Table A2 reports the intervals and the test). A genuine leave-one-Congress-out transfer, training on all others and predicting the held-out Congress, averages 0.45; that is a cross-period transfer quantity, not the within-period quantity the design needs. The cleanest within-period benchmark restricts both training and test to the same recent era (116–118) and recovers at 0.60 on those bills. The claim is therefore that within a period, a single recent era of adjacent Congresses, text recovers position at around 0.60, not 0.69, and that per-Congress recoverability is heterogeneous.

Across periods the recovery decays sharply. On the temporal split, text-to-cutpoint recovery falls to 0.36, while the margin of passage remains recoverable at 0.67. Two facts sit underneath this contrast. How lopsided a bill will be is a stable property of its text; where its ideological cutpoint sits is tied to each Congress’s particular geometry, which does not extrapolate. The paper’s estimand is consequently a within-period quantity, a point I return to as a scope condition rather than a defect.

3.1 Does text read content, or only the sponsor?

The natural objection to a text measure of a bill’s position is that it may be reading nothing more than who wrote the bill. A member’s bills carry that member’s ideology; a classifier could recover a bill’s cutpoint by learning its sponsor’s identity and never engaging its content. If so, the text measure would add nothing to what a roster of sponsors already provides, and the eventual claim to have mapped agenda control by *content* would collapse.

The objection is testable, and the test is central to the paper. I build a sponsor-identity baseline from the primary sponsor’s and cosponsors’ own estimated ideal points and party, and ask how much the text adds beyond it. Sponsor identity is genuinely informative: within a Congress it recovers a bill’s cutpoint on its own at 0.59. But text carries substantial position information beyond identity. Regressing the text prediction on the

sponsor baseline and measuring the residual recovery, the increment from text is 0.45 out of sample, and the two together (0.72) exceed either alone — text is reading content, not relabeling the sponsor (Table 2). The increment is conservative by construction: wherever the sponsorship record supplies cosponsors, roughly half of the well-identified bills, they join the baseline, and a stronger identity baseline can only shrink the increment that text earns beyond it.

Table 2. Text’s contribution beyond sponsor identity (cutpoint recovery, r).

	sponsor identity	text	text + sponsor	text increment
within period (random split)	0.59	0.69	0.72	0.45
across periods (temporal split)	0.37	0.35	0.40	0.16

Note: Sponsor-identity baseline: primary sponsor and cosponsor estimated ideal points and party (cosponsor-augmented; cosponsors present for roughly 48% of well-identified bills, primary sponsor otherwise). Text increment: out-of-sample residual recovery of the text prediction after projecting out the sponsor baseline. A primary-sponsor-only baseline yields a larger within-period increment (0.55).

The same comparison across periods locates exactly where the measure decays. On the temporal split the sponsor baseline recovers at 0.37 and now matches the text prediction (0.35); the text increment falls to 0.16. What travels across Congresses is party identity, which is time-stable; what does not travel is the text-specific content signal, which is period-bound. This is the two-channel structure that organizes the rest of the paper. Within a period a bill’s position has a durable party-identity component that its sponsors reveal and a content component that only its text reveals; across periods only the party-identity component survives. The recovery model is strongest, and the content channel is legible, precisely within a Congress — which is where the counterfactual roll is defined.

4. Deploying the recovery model on the bills that never reached a vote

The recovery model is trained where cutpoints are observed, on bills that reached a vote, but it must run where they are not, on bills that never did. Nothing guarantees that a model fit to floor bills behaves sensibly on bills the floor never saw; the never-voted bills could be different enough in language that the mapping breaks. Because their true positions are unobserved, the deployment cannot be checked against a gold-standard cutpoint. It can, however, be checked against an independent signal of position that does not use the bill's text at all: the revealed ideology of the coalition that cosponsored it.

I deploy the voted-trained model on the never-voted House bills that attracted a working coalition, the roughly 24,000 bills with at least ten cosponsors, and compare each bill's text-recovered position to the mean ideal point of its cosponsors. The raw association is 0.43 (text placements track the cosponsoring coalition's ideology), but the raw number conflates content with party, since a bill's cosponsors and its text both reflect the sponsor's party. The sharper question is whether text adds position information beyond the party identity of who signed on. Holding the primary sponsor's ideology fixed, the partial association between the text placement and the cosponsor coalition's position is 0.20, against a comparable within-period partial of 0.28 on the voted bills where the same quantity can be computed, roughly three-quarters of the voted-bill signal. The content-beyond-party channel is present on the never-voted bills too, at the same order of magnitude as on the floor, and it is present in every one of the sixteen Congresses. The model does not merely re-read party when it leaves the floor; the within-period content channel extends to the bills the study is about.

Two boundaries of this validation deserve explicit statement. The cosponsorship benchmark exists only where a coalition exists, so the check covers the organized portion of the universe (the bills with at least ten cosponsors) and says nothing directly about the

lone-sponsor introductions in the pool's quieter reaches; for those, the deployment relies on the hardening checks of Section 7, which strip out imputed levels and re-derive the headline from cleavage direction alone. And the check is deliberately period-by-period: each Congress's recovery rests on its own floor-validated mapping, and the partial association never flips sign in any Congress, ranging from 0.15 to 0.36 in magnitude, so the deployment does not borrow strength across eras it cannot certify. Every never-voted bill's roll probability is computed under its own Congress's mapping and floor composition before any average is taken (Appendix Table A4 reports the per-Congress decomposition); that is the footing the within-period estimand requires.

5. The agenda-control wedge

5.1 Validation and magnitude

Before its output can carry a headline, the counterfactual calculation has to reproduce what the floor actually did. Two checks validate it. Fed the true item parameters estimated from the recorded votes, the spatial vote model's predicted pattern of majority rolls tracks the observed pattern across Congresses essentially perfectly (correlation 1.000). The reconstruction is a within-sample consistency check, not a forecast (the parameters come from the very votes being reconstructed), but that is exactly what it needs to be: it certifies that the vote-to-roll machinery introduces no distortion of its own, and that the roll, a joint threshold event on the chamber and the majority party, is generated by the spatial structure rather than by idiosyncratic coalitions the model cannot represent. And the model independently recovers the fact the literature is built on — on bills that reached a passage vote, the majority party was rolled 1.0 percent of the time, the canonical Cox–McCubbins figure. Whatever error the wedge carries must therefore enter through the text imputation, and the comparison is built so that the component of that error common to the two pools cancels: floor bills are evaluated with the same text-imputed

parameters as never-voted bills, and the wedge is the difference. Error specific to one pool would not cancel; the artifact checks of Section 7.1 test for it directly.

Deployed on the full never-voted universe, the recovery model yields the central estimate. Text-imputed roll-proneness across the 122,286 never-voted bills is 13.3 percent; the same model assigns the bills that reached a passage vote 3.7 percent. The agenda-control wedge is +0.096, with a 95 percent interval of +0.084 to +0.104 (Table 3): the bills that never reached the floor carried roughly three and a half times the majority-roll risk of the bills that did. Set against the floor's observed roll rate of one percent, the never-voted universe's 13.3 percent reads as an order of magnitude higher, though the two numbers are on different scales: the model places the same floor bills at 3.7 percent, and it is the imputed-to-imputed difference that the wedge reports. These figures propagate the model's uncertainty — item parameters come from a bootstrap ensemble that disagrees with itself where the text prediction is least reliable, and a smoothed vote step replaces the knife-edge passage threshold; the deterministic version of the same calculation gives +0.114.

The headline is deliberately the conservative reading: it places bills by their text alone, and the content channel is the smaller of the two channels through which a bill's position flows (Section 3.1). Re-imputing positions three ways on the never-voted bills with full sponsorship records makes the decomposition explicit. Sponsor identity alone implies a wedge of +0.32 — the party channel is the larger driver of hidden roll risk, exactly as it is the larger component of bill position — but sponsor-only placement is also the worst-validated of the three (cutpoint recovery 0.53), since it locates every bill at its author's ideology and overstates the extremity of bills written by party-extreme members. Text alone implies +0.11 on the same subset. The joint placement, the best-validated of the three (0.65), puts the wedge at +0.26: below the sponsor-only figure because text pulls party-extreme placements back toward what the bills actually say, above the text-only figure because sponsorship carries real position information that text does not exhaust.

Two numbers therefore state the magnitude. Hidden agenda control amounts to roughly ten points of roll-proneness (the content channel, with all sponsor information withheld), and the best joint estimate is twenty-six points. Adding sponsor information raises the wedge; the content-only headline is a lower bound, not an inflated proxy.

Table 3. The agenda-control wedge, by placement channel.

placement of			
never-voted bills	wedge	cutpoint recovery	reading
text only — the	+0.11	0.61	conservative
content channel			lower bound (the headline)
sponsor identity	+0.32	0.53	largest,
only — the party			worst-validated
channel			
text + sponsor —	+0.26	0.65	best-validated;
joint			the best estimate

Note: Headline full-universe estimate: +0.096 [95% CI +0.084, +0.104], probabilistic (bootstrap-ensemble parameters + smoothed vote step); its deterministic counterpart is +0.114. The channel comparison uses the deterministic vote step on the 121,028 never-voted bills with sponsorship records (where the text-only wedge is +0.112). Floor baseline: text-imputed roll-proneness of voted test bills, 0.037; observed floor roll rate, 0.010.

5.2 Where the screening happens: the gate gradient

A bill that never received a passage vote stalled somewhere, and the version record shows where. Of the never-voted universe, 111,493 bills died in committee, leaving only introduced text; 6,494 cleared committee (a committee reported them to the floor) and then never received a vote, stalling at the scheduling stage; 2,517 passed the House with no recorded vote at all; the remaining 1,782 stalled in mixed cross-chamber stages. Array-

ing roll-proneness across these stages yields the filter's profile (Table 4). Bills that died in committee would have rolled the majority party 13.8 percent of the time; bills that cleared committee but never reached a vote, 7.1 percent; bills that received a recorded passage vote, 3.7 percent; and on those the party actually lost 1.0 percent. Roll risk falls step by step as bills advance: every majority-controlled gate passes forward a pool less threatening to the party than the pool it holds back. The deepest screening occurs at the committee stage, where the cartel account has long located the majority party's first and cheapest veto; the scheduling gate then removes a further increment. Among bills a committee had already judged worth reporting, those that never reached the floor carry nearly twice the roll risk of those that did, a floor-gate wedge of +0.034 whose sampling interval sits clear of the floor baseline.

The profile answers the strongest objection to the wedge: that the never-voted pool is mostly legislative dead weight (commemoratives, duplicates, position-taking vehicles nobody expected to move), so that its roll-proneness reflects what members introduce rather than what the majority party screens. Three facts resist that reading. The committee-stage reservoir is not dead weight: bills that died in committee attracted sixteen cosponsors on average, slightly more than the bills committees reported out. The canonical dead-weight category itself runs opposite to the objection: commemorative bills sit near the bottom of the roll-risk distribution, not its top. The screening signature persists at a gate where a dead-weight story does not apply: committee-reported bills are serious legislation by revealed institutional action, and their blocked remainder still nearly doubles the floor's roll risk. And the gradient's shape is the discriminating pattern of Section 2. A filter indifferent to the majority party's exposure would be flat across gates. A filter screening on whether a bill can pass would produce the opposite gradient, because a bill that cannot command a floor majority cannot roll the majority party: the pool such a filter holds back is safe by construction, and roll risk would rise rather than fall as bills advance. A filter screening on quality would not sort bills by majority-party roll risk at all. Sequential

screening on the roll dimension is what the cartel account predicts; it is what the profile shows.

Table 4. Roll-proneness by the stage at which a bill stalled.

stage reached	n	roll-proneness	wedge vs floor
			baseline
died in committee (introduced only)	111,493	0.138	+0.100
cleared committee, never scheduled	6,494	0.071 [0.067, 0.076]	+0.034
passed House without a recorded vote	2,517	0.082	— (not blocked)
other / cross-chamber stages	1,782	0.125	+0.088
reached a recorded passage vote	—	0.037	baseline
observed majority roll on passage votes	—	0.010	—

Note: Probabilistic (v2) roll-proneness by version stage, proxied by the type of each bill's longest surviving text: a Reported print proves committee clearance, and any

committee-cleared bill whose introduced print happens to be its longest text lands in the died-in-committee stratum, diluting rather than producing the gradient. Intervals reflect bill sampling conditional on the fitted ensemble; parameter error is common across strata and largely differences out of within-run comparisons (the pooled headline interval carries it in full). The passed-without-rollcall stratum was not blocked, and the cross-chamber/other stratum mixes stages including a few hundred enacted bills; excluding both leaves the pooled figure essentially unchanged (13.4 rather than 13.3 percent). Mean cosponsor counts: died-in-committee 15.9, committee-reported 13.6.

6. Who and what the majority party blocks

The wedge and the gate gradient establish that negative agenda control is large and that it operates sequentially, but they leave the substantive question untouched: on which bills, and in which areas of policy, does the majority party's blocking power actually concentrate? This is the question the roll-rate tradition cannot reach, because a blocked bill casts no vote from which its would-be floor position could be read, and it is the question a text-based recovery is built to answer: first the coalitions and sponsors whose bills carry the highest counterfactual roll risk, then the policy-domain map of agenda control, where reading a bill's content rather than its sponsor's identity earns the method its keep.

6.1 Who: organized coalitions, not lone sponsors

The bills the majority party most needs to block are the ones that could actually reach the floor, and a bill's cosponsor count is the visible record of the coalition supporting it. Roll-proneness among the never-voted bills rises monotonically with the size of that coalition. Bills introduced with no cosponsors would have rolled the majority party 8.6 percent of the time; the figure climbs through 12.2, 13.8, 15.7, and 18.5 percent as the coalition widens, reaching 21.0 percent for the bills carrying fifty or more cosponsors (Table 5). The

most roll-prone blocked bills are not obscure lone-sponsor proposals; they are broadly cosponsored measures with organized demand behind them, which is exactly what the cartel logic of Section 2 predicts negative agenda control exists to contain.

Table 5. Roll-proneness of never-voted bills by size of cosponsoring coalition.

cosponsors	roll-proneness
0	0.086
1–4	0.122
5–9	0.138
10–24	0.157
25–49	0.185
50+	0.210

Note: Probabilistic (v2) roll-proneness over the full never-voted pool ($n = 122,286$), by cosponsor count — the same machinery as the headline and Table 6. The deterministic estimate runs uniformly higher (10.2 to 23.3 percent) with the identical monotone ordering.

This gradient is also the answer to the objection that these bills were never going anywhere in the first place. A bill introduced with fifty or more cosponsors under a majority party that opposes it is not the absence of legislative demand; it is precisely what blocked-but-organized demand looks like. The coalition assembled, the demand for floor consideration was organized and visible, and the vote never came. The blocking is the observation, not the missing datum. And the roll risk of these bills climbs with coalition size, which is the fact that matters for the estimand: whatever mix of sincere lawmaking and position-taking produces these introductions, the pool the majority party holds back is weighted toward the bills that would divide it on the floor, not toward noise.

Who sponsors a bill remains, as elsewhere, a powerful observable predictor of whether it would roll the majority party, and here the text measure only confirms the party signal the literature already reads. Minority-sponsored never-voted bills carry a roll-proneness

of 0.221, against 0.094 for majority-sponsored ones, and minority sponsors account for 42.4 percent of the pool. Bills the model flags as would-roll skew toward out-party sponsors — their mean first-dimension NOMINATE score is -0.119 , on the liberal side, against $+0.029$ for bills not flagged, with predicted rolls concentrating in Republican-majority Congresses. None of this is novel; it is the Cox–McCubbins party signal, recovered as validation rather than as discovery. The contribution lies not in confirming that blocked bills lean toward the minority but in showing, next, that agenda control has a structure the sponsor’s party alone does not supply.

6.2 What: the policy-domain map

Where the majority party’s blocking power concentrates, bill for bill, across the substance of policy is the finding only text can produce, since a bill’s sponsor tells you who wrote it but only its words tell you what it is about. Sorting the never-voted universe by policy domain and averaging the counterfactual roll probability within each yields a map with wide and orderly variation (Table 6). Against a floor baseline of 0.037 (the same model’s roll-proneness for the bills that did reach a vote), the most roll-prone domains cluster near a fifth of bills rolling the majority party: Families at 0.223, Emergency Management and Labor and Employment at 0.205, Arts, Culture, and Religion at 0.200, Civil Rights and Liberties at 0.198, Transportation and Public Works at 0.195, Agriculture and Food at 0.192, Environmental Protection at 0.177, and Education at 0.175.

Health, the largest single domain by volume, sits mid-table at 0.172, so the domains above it are more roll-prone than the domain that holds the largest share of the pool. At the low end the ordering is just as legible. Commemorations register at 0.097, Economics and Public Finance at 0.087, Law at 0.076, Congress at 0.064, Taxation at 0.058, and Foreign Trade and International Finance at 0.033 — a low-risk tail that approaches, and in trade falls below, the roll risk of the bills the majority party actually let onto the floor.

Table 6. Roll-proneness of never-voted bills by policy domain.

policy domain	roll-proneness	n
Families	0.223	821
Emergency Management	0.205	1,375
Labor and Employment	0.205	2,965
Arts, Culture, Religion	0.200	613
Civil Rights and Liberties	0.198	1,225
Transportation and Public Works	0.195	4,075
Agriculture and Food	0.192	2,595
Environmental Protection	0.177	2,879
Education	0.175	4,909
Health (<i>mid-table, largest domain</i>)	0.172	12,697
...		
Commemorations	0.097	3,340
Economics and Public Finance	0.087	1,881
Law	0.076	1,369
Congress	0.064	5,704
Taxation	0.058	10,769
Foreign Trade and International Finance	0.033	5,859

Note: Probabilistic (v2) OOD-aware roll-proneness (bootstrap-ensemble parameters + smoothed vote step), by policy domain; selected top and bottom rows shown. Floor baseline 0.037. The most out-of-range domains attenuate under the OOD-aware estimate relative to the deterministic one (Families 0.270 $\rightarrow$ 0.223), but the ordering is essentially unchanged (rank agreement $\rho = 0.955$).

The pattern is not a single “social versus fiscal” contrast but two coherent concentrations. Agenda control concentrates most strongly on the social, labor, and civil-rights domains — Families, Labor and Employment, Civil Rights and Liberties — the questions on which a majority party’s members most visibly divide by district. It concentrates nearly

as strongly on the constituency-facing distributive domains: Transportation and Public Works, Agriculture and Food, Emergency Management, the particularistic-spending areas where a bill benefiting one set of districts asks members from other districts to pay for it.

That the distributive domains rank high is not a complication of the account but a confirmation of it, since particularistic spending splits a majority party whose members represent different places exactly as social cleavages do. Reading the map as purely social or redistributive would miss half of what the majority party guards; the constituency-facing spending domains belong on the same list, and they get there through the same mechanism of intra-party geographic conflict. The weak end is equally coherent. Taxation, Congress, Foreign Trade, Economics and Public Finance, and Law occupy the low end; read through the cartel account, these are the domains on which the majority party had little to fear from a floor vote, and so little need of a scheduling veto.

The decisive question is whether this map reflects what bills are about or merely who writes them, the discriminating test of Section 2 between the content account and the party-composition alternative. Three results place the weight on content. The cross-domain correlation between roll-proneness and minority-sponsor share is +0.445, so party composition accounts for only about a fifth of the cross-domain variance in roll-proneness, and four-fifths lies elsewhere. Holding sponsor party fixed at the minority, the domain ranking barely moves: the rank correlation between the all-bills and minority-only orderings is 0.947. Within minority-sponsored bills alone, domain still swings roll-proneness across thirty-two points, from Labor and Employment at 0.360 down to Foreign Trade at 0.044 (Appendix Table A5).

A secondary party interaction does exist and deserves explicit statement. The minority-sponsored and majority-sponsored domain rankings agree only at $\rho = 0.336$, so the domain ordering is not stable across sponsor party. The direction of the interaction is sensible rather than contradictory: never-voted bills run more roll-prone under minority

sponsorship, as the sponsor-party contrast above showed, and a socially charged bill written by an out-party member sits further from the majority party than the same subject written from within. But the interaction overlays a structure it does not create. The map's claims rest on the minority-sponsored stratum, where most of the blocked pool's roll risk is concentrated, and within that stratum domains sort roll risk on their own content, which is precisely what a party-composition story predicts should not happen.

One independent check that a reader might expect, comparing the domain map against how divisively each domain's bills actually split the floor, cannot serve as corroboration here, and the reason is the paper's argument in miniature. Where blocking succeeds, a domain's most divisive bills never reach the floor, so the divisions observed on passed bills understate the domain's underlying divisiveness by exactly the amount the majority party screened out. The floor-division ranking agrees with the recovered map only weakly, at a rank correlation of +0.27, consistent with heavy and content-targeted screening rather than evidence against the map. Floor-based checks cannot adjudicate this map, because the selection they would have to see through is the very object under study.

Face validity runs the right way at the boundaries. Commemorations, the canonical dead-weight category, sits near the bottom of the roll-risk distribution: the model rates the bills nobody expects to divide the chamber as low-risk, consistent with Section 5.2's rebuttal of the junk-bill objection. That the model reaches this verdict on its own, from the text of thousands of commemorative resolutions and without any hand-coding of which categories should count as trivial, is the kind of check the method has to pass. The domains it rates as most dangerous, the socially charged core described above (families, labor, civil rights), are, in turn, the ones a reader of American legislative politics would name first if asked which bills a cross-pressured majority party would least want on the floor.

The upshot is a claim the roll-rate tradition is structurally unable to make. Negative

agenda control is organized by what bills are about, concentrated on the domains that split the majority party from within, and this policy map lies out of reach for any measure that reads a bill only through its sponsor's party and the votes it happened to receive. The floor-only view can say that the majority party is rarely rolled and that blocked bills tend to come from the minority party; it cannot say that the blocking falls on families and labor and roads rather than on taxes and trade, because the never-voted bills supply no votes, and so no measure of what each one would have done to the majority party on the floor. Recovering position from text is what makes the substance of party power visible.

7. Robustness and scope

7.1 The wedge is not an artifact

The first concern is that the model invents extremity on unfamiliar text, placing never-voted bills at implausible cutpoints precisely where their words are least informative. The headline is built to absorb this. It is the out-of-distribution-aware probabilistic estimate: item parameters come from a bootstrap ensemble that disagrees with itself where the text is least reliable, and a smoothed vote step replaces the knife-edge passage threshold. That construction pulls the estimate down from its deterministic counterpart of +0.114 to +0.096. To see whether even this rests on imputed extremity, I strip the levels out two ways. Re-centering never-voted imputed positions on the voted distribution leaves a wedge of +0.092; using only the sign of each bill's discrimination parameter with magnitudes typical of the gold-standard estimates leaves +0.097. The wedge is driven by which side of the cleavage a bill is on, not by how extreme the model places it.

A second concern is that the whole result hinges on the text machinery. If the phenomenon appeared only under one particular text representation, it would be a property of my model rather than of the House. A completely text-free deployment answers this directly. Placing each never-voted bill by its cosponsorship coalition alone, using no words

at all, yields a wedge of +0.35, larger than the text headline. Text is the conservative channel, as Section 5 makes explicit, and the majority-screening pattern appears through an entirely independent signal. The phenomenon does not depend on the text model.

A third concern starts from an admission the paper makes about itself. Section 3 reports that per-Congress recovery is genuinely heterogeneous, ranging from 0.21 to 0.76, which invites the worry that the wedge is manufactured in the Congresses where the model works worst. It is not. Across Congresses the wedge is weakly positively correlated with recovery quality, measured as within-Congress recovery on voted bills, an in-distribution proxy for how well the model reads each Congress ($r = +0.15$; rank correlation +0.20). In the six Congresses with the highest within-Congress recovery (0.60 and above), the never-voted-weighted wedge is +0.107, against +0.082 in the ten below. The screening signal is, if anything, sharper where the model reads bills more accurately, which is the opposite of what an imputation-noise story requires.

A fourth concern is that later stages of the process select on drafting quality rather than on roll risk, so the gate gradient of Section 5.2 might reflect competence rather than exposure. Two facts separate the two. The gate gradient persists within the committee-reported stratum, among bills of revealed institutional seriousness that a committee judged worth sending forward, where a quality story has little room to work. And Commemorations, the canonical dead-weight domain, is among the least roll-prone domains in the pool (Section 5.2): if the score were proxying low quality or low seriousness rather than partisan exposure, the least serious domain in the corpus would rank high on it, and it ranks low. Screening on general passability would sort bills by quality; it would not sort them specifically by the majority party's roll risk, which is what the profile does.

The recovery model carries its own validation battery, reported with the model in Section 3 and in the Online Appendix rather than here. It includes a permutation test that rejects the sign-label null at $p \approx 5 \times 10^{-4}$, an independent NOMINATE-midpoint gold standard that recovers at 0.59 within and 0.34 across periods against 0.69 and 0.36

on the primary gold standard, the same pattern at modestly lower strength, genuine two-dimensional recovery, and a full-text long-context fine-tune that improves on the bag-of-words model by only +0.03 within period. None of it disturbs the recovery the wedge is computed on.

7.2 What the wedge cannot identify

The wedge is a within-period quantity by design, and three questions readers will reasonably bring to it lie outside what this design can answer. Naming those boundaries is part of stating the estimand precisely. The measure is cross-sectional within a Congress because the content channel is period-bound. Section 3.1 shows the text increment falling from 0.45 to 0.16 across Congresses: only the party-identity channel travels over time, while the content signal that lets me read a bill's position beyond its sponsor is period-bound. The wedge inherits that scope. It compares never-voted against voted bills within the same floor geometry, and it does not extend to comparisons of one geometry against another.

For that reason I make no claim about whether agenda control grew with polarization, which is the question the time series most invites. The per-Congress wedge tracks which party holds majority control, averaging 0.125 in Republican-majority Congresses against 0.027 in Democratic ones, so any upward or downward drift confounds agenda control with majority identity. The confound compounds because the Democratic-era never-voted pools sit far outside the model's training range, with a mean imputed cut-point of +1.18, where recovery is least reliable. The series is not cross-era comparable, and I do not interpret it.

I also cannot identify whether agenda control shields electorally marginal majority members, a natural extension of the cartel logic that this design is not equipped to test. A majority member's exposure to the blocked would-roll bills is nearly deterministic in ideological moderation (within Congresses, exposure and ideological extremity correlate at

−0.96), and the electorally marginal are disproportionately the moderates, so the measure cannot separate protection from position. Controlling for ideology, the apparent protection gradient collapses to almost nothing, as that collinearity guarantees it must. Answering the member-protection question needs a genuine identification strategy, within-member panels or instruments for exposure, which I leave to future work.

Finally, the wedge is a footprint, not an attribution. It characterizes what the chamber’s filter screens on, taken over the universe of bills introduced; it does not assign any single bill’s death to a designated blocking act. A bill can stall for many reasons that have nothing to do with a majority-party veto, and the estimand discipline of Section 2 keeps the claim at the level of the process rather than the individual case. What the wedge measures is the shape of the screen, and that is the quantity I interpret.

8. Conclusion

The bills the majority party never let reach a vote carried several times the majority-roll risk of the bills it did; that screening is visible at every majority-controlled gate, from committee through scheduling, and it concentrates in the policy domains that split the majority party. Read together, these three facts relocate negative agenda control from the floor, where it is inferred, to the reservoir of blocked legislation, where it operates.

The immediate lesson is for how the field measures party power. Floor-based measures — roll rates, ideal-point scalings, the pivot models built on recorded votes — take the observed floor as their evidentiary universe, but that floor is the filtered output of the very power they are trying to detect. The consequence is not noise but systematic bias: the measure understates agenda control and mislocates it, reading party power as weak precisely where the theory says it is strongest. A one percent floor roll rate coexists with a never-voted reservoir whose model-read roll risk runs an order of magnitude higher, and no quantity computed on the floor can register that gap.

The direction of the bias is not incidental. Because the filter removes the most threatening bills first, what reaches the floor is the least informative slice about the majority party's exposure, so a measure built only on recorded votes converges on a picture of a party that rarely faces danger, the exact impression the gatekeeping produces. When the observed sample is selected by the mechanism under study, more careful measurement of the sample does not fix the problem; recovering the unobserved cases does.

For the long argument between the cartel and majoritarian accounts, the counterfactual roll supplies a discriminating pattern that floor-based tests could not deliver in principle. A majoritarian filter conditions floor access on a bill's capacity to command a chamber majority and is otherwise indifferent to which party the bill endangers. What I find instead is exclusion that tracks majority-party roll risk conditional on passability, that operates at each successive gate, and that retains its policy structure with sponsor party held fixed. Each of those features is what partisan gatekeeping predicts and a majority-rule filter should not generate. This is within-period, cross-sectional evidence, not a design that observes the same institution before and after a change in control; I claim that the footprint of the House filter has the shape the cartel account requires, not that I have isolated an act of blocking.

The results also carry a caution for text-as-data measurement of political content. Within a Congress, bill text reveals position beyond the identity of who sponsored it: the content channel is real and separable from the party channel. But it is period-bound. Across Congresses the content signal decays while party identity persists, because a bill's words map onto a chamber geometry that is specific to its moment, whereas the partisan meaning of authorship is stable. Any measurement that trains a text model in one institutional era and deploys it in another inherits this asymmetry. What travels is party; what is local, and must be estimated where it is used, is content. The recovery here is trustworthy exactly because it is confined to the period in which the mapping was learned.

Two extensions follow directly. The first is the Senate, where the filter is built from dif-

ferent parts (floor recognition, the sixty-vote cloture threshold, unanimous-consent practice), so the counterfactual roll would recover a differently shaped agenda control; the model already recovers the positions of Senate-origin bills that reached the House floor about as well as House-origin ones. The second is the question this design cannot answer: whether agenda control is aimed at protecting electorally marginal members. Establishing that requires following individual legislators over time — a within-member panel asking whether the blocking of the bills that would cross-pressure a member intensifies as her seat grows competitive. The cross-sectional footprint is silent on that attribution.

Three limits bound the claims. The estimand is within-period throughout, since the content channel does not extrapolate across Congresses, and every wedge should be read as a snapshot of one chamber's filter rather than a trend. Never-voted status marks the output of the majority party's filter, not a bill-by-bill veto; the wedge is the footprint of that process over the universe of introduced bills, not the causal effect of a designated act of exclusion. And recovery quality is heterogeneous across Congresses, so the pooled figures should be understood as an average over eras the model reads better and worse.

The study of party power has largely been the study of what majority parties let happen: which bills passed, who won the recorded votes, how rarely the majority party lost. That evidence was always the residue of a prior decision made offstage, about which bills would be permitted to reach a vote at all. The bills that never got there are where the power actually operates, and until now they left no measurable trace. Reading their would-be floor positions from their text turns that silence into evidence — the majority party's most consequential decisions, the ones it made by ensuring nothing happened, can finally be read alongside the votes it chose to hold.

References

- Aldrich, J. H. (1995). *Why Parties? The Origin and Transformation of Political Parties in America*. University of Chicago Press, Chicago.
- Clinton, J., Jackman, S., and Rivers, D. (2004). The statistical analysis of roll call data. *American Political Science Review*, 98(2):355–370.
- Cox, G. W. and McCubbins, M. D. (1993). *Legislative Leviathan: Party Government in the House*. University of California Press, Berkeley.
- Cox, G. W. and McCubbins, M. D. (2005). *Setting the Agenda: Responsible Party Government in the U.S. House of Representatives*. Cambridge University Press, New York.
- Delano, R., Rudkin, A., and Kim, I. S. (2025). Bulk ingestion of congressional actions and materials dataset. *Scientific Data*, 12:1567.
- Den Hartog, C. and Monroe, N. W. (2011). *Agenda Setting in the U.S. Senate: Costly Consideration and Majority Party Advantage*. Cambridge University Press, New York.
- Gerrish, S. and Blei, D. M. (2011). Predicting legislative roll calls from text. In *Proceedings of the 28th International Conference on Machine Learning (ICML)*.
- Grimmer, J. and Stewart, B. M. (2013). Text as data: The promise and pitfalls of automatic content analysis methods for political texts. *Political Analysis*, 21(3):267–297.
- Imai, K., Lo, J., and Olmsted, J. (2016). Fast estimation of ideal points with massive data. *American Political Science Review*, 110(4):631–656.
- Krehbiel, K. (1998). *Pivotal Politics: A Theory of U.S. Lawmaking*. University of Chicago Press, Chicago.
- Poole, K. T. and Rosenthal, H. (1997). *Congress: A Political-Economic History of Roll Call Voting*. Oxford University Press, New York.

Rohde, D. W. (1991). *Parties and Leaders in the Postreform House*. University of Chicago Press, Chicago.

Yano, T., Smith, N. A., and Wilkerson, J. D. (2012). Textual predictors of bill survival in congressional committees. In *Proceedings of the 2012 Conference of the North American Chapter of the Association for Computational Linguistics (NAACL-HLT)*.

A Data construction and the version-at-vote rule

The analysis links two sources: Voteview roll-call records for the House, Congresses 103–118, and the BICAM corpus of version-level bill text (Delano et al., 2025). Of 20,236 House roll calls in the period, 8,324 are eligible final-passage votes (question types where the vote is on the bill or resolution text itself: passage, the suspend-the-rules passage variants, resolution agreement, and conference-report adoption) on 8,021 distinct bills; pre-decision text is recovered for 8,274 roll calls (99.4 percent), with Congressional Research Service summaries available for 96.6 percent and coverage of the close-vote stratum at 99.3 percent. The never-voted universe comprises 122,286 House bills introduced with text that never received a final-passage roll call.

Because a bill exists in multiple textual versions, version selection is where outcome leakage can enter: the “Engrossed in House” version is produced after passage, so any rule that selects the latest available text conditions on the outcome. A naive “latest version dated on or before the vote” rule selects post-passage text in 84 percent of audited cases, because engrossment carries the vote date. The rule used throughout (Rule A) instead selects the latest pre-decision version: Introduced, Reported, Referred, Placed on Calendar, or Committee Print, dated on or before the vote. In a fifty-vote manual audit, Rule A selects a genuinely pre-decision version in 46 of 50 cases, with the four exceptions arising from cross-chamber edge cases that are flagged and retained; a second audit after the BICAM link found 47 of 50 clean.

Two sample restrictions apply to the item-recovery analyses. Cutpoints are defined as the ratio of item parameters and are unstable when the discrimination is near zero or the vote is near-unanimous, so item recovery is evaluated on well-identified items: absolute discrimination above 0.3 and a yea-share strictly inside the unit interval. This retains 4,039 of 8,274 roll calls (48.8 percent); the excluded half consists of lopsided or ideologically unstructured votes for which a cutpoint is not a meaningful quantity. Cutpoints are winsorized to the interval from minus two to plus two before evaluation; 25.3 percent of

well-identified cutpoints sit at the bound (29 percent in the temporal test set, asymmetrically at the upper bound), which is why rank correlations, which are insensitive to the capping, are reported alongside Pearson correlations in the main text.

B The gold target and its independent validation

The gold parameters are estimated separately within each Congress by expectation-maximization item response theory on the full House vote matrix. Two independent checks against NOMINATE, an entirely separate estimation tradition on the same votes, calibrate their credibility. Member ideal points correlate with first-dimension DW-NOMINATE at a mean of 0.978 across the sixteen Congresses (range 0.967 to 0.984). Item cutpoints correlate with the NOMINATE cutting-line midpoint at a mean of 0.751, with per-Congress values reported in Table A1; the agreement is high in most Congresses and weakest in the 117th (0.52) and 118th (0.56), with the 115th close behind (0.59), which is one reason per-Congress recovery figures in those Congresses carry extra target noise.

Table A1. Gold-target validation against NOMINATE, by Congress.

Congress	cutpoint vs. NOMINATE	ideal points
	midpoint	vs. DW-NOMINATE
103	0.86	0.98
104	0.75	0.97
105	0.76	0.97
106	0.79	0.98
107	0.83	0.98
108	0.83	0.98
109	0.81	0.98
110	0.82	0.98
111	0.83	0.98
112	0.68	0.98
113	0.70	0.98
114	0.95	0.98
115	0.59	0.97
116	0.73	0.98
117	0.52	0.98
118	0.56	0.98

Note: Means 0.751 and 0.978. Cutpoint agreement between two independent estimators of the same latent quantity; the weaker late-Congress values (115th, 117th) indicate lower gold-target reliability there rather than recovery-model failure.

C The recovery-model validation battery

Four checks, all run with a train-only TF-IDF vocabulary fit so that no vocabulary information leaks from test to train, on both frozen splits.

First, an independent gold. Replacing the emIRT cutpoint with the NOMINATE

cutting-line midpoint, a target the model was never trained toward, text recovers position at 0.59 on the random split and 0.34 on the temporal split, tracking the primary-gold figures of 0.69 and 0.36. The recovery is not an artifact of the emIRT machinery. The train-only refit also reproduces the published random-split headline exactly (0.690), confirming that no vocabulary leakage inflates it.

Second, a permutation null. Shuffling the assignment of texts to items and re-fitting, the null recovery distribution is centered at zero with a 97.5th percentile of 0.07 (random) and 0.06 (temporal); the observed 0.690 and 0.355 sit far outside it, with permutation p-values of 0.0005 on both splits. The temporal recovery is modest in absolute terms but is unambiguously signal.

Third, two-dimensional recovery. The chamber’s cutting lines are genuinely two-dimensional: only 23 percent of true cutting-line normals lie within fifteen degrees of the first-dimension axis. Text recovers the second-dimension tilt at 0.67 on the random split, and the mean angular error between predicted and true cutting-line normals is 26 degrees against a 45-degree random baseline. The recovery is not a degenerate first-dimension artifact. Both quantities decay across periods (0.17 and 36 degrees), consistent with the period-bound content channel of Section 3.1.

Fourth, subgroup stability. Within the random split, recovery holds across vote-question types (0.56 to 0.74) and across chamber of origin (0.69 House-origin, 0.66 cross-chamber); no subgroup collapses.

D Per-Congress recovery heterogeneity

The pooled random-split recovery of 0.69 exploits between-Congress cutpoint spread, so the main text foregrounds within-Congress figures. Table A2 reports within-Congress recovery estimated one Congress at a time on the random split’s test bills.

Table A2. Within-Congress cutpoint recovery by Congress (random split), with 95 percent

confidence intervals.

Congress	r	95% CI	n	Congress	r	95% CI	n
103	0.25	[−0.07, 0.53]	39	111	0.66	[0.49, 0.78]	64
104	0.36	[0.04, 0.61]	37	112	0.60	[0.37, 0.77]	43
105	0.21	[−0.11, 0.49]	39	113	0.66	[0.45, 0.80]	43
106	0.37	[0.09, 0.60]	44	114	0.41	[0.18, 0.60]	62
107	0.54	[0.16, 0.78]	23	115	0.37	[0.09, 0.60]	45
108	0.65	[0.42, 0.81]	37	116	0.51	[0.25, 0.70]	44
109	0.34	[0.08, 0.56]	52	117	0.38	[0.20, 0.53]	103
110	0.76	[0.63, 0.84]	66	118	0.72	[0.57, 0.82]	62

Note: Intervals are 95 percent Fisher-z confidence intervals. With test sets of 23–103 bills the individual estimates are imprecise, but a homogeneity test rejects a common within-Congress recovery ($Q(15) = 42.8$, $p < 0.001$): the heterogeneity is real, while individual Congress rankings remain uncertain. Restricting both training and test to the same recent era (Congresses 116–118) gives the cleanest within-period benchmark, 0.60 ($n = 209$); a leave-one-Congress-out transfer averages 0.45. Section 7.1 shows the wedge is, if anything, larger in the Congresses where recovery is better.

E Model capacity: the full-text fine-tune

To test whether the bag-of-words model leaves recoverable signal on the table, two 4,096-token Longformer encoders, one general-domain and one legal-domain, were fine-tuned end-to-end on the full bill text with a multitask regression head (cutpoint, discrimination, yea-share), on both frozen splits. Because both encoders share the same length budget, their comparison isolates domain pretraining with length held constant. Truncation at 4,096 tokens affects 31 percent of documents for the vanilla encoder and 23 percent for the legal encoder.

Table A3. Full-text fine-tune versus the bag-of-words model (cutpoint recovery).

model	within period	across periods
bag-of-words + ridge (main text)	0.69	0.36
Longformer, general-domain	0.72	0.38
Longformer, legal-domain	0.71	0.44

Note: Within-period gains of at most +0.03: the linear model sits near the recoverable ceiling. The temporal decay survives fine-tuning at scale, so the period-boundedness of the content channel is intrinsic, not a capacity limit. The legal-domain encoder’s temporal edge over bag-of-words (+0.09, unpaired test $p = 0.022$) is real but modest; second-dimension recovery under fine-tuning is negative across periods, confirming that the chamber’s two-dimensional geometry does not extrapolate.

F The wedge machinery, v1 and v2

The deterministic (v1) pipeline predicts each bill’s cutpoint and discrimination from text, computes each member’s yea probability from the item response function at the member’s gold ideal point, and declares a bill passed when its expected yea count exceeds half the chamber and a majority roll when it passes while the majority party’s expected internal yea share falls below one half. Two features make v1 brittle: a knife-edge threshold that flips a bill’s roll status on small cutpoint errors, and single point predictions on out-of-distribution text.

The probabilistic (v2) estimate, used for the headline, corrects both. Parameter uncertainty is propagated by a bootstrap ensemble of forty text models (training rows resampled, ridge refit at fixed penalty and vocabulary); on out-of-distribution bills the ensemble disagrees with itself, widening the implied uncertainty exactly where the point prediction is least trustworthy. The vote step is smoothed by treating total yeas and majority-party yeas as sums of independent Bernoulli votes, jointly approximately bivariate normal, and

computing the exact orthant probability that a bill passes while a majority of the majority party opposes, by Gauss–Legendre quadrature; no Monte Carlo noise enters. The full-universe wedge is +0.096 (interval +0.084 to +0.104) under v2 against +0.114 under v1.

Table A4 reports the per-Congress v2 decomposition.

Table A4. Per-Congress v2 wedge (never-voted roll-proneness minus floor baseline).

Congress majority		wedge	interval	Congress majority		wedge	interval
103	D	+0.03	[+0.01, +0.06]	111	D	+0.03	[+0.02, +0.05]
104	R	+0.06	[−0.02, +0.12]	112	R	+0.15	[+0.10, +0.19]
105	R	+0.10	[+0.07, +0.14]	113	R	+0.15	[+0.12, +0.17]
106	R	+0.09	[+0.05, +0.13]	114	R	+0.19	[+0.16, +0.22]
107	R	+0.13	[+0.09, +0.15]	115	R	+0.19	[+0.17, +0.22]
108	R	+0.13	[+0.10, +0.16]	116	D	+0.03	[+0.01, +0.04]
109	R	+0.04	[+0.02, +0.06]	117	D	+0.00	[+0.00, +0.01]
110	D	+0.04	[+0.02, +0.05]	118	R	+0.15	[+0.10, +0.21]

Note: Ensemble percentile intervals. The per-Congress series tracks which party holds majority control (Republican-majority mean 0.125, Democratic-majority mean 0.027) and the Democratic-era never-voted pools sit far outside the recovery model’s training range; Section 7.2 explains why the series is not interpreted over time.

Three hardening checks address the residual concern that the wedge is driven by

imputed extremity. Re-centering the never-voted imputed cutpoint distribution on the voted distribution (a location shift of -0.48) leaves a wedge of $+0.092$. Replacing text entirely with a cosponsorship-based placement, available for the 87,362 never-voted bills with matched cosponsors, yields $+0.352$ against its own floor baseline; the cosponsorship channel independently confirms the level shift in never-voted positions ($+0.48$ on its own scale). Replacing each bill's imputed discrimination with its sign at a magnitude typical of the gold-standard estimates yields $+0.097$. The wedge is driven by which side of the cleavage bills are on, not by imputed magnitudes.

G The policy-domain map in full

Table A5 reports the complete domain table underlying Table 6 of the main text: v2 roll-proneness by policy area for every domain with at least 300 never-voted bills, with the domain's minority-sponsor share and mean imputed cutpoint.

Table A5. Never-voted roll-proneness by policy domain, complete (v2).

policy domain	roll-proneness	n	minority share	mean cutpoint
Families	0.223	821	0.50	+1.20
Emergency Management	0.205	1,375	0.43	+1.42
Labor and Employment	0.205	2,965	0.46	+0.90
Arts, Culture, Religion	0.200	613	0.47	+1.01
Civil Rights and Liberties	0.198	1,225	0.52	+0.29
Transportation and Public Works	0.195	4,075	0.44	+1.00
Agriculture and Food	0.192	2,595	0.44	+1.06
Environmental Protection	0.177	2,879	0.40	+0.83
Education	0.175	4,909	0.51	+1.18
Native Americans	0.174	1,210	0.40	+1.13
Health	0.172	12,697	0.43	+1.05

Public Lands and Natural Resources	0.163	5,757	0.38	+1.01
Animals	0.161	565	0.35	+0.75
Science, Technology, Communica- tions	0.160	2,329	0.44	+1.16
Housing and Community Development	0.160	1,598	0.42	+1.61
Water Resources	0.153	1,206	0.40	+1.27
Development International Affairs	0.146	7,663	0.45	+0.37
Social Welfare	0.143	1,930	0.46	+0.80
Crime and Law Enforcement	0.142	5,958	0.47	+0.12
Immigration	0.139	3,055	0.47	+0.31
Private Legislation	0.133	1,210	0.50	+0.39
Commerce	0.131	2,786	0.45	+1.35
Government Operations and Politics	0.128	8,079	0.42	+0.48
Finance and Financial Sector	0.118	3,658	0.40	+0.46
Energy	0.114	2,754	0.41	+0.68
Armed Forces and National Security	0.111	8,211	0.46	+0.53
Sports and Recreation	0.103	461	0.47	+0.57
Commemorations	0.097	3,340	0.49	+0.23
Economics and Public Finance	0.087	1,881	0.40	+0.32
Law	0.076	1,369	0.36	+0.19
Congress	0.064	5,704	0.32	-0.03
Taxation	0.058	10,769	0.38	-0.09
Foreign Trade and International Finance	0.033	5,859	0.40	+0.30

Note: Floor baseline 0.037. Under the deterministic estimate the most out-of-range domains run higher (Families 0.270) but the ordering is essentially unchanged (rank agreement 0.955 between the two estimates). Within the minority-sponsored stratum alone the

ranking is essentially the same (rank correlation 0.947 with the all-bills ordering), with Labor and Employment at 0.360 and Foreign Trade at 0.044.

H Gate strata detail

The version-stage strata of Section 5.2 are defined by the House version type of each never-voted bill's retained text: introduced-only bills died in committee ($n = 111,493$); "Reported in House" bills cleared committee and stalled at scheduling ($n = 6,494$); a small stratum passed the House with no recorded vote ($n = 2,517$, voice votes and unanimous consent, not blocked); the remainder are cross-chamber and other stages ($n = 1,782$, a mixed stratum that includes a few hundred enacted bills). Because the text extraction retains one version per bill (the longest), stage is a proxy: the load-bearing direction is safe, since a bill whose retained text is a Reported print necessarily cleared committee, while a committee-cleared bill whose introduced print happens to be longer lands in the died-in-committee stratum — and because committee-cleared bills are less roll-prone, such contamination dilutes the gradient rather than producing it. Stratum intervals reflect bill sampling conditional on the fitted ensemble; parameter uncertainty is shared across strata within a run and largely differences out of the gradient comparison, while the pooled headline interval of Section 5.1 carries it in full. Mean cosponsor counts are 15.9 for the committee-died stratum and 13.6 for the committee-reported stratum. Measured against the pooled floor baseline, the reported-stratum roll-proneness exceeds it in ten of sixteen Congresses; five of the six exceptions are Democratic-majority Congresses in which the never-voted pool as a whole is least roll-prone, so the shortfall reflects the era rather than the stratum, and the pooled reported-stratum interval excludes the floor baseline. Excluding the passed-without-rollcall stratum moves the pooled never-voted roll-proneness from 13.3 to 13.4 percent.