

PACs Follow the Winner: Access-Seeking Across the Donor Spectrum in U.S. House Campaign Contributions

Kisoo Kim*

July 2, 2026

Abstract

Whether organized political giving buys legislative access or expresses policy preferences has framed three decades of debate. I apply a sharp regression discontinuity to the win-loss margin from the donor's perspective, using 1.7 million committee-by-race observations of U.S. House elections, 1990-2022, and extend the design to individual donors. When a committee's previously supported candidate narrowly loses, the committee redirects giving toward the winner, curtails support for the loser, leaves total giving unchanged, and reorients its partisan composition toward whichever party wins—the joint signature of access-seeking. On the subsample of committees sponsored by publicly traded firms, where lobbying is observable, the redirection is matched by a parallel, winner-specific jump in directed lobbying. Ideological committees follow the winner at least as strongly as business committees; small individual donors follow weakly and do not realign their partisan giving. Winner-following is a property of strategic donors, not of corporate status.

Keywords: campaign finance; interest groups; access; regression discontinuity; U.S. Congress; lobbying

*Email: kisoo@virginia.edu

1 Introduction

When Charlie Wilson — a two-term Democrat from Ohio’s Sixth District — lost his seat to Republican Bill Johnson in 2010, seventy-five political action committees that had given \$279,000 to Wilson’s reelection campaign began funding Johnson, giving the new Republican incumbent \$390,000 over the next cycle. The cleaner test came in 2012, when Wilson ran against Johnson for the same seat: the same seventy-five committees gave his comeback bid \$24,000 — less than a tenth of what they had given him as the incumbent — while funding Johnson at more than sixteen times that level. The pattern recurs cycle after cycle: the committees that finance U.S. House campaigns behave as if the office matters more than its occupant, rerouting money toward winners they did not previously support and away from candidates they once funded heavily. What that behavior reveals about the motive behind organized giving has framed three decades of debate.

One view holds that contributions purchase legislative access — face time, information, and attention from whoever holds the seat — a tradition running from Welch (1980) and Cassie and Thompson (1998) to Phillips (2024). A second view, articulated in Ansolabehere et al. (2003) and reflected throughout the cross-sectional CFscore literature (Bonica, 2014, 2018), treats contributions as expressions of ideology: donors back candidates whose positions match their own, and aggregate giving is too small to be rational policy investment. A third combines both: committees expand their budgets to cultivate the new winner while maintaining old commitments. The three views collide the moment an election flips the office-holding identity, and the literature has lacked a design clean enough to put them side by side.

I ask: when a committee’s previously supported candidate narrowly loses a U.S. House election, what does the committee do next? The question admits a sharp answer because the three views imply distinct sign patterns across four behavioral outcomes measured in the cycle after the focal race. Access-seeking predicts within-budget substitution: redirection toward the winner, collapsing loyalty to the loser, an unchanged total, and

partisan reorientation that tracks the winner’s party. Ideology predicts behavioral continuity, because preferences do not move when election outcomes do; the hybrid predicts redirection financed by budget expansion. No single outcome distinguishes the three; the joint signature does.

Identification rests on a sharp regression discontinuity in the signed win-loss margin from the donor’s perspective: positive when the committee’s previously supported candidate wins the focal race, negative when that candidate loses. As-if-random assignment of close-race outcomes (Lee, 2008; Calonico et al., 2014) generates quasi-experimental variation in whether the committee’s preferred candidate holds office at all — the office-versus-no-office margin. The data cover 1.7 million committee-by-race observations from the Database on Ideology, Money in Politics, and Elections (Bonica, 2014), focal cycles 1990 through 2020 with outcomes in the following cycle; the discontinuity is a between-committee contrast, and a within-committee event study confirms the redirection directly.

I find that committees whose preferred candidate narrowly loses redirect giving sharply toward the race winner, curtail support for the defeated candidate, leave total giving unchanged, and reorient the partisan composition of their House portfolios toward whichever party wins by roughly seven percentage points — a shift driven primarily by reallocation across the committee’s other House races rather than by the focal-race redirection alone. The pattern matches the access-seeking signature and only that signature. The null on total giving falsifies the budget-expansion hybrid: on the margins expansion would actually move, total giving changes by under five percent, so the redirection is financed within a near-fixed budget. And the winnerward partisan shift — partisan composition signed toward the focal winner’s party — falsifies partisan loyalty directly: committees reorient toward whichever party wins, not toward their own.

A within-design check addresses the default-to-winner reading, on which the committee redirects to the focal winner only because no ideologically similar alternative exists in

its choice set. Replacements are widely available and do receive redirected giving, but the focal-winner channel runs roughly two and a half times that secondary substitution channel — which the design cannot decompose into ideology or pre-existing network ties — and giving to non-focal in-state winners is a precise zero: committees follow the person who took the seat, not the state.

The contributions of the paper are four.

First, I provide the cleanest causal evidence to date that committee giving is access-driven, identified at a margin upstream of the committee-assignment thresholds studied in earlier power-following work (Powell and Grimmer, 2016; Fourinaies and Hall, 2018). The closest comparison, Fowler et al. (2020), runs the same design on the same style of panel but measures firm value rather than subsequent giving; their tightly bounded null and my behavioral redirection are jointly consistent with committees purchasing legislative attention, not cash-back policy favors. And on the subsample of committees sponsored by publicly traded firms, where lobbying can be matched to legislators, the redirection is accompanied by a parallel, winner-specific jump in directed lobbying — the activity through which access is exercised — putting the access reading there on a second behavioral channel, not the giving pattern alone.

Second, the four-outcome joint signature is novel in the close-election literature on contributions, whose donor-side studies pool donor types (Dumas and Shohfi, 2018), study non-partisan municipal contexts (Rueda and Ruiz, 2021), or measure firm value rather than behavior (Akey, 2015; Fowler et al., 2020); with a single outcome, any of the three views can rationalize an estimated discontinuity. The joint design also converts two would-be non-results into evidence: the informative null on total giving establishes within-budget substitution, and reporting the pre-registered partisan-composition outcome under the winner-aligned sign convention turns a within-pool cancellation artifact into a direct pooled falsification of partisan loyalty.

Third, I adjudicate access against ideology with two heterogeneity tests that trade

identification against substantive strength. The race-side moderator — the winner-loser CFscore gap, fixed before the close outcome materializes — yields redirection essentially flat across the race’s own polarization: the cleanest test the Marshall (2022) critique allows, though of the weaker ideology prediction. The PAC-side moderator tests the stronger prediction, the committee’s own ideological cost of funding a distant winner, and shows a tercile-match pattern consistent with a layered ideology channel — but it conditions on a winner trait, and its magnitude is not point-identified under Marshall’s partial-identification bounds. I read the two jointly; the integrated reading is access-dominant, resting on neither test alone.

Fourth, I show that access-seeking is a property of strategic donors rather than of corporate status, relocating the access-versus-expression debate from donor type to donor sophistication. Across Federal Election Commission committee types the redirection is essentially uniform — every organized type falls within a +0.25 to +0.38 band — and non-connected ideological committees, the type an expressive account would predict to be least access-driven, redirect at least as strongly as business committees. Extending the design to individual donors recovers the contrast the committee split cannot: individuals follow the winner less than committees and realign their partisan giving far less, with both responses scaling monotonically in donation size. The purest access signature, partisan reorientation toward the winning party, runs from approximately zero among small donors giving under \$200 to committee levels among the largest givers.

A final pattern concerns change over time: the redirection is largest before the Bipartisan Campaign Reform Act (+0.44), attenuates after it, and remains substantial through the most recent cycles (+0.31 to +0.33). The attenuation parallels the roughly 50 percent decline in donor punishment of extremism that Myers (2025) documents since 2000: donor responsiveness to electoral signals appears to have muted across mechanisms in the post-BCRA institutional environment.

2 Theory

The committees that fund U.S. House campaigns — political action committees alongside party and candidate committees — pool contributions and disburse them to federal candidates. Their behavior has been read in three competing ways, each with sharp implications for what should happen when a previously supported candidate narrowly loses an election.

The access view of campaign contributions, rooted in a campaign-finance literature stretching back to the late 1970s that distinguishes access-oriented donor strategies from electoral-ideological ones (Welch, 1980; Cassie and Thompson, 1998) and developed through three decades of empirical work on PAC giving (Romer and Snyder, 1994; Brunell, 2005; Fourniaies and Hall, 2014, 2018; Barber, 2016; Barber and Eatough, 2020; Phillips, 2024), treats contributions as purchases of a share of a resource-constrained office’s attention: access-oriented donors work with whomever holds the seat and cultivate relationships with both parties (Phillips, 2024; Brunell, 2005). The access value attaches to the office rather than to the legislator personally, so the marginal access return tracks office-holding, not ideological proximity. A closely related anchor on the lobbying side, Hall and Deardorff (2006), models lobbying as a subsidy of legislative effort flowing to *already-aligned* legislators; the access view I test differs in admitting cross-aisle transactions when office changes hands (Brunell, 2005).

The marginal access return to a relationship with a defeated challenger is small — she neither votes nor sits on committees — and the return to the new winner is large for the symmetric reason: when office changes hands, contributions follow whoever can advance the PAC’s agenda. Aldrich et al. (2017) show that leadership-PAC giving responds to majority-party shifts in this way, and Powell and Grimmer (2016) and Fourniaies and Hall (2018) document analogous power-following at the committee margin — without identifying the prior margin my design isolates, the office-versus-no-office threshold itself.

The access view admits two operative sub-mechanisms. Under a push reading, PACs proactively recalibrate their access portfolios when the office-holding identity changes; under a pull reading, the post-race contributions reflect winner-side solicitation — the new office-holder’s staff, party leadership, or fundraising consultants reaching out to the committees that funded the defeated candidate. The two readings carry identical predictions on the four-outcome signature, because contribution data record the transfer, not its initiator; they part ways in how redirection responds to winner-side solicitation capacity. Appendix Section N reports this adjudication, exploiting variation across committee chairs, and I return to it in the Conclusion.

The ideology view, articulated in Ansolabehere et al. (2003) and reflected throughout the cross-sectional CFscore literature (Bonica, 2014, 2018), treats contributions as expressions of policy preferences rather than as investments in legislative production: aggregate giving is too small to be rational policy investment, so contributions function more like consumption — donors give to candidates whose positions match their own to express support, not to earn a return.

A pure ideology account predicts continuity at the win-loss threshold: preferences do not change when an election outcome changes, so a PAC whose preferred candidate barely loses should continue giving to that candidate or to ideologically similar co-partisans, not redirect to the new winner. Meisels (2026) shows that corporate donors disproportionately punish extremist nominees relative to individual donors — prioritization of viable power-holders rather than self-expression; my design extends the contrast to the post-election margin.

A hybrid account combines both motives: ideological commitments keep giving to the loser from collapsing, while access-seeking adds new giving to the winner — which requires expanding the overall budget, because the two legs of the response operate on different recipients. The hybrid is not formalized in the literature, but it is the natural fallback for any reading that allows both motives to coexist (Liu, 2021; Stratmann, 2002,

2005).

These views translate into sign predictions on four co-primary outcomes, all measured at cycle $t+2$, the next House cycle after the focal race: the probability the PAC gives to the focal-race winner, the probability it continues giving to the focal-race loser, the change in the log of its total House giving, and the *winnerward shift* in the partisan composition of its giving — the sign-aligned form of the pre-registered Democratic-share change (Section 3). No two views share the same complete sign pattern; Table T collects the predictions.

The three views can be written as one constrained-allocation problem whose parameters map directly onto the four outcomes, making precise both what the joint signature identifies and what it leaves open. A committee disburses a fixed budget B — set in advance by member dues and organizational constraints — across House candidates c , choosing nonnegative contributions g_c subject to $\sum_c g_c \leq B$. Each candidate yields a return that combines an access component and an ideological component, scaled by a strictly concave function ψ of the amount given ($\psi' > 0$, $\psi'' < 0$):

$$\max_{\{g_c \geq 0\}} \sum_c (V_O o_c + V_I \rho_c) \psi(g_c) \quad \text{subject to} \quad \sum_c g_c \leq B.$$

Here $o_c \in \{0, 1\}$ indicates that candidate c holds office, $\rho_c \in [0, 1]$ is the ideological proximity of c to the committee, and $V_O, V_I \geq 0$ are the committee's valuations of access and of ideological expression. The three views are nested special cases: pure access sets $V_I = 0$, pure ideology sets $V_O = 0$, and the budget-expansion hybrid relaxes the constraint so that B itself rises when the preferred candidate loses.

Let $\omega_c \equiv V_O o_c + V_I \rho_c$ denote candidate c 's allocation weight. The first-order condition equates the marginal return $\omega_c \psi'(g_c)$ across funded candidates to the budget's shadow price: the committee funds c if and only if ω_c clears a threshold, and funds it more generously the larger the weight. The budget thus governs the *scale* of giving while the weights govern its *composition* — an election that leaves B untouched can still reshuffle the entire portfolio through the office terms.

The close-election cutoff flips the office indicators of the focal pair while holding everything else fixed. By the dominant-amount pick rule the preferred candidate p is more proximate than the rival r ($\rho_p > \rho_r$); as-if-random assignment holds B, V_O, V_I , the proximities, and the rest of the choice set balanced across the threshold, switching only (o_p, o_r) . The four outcomes are the resulting discontinuities. Giving to the focal winner and to the focal loser are each governed by the proximity gap scaled by V_I — the office term enters both limits and cancels, so the discontinuity arises as the label “winner” moves from the proximate pick to the less-proximate rival, not from office-following as such — and both are zero when $V_I = 0$, respectively positive and negative when $V_I > 0$. The change in total giving is zero whenever the budget binds, positive only if B expands on loss. The winnerward partisan shift is governed by V_O : reallocating toward whoever holds office tilts the portfolio toward the winner’s party, vanishing when $V_O = 0$ and reversing under partisan loyalty.

[Table T: four-outcome predictions by model — pure access, pure ideology, budget-expansion hybrid, and budget-constrained access-dominant — against the estimated discontinuities.]

The mapping turns each estimate into a statement about a primitive. The precise null on total giving (§4.1) says the budget binds, rejecting the budget-expansion hybrid. The positive winnerward discontinuity (§4.1) requires $V_O > 0$, rejecting pure ideology and partisan loyalty, under which the portfolio’s partisan composition does not track who wins. The same conclusion appears at the candidate level in a within-candidate office premium: the committee funds the rival — the opponent it did not previously favor — about eleven percent of the time when that rival narrowly wins the focal seat and about four percent when the same rival narrowly loses, a gap that cannot arise under pure ideology because the rival’s ideological proximity to the committee is identical either way. Holding office itself carries value. The large winner and loser discontinuities establish the converse: their magnitude requires $V_I > 0$ — not separating access from ideology by

themselves, but ruling out a knife-edge pure-access account in which proximity carries no weight.

Taken together the four outcomes identify $V_O > 0$, $V_I > 0$, and a binding budget, with the relative magnitude V_O/V_I ordered by the substitution test of §4.2: the focal-winner channel is roughly two and a half times the giving redirected toward ideologically similar non-focal candidates, so the office weight dominates whatever pull those similar candidates exert — ideological proximity, a pre-existing relationship, or both — without eliminating it. The object the design identifies is therefore a budget-constrained allocation in which access dominates ideology — $V_O \gg V_I > 0$ with B fixed — rather than the limiting pure-access case $V_I = 0$, which the size of the winner and loser discontinuities and the documented ideology channels rule out. Access-dominant is the precise description; pure access overstates what the evidence supports.

The partisan-composition outcome does the same work against a partisan-loyalty alternative that the headline winner discontinuity cannot rule out. A loyal committee whose pick loses migrates toward a co-partisan winner wherever one exists, tilting its portfolio toward its own party — a negative winnerward shift. A positive pooled winnerward shift is therefore the access fingerprint on the partisan-loyalty margin.

Two implications follow for heterogeneity. First, ideology predicts that the redirection should attenuate in PAC-by-winner ideological distance, because the policy motive for supporting a distant winner is weak. The campaign-contribution access view predicts flatness or amplification: the access value originates in office-holding rather than alignment, and a PAC's information needs are, if anything, greater when the new officeholder's preferences differ from its own (the Hall-Deardorff lobbying subsidy, flowing to already-aligned legislators, predicts otherwise). Because the PAC-side moderator conditions on a winner trait, its magnitude is exposed to the Marshall (2022) compensating-differentials concern and only partially identified, a caveat Section 4.2 develops.

Second, if access is the operative mechanism, the redirection should be flat across the

race-level CFscore gap — PACs chase the office regardless of how ideologically polarized the race is — whereas ideology predicts attenuation in high-gap races. The gap is fixed before the close-election outcome materializes, which places this moderator largely outside Marshall’s objection. That immunity is bought by testing the weaker of the two ideology predictions, so I read the two moderators jointly: the race-side test as the clean-but-weaker anchor, the PAC-side test as the sharper-but-partially-identified companion.

A final implication concerns the boundary of the access channel: contributions are candidate-specific or office-specific, not state-level diffuse transfers. If the redirection were partly a generic state-level access response, giving to non-focal in-state winners would rise on the loss side; the geographic complement I report below tests this prediction directly.

The three views are usually posed as competing accounts of the same donor. They can instead be read as describing different donors, sorted by how strategically they give. Access-seeking is the logic of a donor who treats a contribution as an investment: the return is legislative attention from whoever holds the seat, so the rational investor follows the office. Expression is the logic of a donor who treats a contribution as consumption — an act of support for a candidate or a cause (Ansolabehere et al., 2003). An investor reallocates when the office-holding identity changes; a consumer does not, because the object of consumption is the candidate or the cause, not the office. The two motives therefore predict opposite behavior at the win-loss threshold, and which one dominates should depend on where the donor sits on a spectrum of strategic sophistication. Organized committees — staffed, professionalized, repeat players that maintain standing relationships across many races — sit at the investment end. Small individual donors, who give modest amounts to candidates they identify with, sit at the expressive end. Large individual donors, whose giving approaches the scale and repeat structure of an organized committee, sit in between.

This reframing yields a prediction the within-committee comparison cannot deliver. If

access-seeking is a property of strategic giving rather than of any particular donor type, the four-outcome signature should appear wherever giving is strategic and fade where giving is expressive. Among organized committees, which are uniformly strategic, the signature should be near-uniform across committee types — business, trade, labor, ideological, leadership — even though those types differ sharply in their professed purposes. Across the full donor spectrum, the signature should scale with sophistication: strongest for committees and large individual donors, weakest for small individual donors. The purest access marker — partisan reorientation toward the winning party, which an expressive donor has no reason to perform — should run from near zero at the expressive end to its full magnitude at the strategic end. The within-committee comparison cannot test this, because every committee is already a strategic donor; only an extension across the full sophistication spectrum can.

3 Research Design

My data come from the Database on Ideology, Money in Politics, and Elections (DIME), assembled and maintained by Bonica (2014, 2018), which records every itemized contribution to every federal candidate from 1980 onward, with donor and recipient identifiers, dollar amounts, dates, and recipient ideal-point estimates (CFscores). DIME classifies each contributor as either an individual (`contributor_type = "I"`) or a committee (`contributor_type = "C"`); the committee category comprises every registered committee filer, from corporate, trade, labor, and ideological PACs to party and candidate committees, as distinct from individual donors. My main analysis uses the full committee universe giving to U.S. House candidates in general elections from 1990 through 2022 cycles; a parallel analysis extends the identical design to individual donors (Section 4.3). Coverage of committee giving stabilizes by the early 1990s, before my sample begins, and I end the focal-cycle window at 2020 so that $t+2$ outcomes fit within the 2022 cycle.

The unit of analysis is the committee-by-focal-race observation. For each contested two-candidate U.S. House general election (which I call “the focal race”), I identify every committee that contributed to either of the top two finishers in any cycle up to and including the focal cycle, and assign each a “pick” using a dominant-amount rule: the candidate to whom the committee gave more dollars cumulatively before the focal race. The rule yields a unique pick for the great majority of committee-race pairs; the small subset that gave to both top finishers is handled in the appendix under three alternative rules, which move the four co-primary outcomes by at most 1.2 percentage points (Appendix Section A.2).

The running variable is the focal race’s vote-share margin, signed positive when the PAC’s pick is the eventual winner and negative when the pick is the eventual loser. The cutoff is zero. The signing convention is what makes the design a donor-perspective close-election RDD rather than a candidate-perspective one: PACs whose preferred candidate barely won ($rv > 0$) and PACs whose preferred candidate barely lost ($rv < 0$) sit on opposite sides of the threshold, and the discontinuity at the cutoff identifies the behavioral consequence of the focal-race outcome on the PAC’s subsequent giving.

I estimate the discontinuity using the bias-corrected local-polynomial estimator of Calonico et al. (2014), with MSE-optimal bandwidth selection, a triangular kernel, and standard errors clustered at the focal-race level. De Magalhães et al. (2025) document nominal coverage for this estimator in close-election RDDs with adaptive bandwidth selection. The bandwidth varies by outcome and is reported in Table 1.

My four co-primary outcomes, pre-registered together with the four-model adjudication strategy, are measured at cycle $t+2$, the next House cycle after the focal race: `gave_to_winner_t2`, an indicator for any positive contribution to the focal-race winner; `gave_to_loser_t2`, its analogue for the loser; `delta_log_total_giving`, the change in the log of the PAC’s total House giving (with $\log(1+x)$ handling of zeros); and the change in the partisan composition of the PAC’s House contributions. I pre-registered

the fourth outcome as `delta_dem_share`, the change in the PAC’s Democratic share, and report it directly; because its pooled discontinuity is a cancellation artifact of two opposite-sign within-pool movements, the manuscript’s primary parametrization is the sign-aligned derivative `delta_toward_winner — delta_dem_share` multiplied by +1 in D-winning and -1 in R-winning races — a deterministic transform under which positive values mean movement toward the winner’s party. The family-wise correction applies to the four pre-registered outcomes, and the appendix reports inverse-hyperbolic-sine and unsigned-amount variants of the third. I measure outcomes at t+2 because t+1 is not well-defined for biennial House cycles and t+4 substantially shrinks the sample; the window follows the literature standard (Powell and Grimmer, 2016; Fourniaies and Hall, 2018).

The estimand is a between-PAC contrast, not a within-PAC counterfactual: the discontinuity compares the average behavior of PACs whose pick barely won with that of PACs whose pick barely lost — two populations statistically equivalent on every observable except the focal-race outcome — and is a change in level, not a redirection rate for any individual PAC. The strategic-redirection interpretation rests on the joint structure across the four outcomes rather than on any single discontinuity; any account that fails to match all four signatures is falsified. The within-donor question that the redirection language raises is taken up directly in Section 4.6, which follows the same PAC’s giving to the two fixed focal candidates over time; a within-PAC fixed-effects sensitivity appears in Appendix Section M.

Identification rests on continuity of pre-treatment characteristics through the cutoff (Lee, 2008; Calonico et al., 2014; de la Cuesta and Imai, 2016). Table 2 reports the formal balance test on six baseline covariates measured before the focal race: all six are balanced at the cutoff, with absolute t-statistics below one — the largest, 0.82, on the focal candidate’s CFscore.

[Table 2: Covariate balance at the cutoff.]

The McCrary (2008) density test rejects continuity of the running-variable density at the cutoff ($T = 11.4$), with more mass on the winner side. The rejection is mechanical rather than evidence of manipulation: PACs cannot determine the outcome of close House elections, and because winners attract more PAC dollars than losers, the universe of PACs whose pick won is mechanically larger. As de la Cuesta and Imai (2016) emphasize, the appropriate validity check when the running variable is signed by donor-side choice is covariate balance, which the design passes — including inside the inner-donut window ($|rv| < 0.5$ percentage points), where precise manipulation would concentrate. The full density diagnostic, bandwidth-stability and donut grids, the inner-donut balance table, and a tight null on differential PAC attrition at $t+2$ (validating the zero-treated-as-zero coding) appear in Appendix Section J.

Two further choices warrant flagging. Race-level clustering produces standard errors approximately three times larger than PAC-level clustering, so I report the more conservative race-level inference throughout. Quadratic-polynomial and honest-RD variants are within sampling error of the linear primary (Section 4.7).

Among the heterogeneity tests, the PAC-side ideological-distance moderator conditions on a winner trait and is exposed to the compensating-differentials critique of Marshall (2022); the race-side CFscore-gap and geographic-concentration moderators are pre-determined relative to the close-election outcome, so the critique applies to them in much-attenuated form. Section 4.2 develops the division of labor: the cleanly identified race-side test speaks to the weaker ideology prediction, the partially identified PAC-side test to the stronger one, and I lean on the two jointly.

Two features of the sample depart from prior close-election work on contributions. First, the committee universe is heterogeneous — spanning business, trade, labor, membership, leadership, and non-connected ideological PACs as well as party and candidate committees, with committees positively marked as corporate by Federal Election Commission code a small minority. Rather than restrict to a single donor type, I retain the

full universe and report the redirection by committee type (Section 4.3); pooling types without separating them, as in Dumas and Shohfi (2018), is what confounds the access-versus-ideology contrast. Second, I build a parallel panel of individual donors under identical race definitions and pick rule — 7.9 million donor-by-race observations across 3.0 million donors, focal cycles 1990 through 2018 — so the two designs differ only in the donor universe. The committee sample — 1,710,562 committee-by-race observations, 5,618 contested races, 111,908 unique committees — is to my knowledge the largest behavioral close-election RDD on organized giving to date.

4 Results

4.1 The four-outcome joint signature

Committees sharply redirect their giving toward the race winner when their previously supported candidate narrowly loses. Having one’s chosen candidate lose causes a 35.4-percentage-point drop in the probability of contributing to the focal winner relative to having that candidate win (Table 1, column 1; SE 1.7 pp; 95% CI [0.326, 0.392]; MSE-optimal bandwidth 14.3 pp; effective N approximately 638,000). The magnitude is comparable to the 50–70 percent drop in industry giving after a lost committee assignment (Powell and Grimmer, 2016) and two to three times the committee-join effect of Fourinaies and Hall (2018): losing the seat itself is a stronger power signal than any committee reassignment, and the office-versus-no-office margin lies prior to every committee threshold. The effect rises monotonically in PAC size — small +0.26, medium +0.38, large +0.44 (Appendix Section B.2) — with every tercile strongly positive, so no handful of mega-PACs drives it.

[Figure 1: Two-panel RDD plot, `gave_to_winner_t2` (left) and `gave_to_loser_t2` (right) at the threshold.]

The substitution leg is just as sharp. The probability that the PAC continues giving to

the focal-race loser drops by 21.5 percentage points (Table 1, column 2; SE 2.5 pp; CI [-0.275, -0.178]; bandwidth 7.0 pp): loyalty to the previously supported candidate collapses upon defeat, and the combined picture is substitution — a 35-point gain to the winner and a 22-point loss to the loser. The drop reflects active PAC disengagement rather than mechanical absence from the ballot: at the close-election cutoff, 96 percent of focal losers (PAC-weighted) appear as candidates at $t+2$ and 80 percent return to challenge the same House seat, and restricting the sample to same-seat rematches yields a loser-loyalty drop of -0.220 (SE 0.027; $N = 939,373$), within half a percentage point of the unrestricted -0.215 and well inside one standard error. Mechanical absence concentrates in lopsided races that contribute negligibly to the local-linear estimator (Appendix Section L).

The hybrid account predicts that the redirection is additive — winner support added without loser support subtracted, expanding total giving. The data show no expansion. The change in log total House giving across the cutoff is a precise zero (+0.045 log-units; Table 1, column 3; SE 0.138; $p = 0.76$). Its confidence interval is nonetheless wide — the upper bound corresponds to roughly a 37 percent expansion — but that width is an artifact of the logarithm amplifying committee exit: roughly one in seven committees ceases House giving entirely by the following cycle, and exit is committees *leaving*, the opposite of what a hybrid requires. Measured on the margin the hybrid actually predicts — the surviving committee’s spending envelope — the budget is tightly fixed: a loss changes total giving by -0.3 percent under within-committee fixed effects, by -1.4 percent among committees that remain active, and by -0.1 percent as a share of the prior budget, bounding expansion below roughly five percent on these exit-robust margins while the unconditional log-change remains the wide ± 37 percent (Appendix Section Q). The redirection is financed by reallocation within a near-fixed budget — the budget-constrained allocation of Section 2’s model, in which access competes with loyalty for a fixed envelope.

[Figure 2: RDD plot, $\Delta \log \text{total giving}$ at the threshold.]

The fourth co-primary outcome, the winnerward shift in the PAC’s partisan compo-

sition, is a positive 6.85-percentage-point discontinuity (Table 1, column 4; SE 0.83 pp; $p = 8 \times 10^{-16}$; bandwidth 12.93 pp).¹ The pre-registered raw outcome — the change in the PAC’s Democratic share — is essentially zero in pool (Table 1, column 5; +0.003; $p = 0.54$), a cancellation artifact of two opposite-sign within-pool movements rather than a substantive null on the partisan-loyalty alternative. Two checks discipline the estimate. Because the outcome is a difference, Appendix Section J reports a pre-trend check: the pre-period winnerward drift is +2.9 percentage points, and the difference-in-differences-adjusted shift, netting out each PAC’s own drift, is larger still at +9.49 percentage points (SE 1.21 pp; $p = 3 \times 10^{-14}$). And because the outcome is computed on the PAC’s entire House portfolio, part of the discontinuity could be the focal-race reallocation itself: recomputing it with the focal-race contributions removed yields +5.58 percentage points ($p < 10^{-31}$), 80.5 percent of the full pooled discontinuity — the share computed at the bandwidth at which the focal-excluded outcome is identified (Appendix Table O1) — so the partisan-loyalty falsification rests on portfolio-wide reorientation toward the winner’s party, not on a mechanical artifact (Appendix Section O).

The pooled shift decomposes symmetrically: PACs move toward the winning Democrat when a Democrat wins (+0.072; SE 0.017; $p < 0.001$) and toward the winning Republican when a Republican wins (+0.111; SE 0.016; $p < 10^{-12}$), both confirmed under difference-in-differences adjustment. An account in which PACs migrate toward a co-partisan winner is incompatible with both pools, and the opposite-sign within-pool movements on the raw Democratic-share scale (-6.7 and +6.9 percentage points) are what produce the pooled cancellation.

[Table 1: Main RDD estimates, four co-primary outcomes plus F9 ideology-alternative

¹I report this outcome with the per-observation Democratic-share change multiplied by an indicator that takes value +1 in races won by a Democrat and -1 in races won by a Republican. Positive values therefore indicate movement toward the focal-race winner’s party in every pool. The same convention is used for the robustness grid (Appendix Table B1: logit RDD and race-margin LPM), and the parametrization choice is documented as ADR-006 in the project decision log. The *rdrobust* coefficient on the sign-aligned outcome is reported as the absolute magnitude in the substantive prose (loser-side limit minus winner-side limit) so that positive values represent winnerward redirection by loser-PACs over the winner-PAC baseline.

outcomes.]

After Bonferroni and Benjamini-Hochberg adjustment for the pre-registered four-outcome family, the two large discontinuities survive with adjusted p-values below 10^{-18} and the two nulls remain non-significant by construction (Appendix Section K): three large effects — winner up, loser down, partisan composition winnerward — plus one informative null on total giving.

The four outcomes form a single signature rather than four separate findings, and the model of Section 2 assigns the adjudication load unevenly. A pure-ideology account — including its partisan-loyalty variant — is falsified by the winnerward partisan-composition discontinuity together with the within-candidate office premium: both require office-holding to carry value ($V_O > 0$), which ideology sets to zero, and partisan loyalty in particular predicts a negative or zero winnerward shift where the data show a sharply positive one. The large winner and loser discontinuities do not by themselves falsify ideology — ideology generates the same signs through relabeling (Section 2) — but their magnitude rules out the opposite knife-edge, a pure-access account in which proximity carries no weight. A budget-expansion hybrid is falsified by the null on total giving, near-fixed on precisely the margins expansion would move. And an information account, in which losing reveals the candidate is unlikely to recover, predicts a smaller effect in open-seat races with no incumbency to update on; the open-seat effect is +0.46, larger than the incumbent-reelected +0.42 and the challenger-won +0.35, opposite the prediction (Appendix Section B.2). Access is the only account left standing.

4.2 Access versus ideology

The signature leaves one serious rival standing: ideology. Three tests within the same quasi-experiment adjudicate between them — a direct measurement of where the loser’s dollars would otherwise go, a race-side moderator that is cleanly identified but tests the weaker ideology prediction, and a PAC-side moderator that tests the stronger prediction

but is only partially identified. I take the three in turn and then read them jointly.

The direct measurement targets the default-to-winner reading, under which the PAC redirects to the focal winner only because no ideologically similar alternative exists in its choice set. I track PAC giving to non-focal House candidates whose CFscore lies within 1.38 units of the focal loser's — the 25th percentile of the race-level winner-loser gap, tight enough that the binary outcome does not saturate. The default reading fails on both of its predictions. Replacement candidates are widely available: at baseline, more than three-quarters of PACs give to someone within ideological proximity of the focal loser. And loser-PACs do use them — giving to an ideologically similar non-focal candidate rises 14.0 percentage points at the cutoff (SE 1.7 pp; $p < 10^{-15}$) — yet the focal winner still receives roughly two and a half times that redirection. Looser thresholds, an intensive-margin count, and the full sensitivity curve tell the same story (Appendix Section H); and since noisy CFscores attenuate the estimate (Bonica, 2014), doubling it under a generous reliability ratio narrows but does not close the gap with the +0.354 focal channel. What the secondary channel *is* remains genuinely ambiguous: a PAC gives to a CFscore-adjacent candidate either because it values that candidate's ideology or because it already holds a relationship there. The two are observationally equivalent, so I read F9 as a secondary substitution channel within the PAC's existing relationship space rather than as an identified ideology channel — a distinction that leaves the access interpretation intact either way.

A geographic complement pins the channel to the candidate rather than the state. Tracking giving to non-focal winners in the focal state returns a precise zero (+0.004; SE 2.3 pp; $p = 0.85$): PACs do not redirect toward generic in-state winners but toward the specific person who replaced their preferred candidate, in line with the relationship-specific access view stretching back to Welch (1980) and recently surveyed by Phillips (2024); Hall and Deardorff (2006) supply the closely related lobbying-side framing.

[Figure 3: F9 ideology-alternative redirections compared to the +0.354 focal-winner

reference.]

The race-side moderator conditions on the ideological gap between the focal winner and loser, the absolute CFscore distance fixed before the close outcome materializes. That timing makes it largely immune to the Marshall (2022) concern developed below — immunity bought by testing the weaker of the two ideology predictions: whether redirection flattens across how polarized the *race* is, not whether the PAC’s own ideological cost moderates its giving. Pure ideology predicts attenuation in the gap, since crossing a wider distance is harder; access predicts flatness, because PACs chase the office regardless of how polarized the race is. The data show flatness: redirection across gap quartiles is +0.370, +0.396, +0.363, and +0.320 — a five-point spread on a 35-point base — and the continuous loss-by-gap interaction is statistically zero (-0.011; SE 0.015; $p = 0.48$; full per-outcome detail in Appendix Section G). Even doubling that slope under errors-in-variables in the moderator (Bonica, 2014) leaves it below any plausible ideology-loaded benchmark. The cleanest test the design offers finds no ideology gradient on the prediction it can reach.

The PAC-side moderator conditions on the CFscore distance between the contributing PAC and the focal-race winner, and tests the substantively stronger prediction: whether the PAC’s own ideological cost of funding a distant winner moderates its redirection. This is the margin on which access and ideology compete most directly, but it is also where identification is weakest. Because the moderator is a function of a winner trait, it inherits Marshall’s (2022) compensating-differentials problem — candidates who win close races despite an unfavorable observed trait must be advantaged on unobservables, so conditioning on the trait reintroduces exactly the selection the design removes — and the interaction magnitude is not point-identified. Descriptively, redirection declines monotonically in distance: same-tercile PAC-winner pairs +0.342, adjacent-tercile +0.317, opposite-tercile +0.151, with a continuous interaction of -0.198 (SE 0.014) at face value. But applying Marshall’s §6 partial-identification bounds with the local CFscore variance $\sigma_Z \approx 0.91$ produces intervals that bracket zero at every considered tightness ($K = 0.5$: [-0.65,

+0.26]; Appendix Section F). I therefore retain the tercile pattern as descriptive evidence of a layered ideology channel running alongside the dominant access pattern, without interpreting its magnitude.

[Figure 4: Heterogeneity caterpillar plot combining PAC size, PAC ideology (cfscore terciles), race incumbency type, and era.]

[Figure 5: Race-level cfscore-gap heterogeneity, two-panel (winner redirection by quartile; loser-loyalty drop by quartile).]

The three tests are complementary rather than convergent, and none is individually decisive. The race-side test cleanly rejects the weaker ideology prediction; the PAC-side test finds an ideology gradient real in sign but bounded away from any magnitude large enough to overturn access; and the focal-winner channel runs two and a half times the quantified secondary substitution channel that shadows it. Read together they underwrite the access-dominant reading — $V_O \gg V_I > 0$, not the knife-edge $V_I = 0$ — with the headline +0.354 level unaffected by any of the three. The one non-flat race-side outcome is loser-loyalty, which weakens slightly less in high-gap races (interaction -0.070, $p < 0.001$), a sign-preserving attenuation that does not bear on the winner-side adjudication. An industry split speaks the same way to generality: utilities +0.33, financial services +0.31, telecommunications +0.27, defense +0.35, and non-regulated industries +0.36 (Appendix Section B.2) — no single sector drives the headline, and the access pattern holds across the committee universe rather than concentrating in classically regulated sectors.

4.3 Strategic versus expressive donors

The theory in Section 2 implies a test the committee data can run only halfway and the individual data can complete: access-seeking should be a property of how strategically a donor gives — near-uniform among organized committees, and scaling down toward small expressive donors.

I first split the committee sample by Federal Election Commission committee type

(crosswalk construction and coverage in Appendix Section P). The redirection is essentially uniform across types (Table 3): every organized type redirects at between +0.25 and +0.38, the pooled re-run on the classified sample reproduces the headline exactly (+0.354), and excluding party and candidate committees leaves it essentially unchanged at +0.334. The donor type an expressive account would single out as least transactional — non-connected ideological committees, which exist to advance a policy agenda rather than to cultivate officeholders — redirects among the most, at +0.376, essentially tied with party committees. The type split thus delivers the first half of the theoretical prediction — the signature is near-uniform wherever giving is organized — but cannot deliver the second, because every committee is already a strategic donor.

[Table 3: Redirection by donor type (committee types and individual-donor size bands), with the winner, loser, and partisan-reorientation outcomes.]

To complete the test I re-run the identical design on individual donors, who span the full range of giving sophistication. On the cycles both panels share (focal 1990 through 2018), individual donors redirect toward the winner less than committees do (+0.241 versus +0.358) and realign the partisan composition of their giving far less (a winnerward shift of 0.027 versus 0.073). Within individuals, redirection rises monotonically with donation size: donors giving under \$200 in the focal cycle redirect at +0.148, mid-sized donors at +0.276, and large donors at +0.414, above the committee level. Because the binary outcome is partly mechanical — a donor who gives to more candidates is more likely to give to any one of them — I read the gradient primarily off the partisan-reorientation share, which is invariant to giving breadth and which an expressive donor has no reason to move. On that outcome the gradient is sharpest in meaning: small donors 0.002, statistically indistinguishable from zero; mid-sized donors 0.047; and large donors 0.063, approaching the committee level of 0.073. Small donors follow a specific winner modestly but do not realign their partisan giving at all; large donors realign as much as committees do.

Two checks confirm the gradient reflects motive rather than a breadth artifact: small donors are, if anything, slightly more active at the following cycle than mid-sized donors, so their weak redirection is not attrition; and conditioning on donors who remain active at $t+2$ leaves the gradient intact and if anything sharpens it (small +0.246 versus large +0.626). The realignment marker, which no breadth artifact can generate, runs from approximately zero among small expressive donors to the committee level among the most sophisticated givers.

The gradient relocates the access-versus-expression debate: the operative contrast is not between business and ideological donors, all of which behave transactionally, but between strategic and expressive giving. Access-seeking is the behavior of the sophisticated giver, wherever in the donor population that giver sits.

4.4 A second channel: directed lobbying follows the winner

The evidence to this point is a pattern in *giving*. If it reflects access-seeking, the same redirection should appear in a second, behaviorally independent channel — lobbying, the activity through which the office term V_O of Section 2 is actually exercised. Both channels are observable together on the subsample of committees sponsored by publicly traded firms, whose lobbying can be matched to specific legislative targets through the firm-by-legislator panel assembled by Kim et al. (2026) (focal cycles 2008 through 2016); I re-run the close-election design with the firm’s directed lobbying of the race winner as the outcome.

The discontinuity reproduces. A firm whose backed candidate barely won directs significantly more lobbying expenditure toward that winner over the following cycle than a firm whose candidate barely lost — a local-linear discontinuity of +1.25 log points in directed lobbying dollars (robust SE 0.45, $p = 0.003$), stable across bandwidths from 5 to 20 percentage points. The same firms’ donations, measured in the same data, jump in parallel (a binary +0.17 and a log-dollar +1.41, both $p < 10^{-8}$), so the lobbying result behaves

like the giving headline rather than reflecting an idiosyncratic subsample. Lobbying directed at the *loser* — who is out of office at $t+2$ — shows no discontinuity ($-0.30, p = 0.37$): the effect is specific to the candidate who actually holds the seat, as an access account predicts and a partisan-loyalty account does not.

The result survives the probes a skeptic would run first (Appendix Section R): it holds when the outcome window is restricted to the post-election year alone; it is marginally larger among firms that lobby at all ($+1.42, p = 0.003$), so it is a discontinuity in *how much* a firm lobbies the winner, not in *whether* it lobbies; directed lobbying of the committees on which the winner sits jumps as well ($+0.15, p = 0.001$); firm-side predetermined covariates are balanced at the cutoff; and the race-level margin density is continuous ($p = 0.59$) once the mechanical firm-count replication is removed.

This is mechanism confirmation on the subsample where both channels are observable, not a claim about the full 1990–2020 universe: the firms here are public-company committee sponsors, the window is 2008 through 2016, and the lobbying magnitude is not the giving magnitude. Within those bounds the redirection is no longer only a behavioral signature consistent with access; it is accompanied by a parallel, winner-specific discontinuity in the activity through which access is actually exercised, which is what distinguishes access-seeking from the alternatives the giving pattern alone could not fully separate.

4.5 Era heterogeneity

The access redirection has weakened over time but remains substantial. Splitting the sample at policy-relevant cutpoints yields subsample estimates of $+0.438$ (1990–2002, pre-BCRA), $+0.310$ (2004–2014), and $+0.333$ (2016–2022; eras labeled by outcome-cycle window, with focal cycles running through 2020; Appendix Table B1). The formal era-by-loss interaction is well outside chance (Wald $F = 12.96$; $p < 10^{-5}$). The pre-BCRA effect is roughly 40 percent larger than the 2004–2014 effect, the post-2014 estimate recovers only

partially, and every era's effect remains large and significantly positive: the institutional regime change reduces but does not eliminate access-seeking redirection. The decline is not a compositional artifact of DIME's committee coverage drift: the 1,273 committees continuously observed across all three eras show decline magnitudes essentially identical to the full sample's, though at a uniformly higher level — a separate finding about long-lived committees, not a threat to the era pattern (Appendix Section E, Table E1).

Two readings are consistent with the aggregate era pattern. Post-2002 institutional changes — the Bipartisan Campaign Reform Act soft-money ban and the post-Citizens-United growth of 501(c)(4) and independent-expenditure channels — may have shifted marginal access purchasing into vehicles not captured by direct PAC giving; alternatively, growing partisan polarization changed the propensity to redirect. The within-vs-cross-party era split (Appendix Section E) decomposes the attenuation: within-party redirection — committees whose ideological lean matches the focal winner's party — drops sharply from +0.400 in 1990-2002 to +0.168 in 2004-2014, with partial recovery to +0.209, while cross-party redirection is essentially stable at +0.428, +0.420, and +0.355. The attenuation is thus overwhelmingly a within-party phenomenon, which disposes of one polarization story while leaving another alive. It rules out the version in which rising polarization makes cross-aisle support less palatable — that version predicts the cross-party component to shrink most, and it is the stable one. But a version in which polarization tightens within-party loyalty, with committees increasingly defaulting to their own co-partisans regardless of the focal race, fits the data and is observationally equivalent here to channel substitution — already-aligned committees migrating marginal access flows into 501(c)(4) and super-PAC vehicles while cross-party committees continue through the direct channel. My design does not separate these two accounts.

The attenuation parallels Myers (2025), who reports a roughly 50 percent decline in donor punishment of extremism since 2000: donor responsiveness to electoral signals appears to have muted across mechanisms in the post-2002 environment.

The cross-party component that anchors the era pattern is also the half that most sharply identifies the mechanism. Pooled across the three eras, committees whose ideological lean is opposite the focal winner's party redirect at +0.410, against +0.257 for the within-party component (difference $z = 3.7$). In the allocation model of Section 2, ideological proximity pulls a committee away from a cross-party winner, so a positive cross-party discontinuity isolates the office term V_O : the behavior of a contributor that values incumbency enough to fund a legislator it is predisposed to oppose. Decomposing that component shows the office-seeking is broad-based. Committees in classically regulated industries — utilities, finance, telecommunications, and defense — redirect at +0.360, statistically indistinguishable from the +0.410 of the rest of the committee universe ($z = -1.1$): cross-party office-following is a general property of organized giving, not a behavior concentrated among the committees whose dependence on government is most conspicuous. Nor does geography concentrate the redirection: the cross-party concentration gradient is modest and does not survive a fully flexible specification, so I read it as suggestive only (Appendix Section E.2), and in the full sample the per-tercile discontinuities are statistically indistinguishable across concentration levels (+0.35, +0.38, +0.35), with the continuous interaction — robust to the same flexible specification — implying a discontinuity a few percentage points *smaller* at high concentration: the narrower portfolios of locally concentrated committees, not a district-specific amplification of access (Appendix Section I).

The sharpest evidence sits in ideological space. Were the cross-party effect an artifact of committees whose ideological sign is barely defined, it would vanish among contributors unambiguously opposed to the winner. It does not: sorting cross-party dyads by the ideological distance between committee and winner, redirection falls from +0.359 for the closest third to +0.196 for the most distant — but that most distant third, funding a winner roughly 1.8 CFscore units away, still redirects at +0.196 ($p < 10^{-6}$). Office-seeking overcomes even a large ideological cost, while the attenuation confirms the cost is real

rather than absent. This is the access-dominant region the model of Section 2 identifies — the office weight V_O dominating but not extinguishing the proximity weight V_I — rather than the knife-edge of pure access, and a robust positive effect among the most opposed dyads is one ideology alone cannot generate.

4.6 Within-PAC dynamics

The discontinuity estimates compare different committees on either side of the threshold in a single post-period. To ask whether the redirection is a behavioral response *inside* a given committee, I track each committee’s contributions to the two fixed focal candidates — the eventual winner and the eventual loser — across the two cycles before and after the focal race, restricting to close contests; Appendix Section M.1 reports the specification and full estimates, and Figure 6 plots the trajectories. Among committees whose preferred candidate narrowly lost, giving to the winner — the opponent who defeated the committee’s pick — is essentially absent before the focal race (about two percent, flat across both pre-cycles) and rises to roughly eleven percent two years after the loss, a within-committee increase of eight to nine percentage points ($p < 10^{-50}$) with no detectable pre-trend. Giving to the former pick peaks in the focal cycle and then falls to and below its earlier baseline, so the collapse of loser-loyalty is in part the mechanical end of the campaign cycle, and it is the winner-following rise that carries the clean behavioral signal. The within-committee preference between the two candidates flips accordingly — a swing of about ten percentage points against the pre-campaign baseline ($p < 10^{-15}$): the redirection is visible as a within-donor reorientation, not only as a cross-sectional discontinuity.

[Figure 6: Within-PAC event study. For committees whose pick lost (left panel) and whose pick won (right panel), the probability of contributing to the fixed focal winner and to the fixed focal loser across relative election cycles $\tau = -2$ to $+2$, where $\tau = 0$ is the focal cycle and $\tau = +1$ is the paper’s $t+2$ outcome cycle; the pre-campaign cycle $\tau = -2$

serves as the event-study baseline.]

The trajectories also discipline the interpretation of the headline magnitude. Two years out, committees whose pick won give to that winner at about forty-seven percent, committees whose pick lost at about eleven percent; the gap reproduces the headline discontinuity of §4.1. But the two levels are not the same behavior: the win-side level is largely the committee continuing to fund the office-holder it helped elect, whereas the loss-side level is the genuine act of following a newly powerful office-holder it did not previously support. The decomposition sharpens rather than overturns the central claim: the headline is a level gap dominated by win-side incumbency continuation, and the within-donor switch toward a former opponent — real, precisely estimated, and the smaller part — is the cleaner signature of office-following.

4.7 Robustness

The redirection is stable under the estimator perturbations a technical referee runs first: bandwidths from one to twenty percentage points, quadratic as well as linear local polynomials, donut-hole exclusions from 0.25 to 2 percentage points, PAC-level rather than race-level clustering, and covariate adjustment (+0.353 against +0.354 unadjusted). The smoothness-based honest-RD inference of Armstrong and Kolesár (2018) returns +0.348 at a smoothness bound of $M = 0.1$ and +0.355 at the doubled bound. Appendix Table B1 reports the full grid and Appendix Section B the covariate-adjusted and parametric variants.

More demanding than tuning one estimator are the choices that define the design itself, and the discontinuity survives each of them (Appendix Table B3). It is not an artifact of the extensive-margin indicator: dollars to the winner jump at the threshold, and the winner’s share of the committee’s post-election House giving — a compositional measure that cannot rise merely because a committee funds more candidates — increases by three percentage points. It reproduces exactly when the panel is collapsed to one observation

per race, ruling out any artifact of races with many committees. It is essentially invariant to the pick rule, because earliest-contribution and plurality-of-checks alternatives can disagree only for the roughly one percent of committees that split their giving between both focal candidates. And it is positive and significant in every subsample, smaller though still economically large among committees active in only a single close race — an attenuation that tracks the strategic-donor gradient of Section 4.3. The one specification whose confidence interval does not overlap the primary is the pre-BCRA era subsample — the substantive heterogeneity of Section 4.5, not a robustness concern.

A final sensitivity absorbs all time-invariant committee heterogeneity: a parametric local-linear specification with PAC and cycle fixed effects — identification restricted to committees observed under both win and loss states within the bandwidth — yields a loss coefficient of -0.368 (SE 0.012) on the binary headline and +0.067 (SE 0.006) on the sign-aligned partisan-composition derivative, each within one standard error of the corresponding between-PAC estimate (Appendix Section M). Removing between-PAC composition from the contrast therefore does not attenuate the discontinuity. This fixed-effects contrast is a sub-population sensitivity — identified off the roughly 30 percent of committees observed under both outcomes — rather than a full within-donor counterfactual; the direct within-donor evidence is the event study of Section 4.6, which shows a within-donor switch toward the winner of about nine percentage points against a flat pre-trend.

5 Conclusion

Three decades of debate over what organized political money buys have been framed, implicitly, as a question about donor *type*: whether corporations give transactionally while ideological groups give expressively. The evidence here relocates the question. Access-seeking is not a property of corporate status but of strategic giving. When a committee's previously supported candidate narrowly loses, committees of every stripe follow the

winner — business, trade, labor, leadership, party, and, at least as strongly as any, the non-connected ideological committees an expressive account would single out as least transactional. The sharpest evidence sits at the expressive end of the spectrum: small individual donors follow a winner only weakly and do not realign their partisan giving at all, while large individual donors and organized committees follow and realign in full. The access-versus-expression contrast is real, but it runs along donor sophistication rather than donor type — which is why studies confined to a single donor type have reached conflicting readings of the same behavior.

The behavior that sophistication predicts is within-budget office-following. A committee whose candidate loses does not spend more; it reroutes a near-fixed budget off the defeated candidate and onto the winner, and it tilts the partisan composition of its whole portfolio toward whichever party now holds the seat. On the subsample of committees sponsored by publicly traded firms, where lobbying can be matched to legislators, the redirection is accompanied by a parallel, winner-specific jump in directed lobbying — the activity through which access is actually exercised — moving the access reading there from a labeled behavioral pattern to one corroborated by a second channel. That corroboration is scoped to where both channels are observable; on the full 1990–2020 universe the access reading rests on the four-outcome joint signature. It is consistent with Fowler, Garro, and Spenkuch’s (2020) tightly bounded null on firm value: giving behavior reacts to the office flip while firm value does not — the signature of a subsidy paid for access to a process rather than for a cash-back outcome.

The reading carries normative weight the access-versus-ideology debate has only partly registered. If organized money follows office-holding rather than party, reform arguments premised on money being ideologically aligned with one side miss the system’s central tendency: committees back the side that holds power, symmetrically in Democratic- and Republican-won races alike. The mechanism is access-dominant rather than pure access — a bounded ideology channel runs alongside it, surviving even large

committee-winner ideological distance. The cleanest available test, flatness across the race's own polarization, and the substantively stronger but only partially identified PAC-side test jointly support that access-dominant reading, neither carrying it alone.

Two open questions bound the claim and set the agenda. The first is where the money went: the direct-PAC access channel has attenuated by roughly a quarter since the Bipartisan Campaign Reform Act. Two readings survive my design, which cannot separate them — access flows migrating into the 501(c)(4) and Super-PAC vehicles that grew after Citizens United, or polarization tightening within-party loyalty. A parallel close-election design on independent-expenditure flows, where the channel may have relocated, is the natural next step. The second question is who initiates the redirection. My design records the transfer but not its opener, and two indirect tests point away from winner-side solicitation as the driver: open-seat winners, who arrive with no established donor-relations apparatus, generate the largest redirection of any incumbency type, and interacting the loss with the winner's standing-committee chairmanship — chairs operate the most developed fundraising pipelines in Congress — returns a null whose confidence interval excludes any amplification larger than about half the headline. Whether the relationship is opened by the committee or by the winner's office, and what exactly the access purchased consists of, are distinctions the contribution data cannot draw.

The right description of strategic political giving in the United States is that it follows the winner. The seventy-five committees that moved from Charlie Wilson to Bill Johnson were not converted to Johnson's politics; they were staying with Ohio's Sixth District seat, and their move is the systematic behavior of strategic money wherever office-holding changes hands. Contributions follow whichever office can advance the donor's agenda — wherever the donor gives as an investment in access rather than as an act of expression. The close-election design that isolates the office-versus-no-office margin is generic enough to anchor the next generation of work on access purchase in the post-Citizens United environment.

References

- Akey, P. (2015). Valuing changes in political networks: Evidence from campaign contributions to close congressional elections. *Review of Financial Studies*, 28(11):3188–3223.
- Aldrich, J. H., Ballard, A. O., Rohde, D. W., and Lerner, J. (2017). Does the gift keep on giving? house leadership PAC donations before and after majority status. *Journal of Politics*, 79(4):1133–1145.
- Ansolabehere, S., de Figueiredo, J. M., and Snyder, J. M. (2003). Why is there so little money in U.S. politics? *Journal of Economic Perspectives*, 17(1):105–130.
- Armstrong, T. B. and Kolesár, M. (2018). Optimal inference in a class of regression models. *Econometrica*, 86(2):655–683.
- Barber, M. (2016). Donation motivations: Testing theories of access and ideology. *Political Research Quarterly*, 69(1):148–159.
- Barber, M. and Eatough, M. (2020). Industry politicization and interest group campaign contribution strategies. *Journal of Politics*, 82(3):1008–1025.
- Bonica, A. (2014). Mapping the ideological marketplace. *American Journal of Political Science*, 58(2):367–386.
- Bonica, A. (2018). Inferring roll-call scores from campaign contributions using supervised machine learning. *American Journal of Political Science*, 62(4):830–848.
- Brunell, T. L. (2005). The relationship between political parties and interest groups: Explaining patterns of PAC contributions to candidates for congress. *Political Research Quarterly*, 58(4):681–688.
- Calonico, S., Cattaneo, M. D., and Titiunik, R. (2014). Robust nonparametric confidence intervals for regression-discontinuity designs. *Econometrica*, 82(6):2295–2326.

- Cassie, W. E. and Thompson, J. A. (1998). Patterns of PAC contributions to state legislative candidates. In Thompson, J. A. and Moncrief, G. F., editors, *Campaign Finance in State Legislative Elections*, pages 158–184. CQ Press, Washington, DC.
- de la Cuesta, B. and Imai, K. (2016). Misunderstandings about the regression discontinuity design in the study of close elections. *Annual Review of Political Science*, 19:375–396.
- De Magalhães, L., Hangartner, D., Hirvonen, S., Meriläinen, J., Ruiz, N. A., and Tukiainen, J. (2025). How should we estimate close-election regression discontinuity designs? *Political Analysis*. Forthcoming.
- Dumas, N. K. and Shohfi, K. D. (2018). The behavioral consequences of election outcomes: Evidence from campaign contributions. *British Journal of Political Science*. Advance online publication.
- Fourinaies, A. and Hall, A. B. (2014). The financial incumbency advantage: Causes and consequences. *Journal of Politics*, 76(3):711–724.
- Fourinaies, A. and Hall, A. B. (2018). How do interest groups seek access to committees? *American Journal of Political Science*, 62(1):132–147.
- Fowler, A., Garro, H., and Spenkuch, J. L. (2020). Quid pro quo? corporate returns to campaign contributions. *Journal of Politics*, 82(3):844–858.
- Hall, R. L. and Deardorff, A. V. (2006). Lobbying as legislative subsidy. *American Political Science Review*, 100(1):69–84.
- Kim, I. S., Stuckatz, J., and Freihey, L. W. (2026). Systemic and sequential links between campaign donations and lobbying. *The Journal of Politics*, 88(1):110–124.
- Lee, D. S. (2008). Randomized experiments from non-random selection in U.S. house elections. *Journal of Econometrics*, 142(2):675–697.

- Liu, H. (2021). Campaign contributions and access to congressional offices: Patterns in foreign lobbying data. *Political Research Quarterly*.
- Marshall, J. (2022). Can close election regression discontinuity designs identify effects of winning politician characteristics? *American Journal of Political Science*, 66(2):494–510.
- McCrary, J. (2008). Manipulation of the running variable in the regression discontinuity design: A density test. *Journal of Econometrics*, 142(2):698–714.
- Meisels, M. (2026). Everything in moderation? The effect of extremist nominations on individual and corporate PAC fundraising. *Political Science Research and Methods*, 14(2):353–369.
- Myers, A. C. W. (2025). Do donors punish extremist primary nominees? Evidence from Congress and American state legislatures. *American Political Science Review*, pages 1–21. First View.
- Phillips, C. H. (2024). Interest group strategy and state legislative polarization. Working paper, Vanderbilt University.
- Powell, E. N. and Grimmer, J. (2016). Money in exile: Campaign contributions and committee access. *Journal of Politics*, 78(4):974–988.
- Romer, T. and Snyder, Jr., J. M. (1994). An empirical investigation of the dynamics of PAC contributions. *American Journal of Political Science*, 38(3):745–769.
- Rueda, M. R. and Ruiz, N. A. (2021). When do campaign contributions persist? the role of motivations in non-partisan contexts. Working paper.
- Stratmann, T. (2002). Can special interests buy congressional votes? evidence from financial services legislation. *Journal of Law and Economics*, 45(2):345–373.
- Stratmann, T. (2005). Some talk: Money in politics. a (partial) review of the literature. *Public Choice*, 124:135–156.

Welch, W. P. (1980). The allocation of political monies: Economic interest groups. *Public Choice*, 35:97–120.

This appendix provides supporting analyses referenced in the body of the manuscript. The numbering of tables and figures matches the cross-references in the body.

A Data construction details

A.1 DIME extraction and sample construction

We extract committee contributions to U.S. House general-election candidates from the Database on Ideology, Money in Politics, and Elections (Bonica, 2014, 2018), restricting to `contributor_type = "C"` (committees: the full set of registered committee filers, as distinct from individual donors). The raw DIME file contains every itemized contribution from 1980 onward; we restrict to 1990 forward because committee coverage stabilizes in the early 1990s. The focal cycle is restricted to 1990 through 2020 so that $t+2$ outcomes fit within the 2022 cycle (the most recent in DIME).

For each contested two-candidate U.S. House general election (“the focal race”), we identify every committee that contributed to either of the top two finishers in any cycle up to and including the focal cycle. The construction yields 1,710,562 committee-by-race observations across 5,618 contested races and 111,908 unique committees.

A.2 Pick rule sensitivity

A “dual-pick” PAC×race observation is one where a PAC contributed in cycle t to both top-2 candidates in the same race. Of 19,419 dual-pick PAC×race pairs in the analytic sample (1.1 percent of 1,710,562 PAC×race observations), 3,772 (19.4 percent) involve equal dollar amounts to both candidates (“tie cases”) and 15,647 (80.6 percent) involve unequal amounts. The primary specification applies a dominant-amount rule: keep the row corresponding to the candidate who received more dollars in cycle t , with ties broken lexicographically by `supported`. We assess sensitivity to three alternative rules:

(i) drop-dual (exclude all dual-pick PAC×race observations from the analysis); (ii) replicate (include both candidate-rows for each dual-pick PAC×race, with the running variable sign-flipped per candidate); (iii) share-weighted (include both rows, weighted by each candidate’s dollar share within the dual-pick PAC×race pair, so tie-amount dual-picks contribute weight 0.5 per row and unequal-amount dual-picks contribute weights proportional to dollar share). The drop-dual sensitivity is implemented inside the 27-specification robustness curve (Section B); the replicate and share-weighted sensitivities are run as separate alternative-pick-rule estimations. Because `rdrobust` does not accept observation weights, the share-weighted estimator is implemented as a weighted local-linear regression at the `rdrobust`-selected MSE-optimal bandwidth with race-clustered standard errors; standard errors under the share-weighted rule are cluster-robust but not Calonico–Cattaneo–Titiunik bias-corrected.

Table A2 reports the four pre-registered co-primary outcomes plus the sign-aligned derivative `delta_toward_winner` under each pick rule. Across the three alternatives and the four pre-registered outcomes, the largest absolute deviation from the primary specification is 1.2 percentage points (drop-dual versus primary, on `gave_to_winner_t2`); deviations on `gave_to_loser_t2` and `delta_log_total_giving` are at most 0.7 and 1.0 percentage points respectively, and deviations on `delta_dem_share` and `delta_toward_winner` are below 0.4 percentage points. The body claim that estimates under the three alternatives lie within 1.2 percentage points of the headline (§3) is backed by these numbers.

The headline +0.354 winner-redirection effect, the −0.215 loser-loyalty drop, the precise null on total giving, the cancellation null on raw `delta_dem_share`, and the +6.8 percentage-point sign-aligned `delta_toward_winner` discontinuity are jointly preserved across all four pick rules. The 27-specification robustness curve (Section B) provides additional pick-rule-independent coverage.

Table A1: Pick-rule sensitivity (5 outcomes $\times$ 4 rules).

Reported point estimates with standard errors in parentheses. The primary, drop-dual, and replicate columns use `rdrobust` (Calonico et al., 2014) with a triangular kernel, MSE-optimal bandwidth, and race-clustered HC1 standard errors. The share-weighted column uses a weighted local-linear regression at the `rdrobust`-selected MSE-optimal bandwidth with race-clustered standard errors via `fixest::feols`. The `delta_toward_winner` row reports the absolute value $|\beta|$, matching the manuscript-primary parametrization (§3 footnote; Appendix A.3).

Outcome	Primary (dominant-amount)	Drop-dual	Replicate	Share-weighted
<code>gave_to_winner_t2</code>	+0.354 (0.017)	+0.366 (0.016)	+0.348 (0.017)	+0.356 (0.013)
<code>gave_to_loser_t2</code>	-0.215 (0.025)	-0.217 (0.025)	-0.208 (0.025)	-0.214 (0.021)
<code>delta_log_total_giving</code>	+0.045 (0.138)	+0.055 (0.137)	+0.050 (0.138)	+0.051 (0.087)
<code>delta_dem_share</code>	+0.003 (0.010)	+0.003 (0.010)	+0.002 (0.009)	+0.003 (0.008)
<code>delta_toward_winner</code> ($ \beta $)	+0.068 (0.008)	+0.068 (0.008)	+0.065 (0.008)	+0.068 (0.006)

A.3 Outcome construction

The four co-primary outcomes are constructed at cycle `t+2`: - `gave_to_winner_t2`: indicator for any positive contribution from the PAC to the focal-race winner. - `gave_to_loser_t2`: indicator for any positive contribution from the PAC to the focal-race loser. - `delta_log_total_giving`: $\log(1 + \text{total_giving_t2}) - \log(1 + \text{total_giving_t})$. - `delta_dem_share`: `dem_share_t2` - `dem_share_t` (pre-registered raw partisan-composition outcome, signed in the Democratic direction).

We additionally report the sign-aligned derivative `delta_toward_winner` defined as `delta_dem_share` multiplied by an indicator that takes value +1 in races won by a Democrat and -1 in races won by a Republican. The two parametrizations carry identical statistical content; the sign-aligned derivative aligns the sign of the partisan-composition outcome with the framework’s PAC-perspective treatment indicator (positive value = movement toward the focal-race winner’s party in every pool) and is the manuscript’s primary parametrization for this outcome. The 27-specification robustness curve already uses the same sign-normalization convention for `logit_rdd` and `race_margin_lpm`. Family-wise correction (Bonferroni and BH) is applied to the four pre-registered outcomes; `delta_toward_winner` is reported alongside as a deterministic transform of

`delta_dem_share`, not as a fifth test.

The inverse-hyperbolic-sine transformation of total giving yields a discontinuity of +0.046 (SE 0.135), confirming the log-handling does not drive the null.

B Specification curve

Table B1 reports the full 27-specification grid for the headline `gave_to_winner_t2` outcome.

Specification	Estimate	SE	p-value	h (pp)
primary	0.354	0.017	$< 10^{-100}$	14.3
h = 1	0.352	0.043	1.1e-15	1
h = 2	0.312	0.039	1.7e-18	2
h = 3	0.321	0.035	1.5e-18	3
h = 4	0.334	0.033	5.1e-21	4
h = 5	0.340	0.031	1.6e-24	5
h = 7	0.346	0.028	5.2e-32	7
h = 10	0.353	0.024	1.1e-43	10
h = 15	0.354	0.020	6.8e-67	15
h = 20	0.354	0.018	2.8e-90	20
CER-optimal	0.351	0.019	1.0e-73	9.0
Polynomial p = 2	0.357	0.020	3.6e-73	20.9
Donut 0.25	0.356	0.016	3.0e-117	17.0
Donut 0.50	0.360	0.017	3.1e-98	14.7
Donut 1.00	0.392	0.020	7.3e-92	10.2
Donut 1.50	0.377	0.023	7.5e-61	10.5
Donut 2.00	0.358	0.024	2.3e-50	12.4
Cluster = PAC	0.346	0.006	0.000	7.0
Covariate-adjusted	0.353	0.016	5.8e-115	16.1
Pick: drop dual	0.366	0.016	1.2e-116	14.8
Logit RDD (log-odds)	-1.93	0.074	7.6e-148	14.3
Honest M = 0.1	0.348	0.007	0.000	—
Honest M = 0.5	0.355	0.007	0.000	—
Era 1990-2002	0.438	0.027	3.9e-63	14.2
Era 2004-2014	0.310	0.021	2.0e-49	14.6
Era 2016-2022	0.333	0.032	2.1e-25	19.5
Race-margin LPM	-0.352	0.012 ₄₀	1.0e-195	14.3

(Tables A9–A11 below report the full specification curve for the remaining co-primary outcomes.)

Summary statistics: 100% sign-consistent (after sign-normalizing logit and race-margin LPM), 100% $p < 0.05$, range [0.31, 0.44], coefficient of variation 0.07. Only the pre-BCRA era subsample produces a confidence interval that essentially does not overlap the pooled-primary CI.

B.1 Decompositions by PAC size, PAC ideology, incumbency type, and industry

Per-subgroup `rdrobust` on `gave_to_winner_t2`.

PAC size (terciles of lifetime House giving)

Tercile	Estimate
Small	0.259
Medium	0.380
Large	0.443

Monotone in size. All three terciles are strongly positive, ruling out mega-PAC dominance.

PAC ideology (terciles of real `contributor.cfscore` from DIME, joined ad-hoc per ADR-003)

Tercile	Estimate
D-leaning	0.396
Mixed	0.308
R-leaning	0.346

Strongly positive across the ideological spectrum; the “mixed” group is mildly attenuated but stays well above zero. The `dem_share_t`-based partisan-giving tilt proxy that previously fed this decomposition has been retired in favor of the real CFscore measure

used throughout this paper.

Incumbency type at focal race

Type	Estimate
Open seat	0.459
Incumbent re-elected	0.422
Challenger wins	0.351

Open-seat redirection exceeds incumbent-reelected redirection, opposite the prediction of an information / viability-updating account that requires prior incumbency to update on.

Industry (regulated / non-regulated)

Industry	Estimate
Utilities	0.329
Financial services	0.312
Telecommunications	0.274
Defense	0.349
Non-regulated	0.359

Every regulated industry returns a strongly positive estimate, and the non-regulated baseline is comparable; the access channel is broad-based across the committee universe rather than concentrated in any one regulated sector.

B.2 Substantive design alternatives

Where Table B1 perturbs a single estimator, Table B3 varies the design decisions that define the study — how the outcome is operationalized, which subsample is analyzed, how the panel is constructed, and how each committee’s supported candidate is identified. Each row changes one design choice from the headline specification; all use the MSE-optimal bandwidth, a triangular kernel, and race-clustered standard errors unless the row states otherwise. The pick-rule rows re-stream the raw focal-cycle contributions

to reconstruct the pre-deduplication committee-race table — with both rows present for committees that gave to both focal candidates — and re-resolve the pick under each rule; the race-fixed outcome is reused.

Dimension	Specification	Estimate	SE	p-value	h (pp)
Outcome	Extensive margin: gave to winner (= primary)	0.354	0.017	$< 10^{-100}$	14.3
Outcome	Log dollars given to winner	2.818	0.134	9.0e-100	14.5
Outcome	Winner's share of t+2 House giving (active commit- tees)	0.031	0.002	1.6e-56	21.5
Subsample	Incumbent re-elected	0.422	0.017	2.0e-142	24.3
Subsample	Open seat	0.459	0.027	1.1e-68	15.4
Subsample	Challenger won	0.351	0.036	6.3e-23	10.2
Subsample	Drop post- redistricting cycles (1992, 2002, 2012)	0.357	0.018	3.2e-85	14.6
Subsample	Committees	0.364	0.017	1.6e-99	14.3

Every specification on the binary outcome is positive and significant at $p < 0.001$, spanning the range [0.183, 0.459], and the dollar-denominated and compositional outcomes are likewise strongly positive. The two lower binary estimates — single-race committees (+0.183) and the contribution-weighted panel (+0.239) — remain economically large and highly significant. The single-race attenuation reflects the mechanical narrowing of a less active committee’s giving portfolio: a committee giving to fewer candidates is less likely to give to any specific one, the winner included, which is the same breadth channel the compositional share outcome nets out, and it also tracks the strategic-donor gradient of Section 4.3. The three pick rules agree to within 0.006 because the transaction-count and first-dollar rules reassign only 1,702 and 6,811 of 1,710,562 committee-races respectively — fewer than one in two hundred — since only committees that split their giving between both focal candidates can be assigned differently.

C Bandwidth sensitivity

Estimates across the bandwidth grid h in $\{1, 2, 3, 4, 5, 7, 10, 15, 20\}$ appear in Table B1 above. The bandwidth grid produces a stable estimate range [0.31, 0.35] for the primary outcome, with no monotone trend in either direction. The sensitivity tables for `gave_to_loser_t2`, `delta_log_total_giving`, and `delta_dem_share` show similarly stable patterns (Tables A9, A10, A11 in the body’s referencing convention).

D Donut-hole sensitivity

The donut-hole specification drops observations within $|rv| < d$ for d in $\{0.25, 0.50, 1.00, 1.50, 2.00\}$ percentage points. Estimates range from 0.358 to 0.392, all individually significant at $p < 10^{-50}$. The donut-hole grid addresses the precise-manipulation alternative (in which manipulation is concentrated within a narrow window of the cutoff); the absence

Table A2: Era subsamples on `gave_to_winner_t2`, full sample vs balanced-PAC subset.

Coefficients are `rdrobust` point estimates with race-clustered HC1 standard errors in parentheses below; bandwidth selected per cell by MSE-optimal CCT `mserd`; triangular kernel; local-linear polynomial. The balanced subset restricts to PACs with ≥ 1 PAC $\times$ focal-race observation in each of the three era buckets ($\text{era} \in \{1990\text{--}2002, 2004\text{--}2014, 2016\text{--}2022\}$). Δ rows report the level shift between the balanced subset and the full sample, with the in-SE distance computed against the full-sample standard error.

	1990-2002 (pre-BCRA)	2004-2014	2016-2022 (post-Citizens-United)
Full sample			
$\hat{\beta}_{\text{loss}}$	0.442*** (0.017)	0.308*** (0.005)	0.312*** (0.014)
Bandwidth h (pp)	11.9	8.5	9.4
N obs (within h)	162,038	157,776	128,197
N PACs (within h)	8,907	25,904	21,132
Balanced-PAC subset			
$\hat{\beta}_{\text{loss}}$	0.512*** (0.018)	0.374*** (0.005)	0.328*** (0.025)
Bandwidth h (pp)	13.3	7.7	10.2
N obs (within h)	32,705	29,671	21,170
N PACs (within h)	962	1,111	904
Δ (balanced – full)	+0.071	+0.066	+0.016
(full-sample SE units)	(+4.2 SE)***	(+13.5 SE)***	(+1.2 SE)

of a monotone or large pattern across donut sizes argues against precise manipulation as a threat.

E Era subsamples and Wald interaction

The aggregate era pattern is reported in body §4.5 and shown in Table E1 below: the access redirection has weakened over time but remains substantial in every era, and the formal era $\times$ loss interaction is well outside chance. Table E1 reports the era split for both the full sample and a balanced-PAC subset restricted to the 1,273 committees (1.1 percent of the 111,908 unique committees in the panel) observed in all three eras. The balanced-PAC comparison addresses the adversarial-reviewer concern that the era heterogeneity could partly reflect compositional shifts in DIME’s committee coverage rather than a behavioral change in PAC redirection (Bonica 2014 documents the coverage drift).

Significance: * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$ (two-sided, conventional `rdrobust` infer-

ence for the row-level coefficients; t-statistics for the Δ row are derived against the full-sample SE under independence, conservative because the balanced subset is a strict sub-sample).

Two findings emerge from the side-by-side comparison. First, the *era decline magnitude* is essentially identical in the two samples. From 1990-2002 to 2004-2014, the full-sample redirection drops by 0.134 percentage points of the binary outcome ($0.442 \rightarrow 0.308$); over the same span the balanced-PAC subset drops by 0.138 ($0.512 \rightarrow 0.374$). From 1990-2002 to 2016-2022, the full-sample decline is 0.130 ($0.442 \rightarrow 0.312$) and the balanced-PAC decline is 0.184 ($0.512 \rightarrow 0.328$). The $\text{era} \times \text{loss}$ formal interaction (Wald $F = 12.96$, $p < 10^{-5}$) confirms that the cross-era difference is not chance. The fact that the *gradient* survives unchanged when DIME's coverage drift is held constant — by restricting to PACs continuously observed across all three eras — pins the era heterogeneity on behavioral change rather than on compositional turnover in the coverage panel.

Second, the *level* of redirection is uniformly higher in the balanced-PAC subset across all eras, by roughly seven percentage points in 1990-2002 and 2004-2014 (4.2 and 13.5 full-sample standard errors respectively, both far outside conventional confidence intervals) and by a smaller and statistically marginal 1.6 percentage points in 2016-2022 (1.2 standard errors). Continuously observed PACs are uniformly more access-driven than the cross-section average; the panel's average effect is in part diluted by the entry and exit of more transient PACs whose redirection is somewhat smaller. The level shift is a substantive finding about committee heterogeneity rather than a threat to the era pattern: long-lived committees are *more* access-driven, and the era decline operates within these long-lived committees at essentially the same magnitude observed in the panel.

The audit-trail outcome `delta_toward_winner` (sign-normalized winnerward derivative; `rdrobust` returns right-limit minus left-limit so the $\text{loss}=1$ effect is the sign-flipped value) shows almost no full-vs-balanced separation: era estimates are $-0.058 / -0.086 / -0.048$ in the full sample and $-0.058 / -0.083 / -0.039$ in the balanced subset.

Table A3: Within-vs-cross-party era split on `gave_to_winner_t2`.

rdrobust point estimate (SE), MSE-optimal bandwidth selected per cell, race-clustered HC1 SE, triangular kernel.

Era	Within-party	Cross-party
1990-2002 (pre-BCRA)	+0.400 (0.049)	+0.428 (0.040)
2004-2014	+0.168 (0.043)	+0.420 (0.033)
2016-2022 (post-Citizens-United)	+0.209 (0.073)	+0.355 (0.058)

The partisan-composition channel is essentially flat across both era and sample.

E.1 Within-vs-cross-party decomposition

Two readings of the era attenuation are a priori plausible (body §4.5): post-2002 regulatory channel substitution (BCRA soft-money ban + post-Citizens-United 501(c)(4) growth absorbing access-seeking dollars into vehicles not captured by direct PAC giving) and polarization-driven dampening of cross-aisle redirection. The two readings make opposite predictions about within-vs-cross-party PACs: a polarization story predicts the cross-party component (PAC ideology mismatched with the winner’s party) to shrink most, since polarization makes cross-aisle support less palatable; a channel-substitution story predicts no particular cross-vs-within asymmetry on its own, but is consistent with within-party PACs migrating earlier into alternative vehicles because their access alignment is more substitutable across channels.

We classify each $\text{PAC} \times \text{focal-race}$ observation as `within_party` if the PAC’s CFs-core sign matches the focal-race winner’s party affiliation (D-leaning PACs paired with D winners; R-leaning PACs paired with R winners) and `cross_party` if mismatched. The sample is restricted to non-null PAC CFscore (98.6 percent of the panel; the real `contributor.cfscore` from DIME is the canonical PAC-ideology measure throughout this paper) and to focal races where `winner_party` $\in \{100=\text{Dem}, 200=\text{Rep}\}$. The decomposition is run on the four pre-registered co-primary outcomes plus the sign-aligned `delta_toward_winner` (30 cells total: 5 outcomes $\times$ 3 eras $\times$ 2 components).

The within-party trajectory (+0.400 $\rightarrow$ +0.168 $\rightarrow$ +0.209) drops 58 percent from the pre-BCRA era to 2004-2014, with partial recovery in 2016-2022. The cross-party trajectory (+0.428 $\rightarrow$ +0.420 $\rightarrow$ +0.355) is essentially stable through 2014 (a +0.008 drop, well within the cell standard errors of ± 0.04) and mildly attenuated in 2016-2022. The era attenuation in the pooled `gave_to_winner_t2` estimator is overwhelmingly a within-party phenomenon. The within-party drop from 1990-2002 to 2004-2014 (+0.232) is roughly 28 times larger than the cross-party drop over the same span (+0.008).

This pattern disposes of one polarization story but not another. It rules out the version in which rising polarization makes cross-aisle support less palatable, which predicts the cross-party component to shrink most, whereas the cross-party component is the stable one. But a different polarization mechanism — polarization tightening within-party loyalty, so that committees increasingly default to their own copartisans regardless of the focal race — predicts exactly the observed pattern (a shrinking within-party redirection with a stable cross-party one) and is observationally equivalent here to channel substitution, in which within-party PACs already aligned with the focal-race winner’s party migrated marginal access flows into alternative vehicles after the BCRA soft-money ban while cross-party PACs continued direct-PAC redirection because the cross-aisle access signal is harder to deliver through 501(c)(4) and super-PAC channels (which are typically aligned with the PAC’s home-party network rather than the rival party’s). The within-cross decomposition therefore rules out the cross-aisle-palatability version of polarization but leaves channel substitution and within-party-loyalty tightening as the two live and unseparated readings; cleanly separating them at the mechanism level would require state-level variation in campaign-finance rules within the post-BCRA window, which we leave to future work (body §4.5).

The full decomposition data additionally reports analogous estimates for the remaining co-primary outcomes (`gave_to_loser_t2`, `delta_log_total_giving`, `delta_dem_share`) and the sign-aligned derivative `delta_toward_winner`. Notable

patterns from the secondary outcomes: (i) within-party loser-loyalty drop is consistently larger (more negative) than cross-party loser-loyalty drop in every era (1990-2002: -0.313 vs -0.173 ; 2004-2014: -0.189 vs -0.118 ; 2016-2022: -0.264 vs -0.132), substantively coherent with within-party PACs having a stronger access incentive to follow the winner of “their” party; (ii) `delta_dem_share` shows opposite signs across components by construction (within-party PACs by definition align with the winner’s party direction); (iii) `delta_toward_winner` shows similar magnitude winnerward shifts across components within each era (no within-cross asymmetry on the sign-normalized partisan-composition outcome), consistent with both kinds of PACs moving toward the focal-race winner’s party in the loser-side pool relative to the winner-side pool.

E.2 Cross-party decomposition by sector, geography, and ideological distance

The within-vs-cross-party split of E.1 shows that the cross-party component carries the era-stable, model-diagnostic half of the redirection. This subsection pools that component across the three eras and decomposes it along three predetermined moderators, sharpening what body §4.5 reports. The motivation is theoretical: in the allocation model (body Section 2) ideological proximity pulls a committee away from a cross-party winner, so a positive cross-party discontinuity isolates the office term V_O . The decomposition asks whose office-seeking generates it.

Subgroup-difference z-tests (treating the disjoint subsamples as independent): cross – within $+0.153$ ($z = 3.7$, $p < 0.001$); regulated – non-regulated -0.049 ($z = -1.1$, $p = 0.30$); geographic high – low $+0.080$ ($z = 1.8$, $p = 0.07$); ideological far – near -0.164 ($z = -2.6$, $p = 0.008$). All fourteen subgroup estimates are individually significant at $p < 0.001$, the least significant being the far-distance tercile at $p = 4 \times 10^{-7}$.

The cross-party component is both larger than the within-party component ($+0.410$ versus $+0.257$) and, by the model mapping, the cleaner office-following signal. Its de-

Table A4: Cross-party redirection decomposed by sector, geographic concentration, and committee-winner ideological distance (pooled across eras).

rdrobust point estimate on `gave_to_winner_t2` toward the focal winner; race-clustered HC1 standard error in parentheses; MSE-optimal bandwidth selected per cell; triangular kernel; local-linear. The within-party estimate is retained as the decomposition reference. Industry is the regulated (utilities/finance/telecom/defense) vs non-regulated regex classification of body Appendix Section B.2; geographic concentration is the `pre_state_share` tercile; ideological distance is the committee-winner absolute CFscore distance terciled within the cross-party subsample.

Cut	Subgroup	Estimate (SE)
Reference	Cross-party (all)	+0.410 (0.024)
	Within-party (all)	+0.257 (0.034)
Sector	Regulated (utilities/finance/telecom/defense)	+0.360 (0.041)
	— utilities	+0.367 (0.041)
	— finance	+0.307 (0.056)
	— telecommunications	+0.412 (0.064)
	— defense	+0.431 (0.061)
	Non-regulated	+0.410 (0.023)
Geographic concentration	Low tercile	+0.354 (0.034)
	Mid tercile	+0.411 (0.029)
	High tercile	+0.434 (0.028)
Ideological distance	Near (median 0.72)	+0.359 (0.050)
	Mid (median 1.14)	+0.322 (0.036)
	Far (median 1.76)	+0.196 (0.036)

composition by moderator shows the office-following is broad-based. Committees in classically regulated industries redirect at +0.360, statistically indistinguishable from the +0.410 of the rest of the committee universe ($z = -1.1$), and the four regulated sub-industries bracket the non-regulated estimate rather than exceeding it; cross-party office-following is a general property of organized giving, not a regulated-sector phenomenon. Geographic concentration adds a gradient in the access-predicted direction (high-concentration +0.434 versus low +0.354), but the gradient is modest ($z = 1.8$) and, as shown below, does not survive a fully flexible specification, so it is reported as suggestive. A two-way sector-by-geography cut confirms the gradient is industry-independent: it holds within both regulated and non-regulated committees, and the largest cell is non-regulated committees at high concentration (+0.432), not regulated committees at high concentration (+0.401), so the gradient is not a regulated-sector effect in disguise.

The ideological-distance decomposition is the sharpest of the three and the one that

ties the pattern to the object the model identifies. Sorting cross-party dyads by the absolute CFscore distance between committee and winner, redirection attenuates from +0.359 in the nearest tercile (median distance 0.72) to +0.196 in the farthest (median 1.76; far – near $z = -2.6$). The attenuation confirms that ideological distance carries a real cost, but the farthest tercile — committees funding a winner nearly two CFscore units away, across most of the partisan spectrum — still redirects at +0.196 ($p = 4 \times 10^{-7}$). Office-seeking overcomes a large ideological cost without being extinguished by it, which is the access-dominant region the model identifies (office weight V_O dominant, proximity weight V_I positive) rather than the knife-edge of pure access. The robustly positive redirection among the most opposed dyads also rules out the reading that cross-party redirection is an artifact of committees whose ideological sign is barely defined.

Two methodological points bound the reading. First, the subgroup tercile `rdrobust` estimates — each in the headline sign convention and directly comparable to the +0.410 anchor — are the primary evidence; the fully-interacted continuous specification `loss × moderator × rv` (kernel-weighted, race-clustered) is used only to check whether a tercile gradient survives a flexible functional form. On that check the sector result is confirmed (the regulated interaction is +0.042, $p = 0.05$, marginally reducing redirection), the geography gradient does not survive (loss-by-state-share interaction +0.018, $p = 0.54$), and the distance attenuation is marginal (loss-by-distance +0.071, $p = 0.09$) even though the subgroup terciles show it clearly; the `rv-slope-restricted` interaction used elsewhere in the project is avoided here because it manufactures a spurious geography interaction that vanishes once the running-variable slope is allowed to vary by the moderator. Second, the subgroup-difference z -tests treat disjoint subsamples as independent; shared races across terciles make them mildly anti-conservative, so the cross – within ($z = 3.7$) and far – near ($z = 2.6$) contrasts are read as robust and the geographic high – low ($z = 1.8$) as suggestive. The meaning of “cross-party redirection” is inherited from the within-vs-cross decomposition of E.1; this subsection splits that same quantity by moderator and

does not re-derive its identification.

F PAC-side ideology channel (F14)

The PAC-side ideological-distance moderator characterizes a layered ideology channel that runs alongside the dominant access redirection. We classify PAC-winner pairs by tercile-tercile match in CFscore space (using real `contributor.cfscore` from the DIME committee master file at 100% non-null coverage, 98.6% panel join coverage) and by continuous distance. The moderator conditions on a winner trait and is exposed to the Marshall (2022) compensating-differentials concern, so its magnitude is only partially identified; we read this test as a substantively sharper but partially identified companion to the cleanly identified race-side test of Section G — it speaks to the stronger ideology prediction (whether a PAC’s own ideological cost moderates its redirection), while the race-side test cleanly rejects the weaker one.

Tercile match-distance

Match	Estimate	SE	p-value
Same tercile	0.342	0.022	$< 10^{-50}$
Adjacent tercile	0.317	0.020	$< 10^{-50}$
Opposite tercile	0.151	0.029	1.7e-7

Continuous interaction within MSE-optimal bandwidth (race-cluster SE)

`loss × |pac_cfscore - winner_cfscore| on gave_to_winner_t2: un-`
 bounded point estimate $\beta = \mathbf{-0.198}$ (SE 0.014; 95% CI [-0.226, -0.171]). However, this PAC-side interaction conditions on a winner trait and is exposed to the Marshall (2022) compensating-differentials concern; we report and apply the §6 partial-identification bounds in subsection F.1 below before interpreting the magnitude.

F.1 Marshall (2022) §6 partial-identification bounds

Marshall (2022, §6) shows that politician-characteristic regression discontinuity heterogeneity coefficients are biased by compensating differentials at the cutoff: among close races where the winner is ideologically far from the contributing PAC, unobserved race characteristics (candidate quality, district shocks, opponent weakness) systematically differ from those in close races where the winner is ideologically aligned. The estimator therefore conflates the heterogeneity of interest with these compensating differentials. Under the Cauchy-Schwarz-style assumption that the conditional bias satisfies $|\text{Cov}(W, Z \mid rv \approx 0)| \leq K \cdot \text{Var}(Z \mid rv \approx 0)$ for some scalar K characterizing the strength of the unobserved channel, the true coefficient is bounded by

$$[\hat{\beta} - K \cdot \sigma_Z, \hat{\beta} + K \cdot \sigma_Z]$$

where σ_Z is the local standard deviation of the conditioning trait at the cutoff. We compute σ_Z with kernel weights at the same MSE-optimal bandwidth used to estimate $\hat{\beta}$.

K	σ_Z	$[\beta_{lo}, \beta_{hi}]$	95% CI	
			envelope	Excludes 0?
0.5	0.9118	[-0.6534, +0.2584]	[-0.6814, +0.2865]	No
1.0	0.9118	[-1.1093, +0.7143]	[-1.1373, +0.7424]	No
2.0	0.9118	[-2.0211, +1.6262]	[-2.0492, +1.6542]	No

The bounds bracket zero even at $K = 0.5$ (the most conservative bias allowance). The point estimate $\beta = -0.198$ is therefore not point-identified under Marshall’s partial-identification framework, and we do not interpret its magnitude as a quantified ideology channel. The tercile-match pattern (same / adjacent / opposite) and the 9-cell decomposition above are retained as evidence of a layered ideology channel, real in sign but

bounded in magnitude, and the body’s language is calibrated accordingly: the PAC-side test speaks to the substantively stronger ideology prediction but only partially identifies it, so the access-dominant reading rests on the cleanly identified race-side test (Section G) and the bounded PAC-side test jointly rather than on either alone.

Full 9-cell decomposition (*gave_to_winner_t2*, PAC tercile $\times$ winner tercile)

	WIN_T1 (left)	WIN_T2 (mid)	WIN_T3 (right)
PAC_T1 (left)	+0.355	+0.471	+0.157
PAC_T2 (mid)	+0.317	+0.408	+0.226
PAC_T3 (right)	+0.144	+0.366	+0.266

Diagonal (same-tercile) cells exceed opposite-tercile cells by roughly $2\times$ on both extreme-PAC rows: PAC_T1 with WIN_T1 (+0.355) > PAC_T1 with WIN_T3 (+0.157); PAC_T3 with WIN_T3 (+0.266) > PAC_T3 with WIN_T1 (+0.144).

Why the binary spec disagrees with tercile/continuous. A binary sign-match parametrization (`compatible if sign(pac_cfscore) = sign(winner_cfscore)`) collapses ideology magnitude to direction. Committees cluster heavily at moderate-R (median *pac_cfscore* $\approx$ +0.35; PAC_T2 holds 65% of all panel rows). Binary “compatible” is dominated by weak-R PAC $\times$ weak-R winner pairs where the +0.30 access floor drives the average. Binary “incompatible” is the small set of D-leaning PACs facing R winners (or vice versa) where selection into the off-sign group correlates with PAC types whose access incentive is larger, producing an aggregate gap (+0.402 vs +0.264) that points opposite the tercile and continuous specifications. Tercile and continuous moderators preserve magnitude information and consistently show distance attenuates redirection.

The *dem_share_t*-based proxy version of this binary specification has been retired; the tercile and continuous specifications on real CFscore are the preferred *PAC-side* parametrizations relative to the binary alternative. The cleanly identified adjudication of access against ideology is the *race-side* test in Section G, which tests the weaker of the two ideology predictions but is least exposed to the Marshall critique; this PAC-side test is its

substantively sharper, partially identified companion.

G F8: Race-level CFscore-gap heterogeneity

Per-quartile rdrobust on gave_to_winner_t2:

Gap quartile	Median gap	Estimate	SE	95% CI
Q1 (low gap)	1.01	0.370	0.036	[0.314, 0.453]
Q2	1.59	0.396	0.029	[0.349, 0.464]
Q3	1.89	0.363	0.030	[0.305, 0.424]
Q4 (high gap)	2.20	0.320	0.028	[0.261, 0.369]

Per-quartile rdrobust on gave_to_loser_t2:

Gap quartile	Estimate	SE	95% CI
Q1	-0.259	0.041	[-0.358, -0.197]
Q2	-0.182	0.041	[-0.277, -0.116]
Q3	-0.195	0.043	[-0.294, -0.127]
Q4	-0.152	0.036	[-0.232, -0.092]

Continuous loss $\times$ cfscore_gap interaction (within MSE-optimal bandwidth, race-cluster SE): - gave_to_winner_t2: -0.011 (SE 0.015, $p = 0.48$); CI [-0.041, +0.019]. - gave_to_loser_t2: -0.070 (SE 0.018, $p < 10^{-4}$); CI [-0.105, -0.035]. - delta_log_total_giving: -0.036 (SE 0.074, $p = 0.62$); CI [-0.181, +0.108]. - delta_dem_share (raw, audit trail): -0.001 (SE 0.006, $p = 0.85$); CI [-0.013, +0.011]. - delta_toward_winner (sign-aligned, manuscript-primary): +0.015 (SE 0.006, $p = 0.006$); CI [+0.004, +0.026].

The loser-loyalty interaction is the only outcome with a meaningful interaction; it is sign-preserving (PACs still drop the loser, just slightly less when the race was ideologically polarized) and does not change the access verdict on the winner side.

H F9: Ideology-alternative shadow effect

Outcomes are constructed in cycle $t+2$ as indicators that the PAC made any positive contribution to a non-focal House candidate within ε CFscore units of the focal loser's CFscore, for four threshold values, plus an intensive count and a geographic complement:

- `gave_to_ideo_alt_e10_t2`: $\varepsilon = 1.00$ (a strict sub-set of the tight threshold).
- `gave_to_ideo_alt_tight_t2`: $\varepsilon = 1.38$ (the 25th percentile of the race-level winner-loser CFscore gap; reported as the manuscript-primary tight threshold).
- `gave_to_ideo_alt_e15_t2`: $\varepsilon = 1.50$ (intermediate between tight and median).
- `gave_to_ideo_alt_median_t2`: $\varepsilon = 1.75$ (the race-level median CFscore gap).
- `n_ideo_alt_tight_t2`: count of contributions meeting the tight-threshold criterion (intensive margin).
- `gave_to_focal_state_winner_t2`: indicator for any contribution to a non-focal House winner in the focal state (geographic complement).

						Implied
rdrobust						loss=1
Outcome	point	SE	95% CI	h (pp)	N_h	effect
gave_to_ideo_alt_t2 ($\varepsilon=1.00$)	+0.215	0.014	[-0.242, -0.186]	24.9	1,015,627	+0.215
gave_to_ideo_alt_t2 ($\varepsilon=1.38$)	+0.140	0.017	[-0.168, -0.102]	16.8	732,116	+0.140
gave_to_ideo_alt_t2 ($\varepsilon=1.50$)	+0.106	0.016	[-0.134, -0.070]	16.1	709,217	+0.106
gave_to_ideo_alt_t2 ($\varepsilon=1.75$)	+0.054	0.016	[-0.081, -0.018]	15.2	678,672	+0.054
n_ideo_alt_tight (count)	+24.8	2.70	[-30.5, -19.96]	12.2	552,320	+24.8
gave_to_focal_state_winner_t2 (geo comple- ment)	+0.004	0.023	[-0.040, +0.049]	17.3	757,834	+0.004 (precise zero)

Sign convention: rdrobust returns winner-PAC right-limit minus loser-PAC left-limit.

Negative coefficient $\rightarrow$ loser-PACs give MORE on this outcome $\rightarrow$ ideology-replacement channel is positive in `loss=1` terms.

Sensitivity to the ε threshold: the magnitude of the discontinuity shrinks smoothly and monotonically as ε is loosened (+0.215 at $\varepsilon = 1.00$, +0.140 at $\varepsilon = 1.38$, +0.106 at $\varepsilon = 1.50$, +0.054 at $\varepsilon = 1.75$), and every threshold preserves the negative sign and statistical significance at conventional levels (all four p-values below 0.01, the e10 and tight thresholds below 10^{-15}). The monotonic shape reflects the mechanical fact that loosening the similarity threshold incorporates more candidates whose CFscores are further from the focal loser's, attenuating the discontinuity toward the broader-redirection baseline. The

manuscript-primary $\varepsilon = 1.38$ (25th-percentile cutoff) sits in the middle of a well-behaved sensitivity curve rather than at any extremum; the broader $\varepsilon = 1.75$ reading reported in the same body table reinforces the same channel at the median cutoff. The headline ideology-replacement reading is robust to the choice of similarity threshold.

Saturation diagnostic: the binary ideology-alt outcomes saturate near the cutoff (mean rising from 0.75 at $\varepsilon = 1.00$ to 0.81 at $\varepsilon = 1.75$), confirming that replacement options are widely available across thresholds. The intensive count margin (`n_ideo_alt_tight_t2`) is the cleanest detection of the ideology-shadow channel at the tight threshold and is reported separately.

I F10: Geographic-concentration heterogeneity

Moderator: `pre_state_share` = lifetime PAC giving share to focal state, cycles $\leq t$.

Tercile cuts within focal cycle. Tercile-median values: low 0.012, mid 0.042, high 0.239 (heavily right-skewed).

Per-tercile rdrobust on `gave_to_winner_t2`:

Tercile	Estimate	SE	95% CI	h (pp)
Low	0.348	0.024	[0.300, 0.394]	22.1
Mid	0.378	0.018	[0.349, 0.420]	13.6
High	0.347	0.020	[0.315, 0.393]	12.2

Continuous loss $\times$ `pre_state_share` OLS interactions (within MSE-h, race-cluster SE):

Outcome	Interaction			
	coefficient	SE	p-value	95% CI
gave_to_winner_t2	+0.124	0.009	7.5e-40	[+0.105, +0.142]
gave_to_loser_t2	-0.050	0.011	1.1e-05	[-0.072, -0.028]
delta_log_total_giving	-0.782	0.081	4.1e-22	[-0.940, -0.623]
delta_dem_share	-0.002	0.008	0.78	[-0.018, +0.014]
(raw, audit trail)				
delta_toward_winner	+0.077	0.007	2.2e-29	[+0.064, +0.091]
(sign- aligned, manuscript- primary)				
gave_to_ideo_alt_tight_t2	-0.038	0.011	9.3e-04	[-0.060, -0.015]

Because `loss` is the low-`gave_to_winner` side, the discontinuity at moderator value

m is $\tau(m) = -(\beta_{loss} + \beta_{loss \times share} \cdot m)$, so a *positive* interaction coefficient shrinks the win-minus-loss gap. Correctly decoded, the three binary-outcome interactions say the same thing: higher geographic concentration attenuates every discontinuity — the winner-redirection gap falls from +0.372 at the low-tercile median to +0.345 at the high-tercile median and +0.315 at `pre_state_share` = 0.50, the loser-loyalty drop is shallower, and the ideology-alternative substitution weaker. This matches the essentially flat per-tercile estimates above and tracks the narrower giving portfolios of locally concentrated committees (the moderator main effect on winner-giving is -0.18), the same breadth mechanic as the single-race attenuation in Table B3 and the donor-size gradient in body Section 4.3 — not a district-specific amplification of access. (The `delta_toward_winner` moderated magnitude does rise with concentration, but mechanically: the more concentrated a committee’s giving in the focal state, the larger the share of its portfolio the focal race moves.)

A fully-interacted verification (`loss*share*rv`, allowing the running-variable slope to differ by moderator level — the specification under which the cross-party geography gradient of Section E.2 vanishes) leaves all three binary-outcome interactions intact: `gave_to_winner_t2` +0.118 ($p \approx 5 \times 10^{-10}$) against +0.124 restricted, `gave_to_loser_t2` −0.080 ($p \approx 4 \times 10^{-4}$), `gave_to_ideo_alt_tight_t2` −0.066 ($p \approx 4 \times 10^{-3}$). The interaction is therefore statistically robust; what it supports is the uniform-attenuation reading above, and the body accordingly reports the geography moderator as evidence *against* district-specific concentration of the redirection.

J Validity: McCrary density and bandwidth-stability summary

The McCrary (2008) density test (Table J1):

Test	T-statistic	p-value	density (left)	density (right)
McCrary	11.40	4.4e-30	0.0121	0.0149
(rddensity				
local-				
quadratic)				

Interpretation: rejects continuity, but mechanically. PACs cannot determine House election outcomes; the density asymmetry reflects the mechanical fact that winners attract more PAC dollars across cycles, so any given PAC is more likely to have backed an eventual winner pre-focal. Per de la Cuesta and Imai (2016), McCrary is not the appropriate validity test in donor-perspective close-election RDDs; covariate balance is.

Bandwidth stability for the four co-primary outcomes is reported in the specification curve (Section B). Estimate range for `gave_to_winner_t2` is [0.31, 0.35] across h in $\{1, 2, \dots, 20\}$.

Covariate balance (Table 2 in body): all six baseline covariates $|t| < 1$ at the cutoff;

Table A5: Inner-donut covariate balance ($|rv| < 0.5$ percentage points, 27,783 PAC-race observations across 69 close races).

diff is the loss-side minus winner-side covariate mean; SE is race-clustered HC1.

Covariate	Mean (winner side)	Mean (loser side)	Diff (loser – winner)	SE	t	p	Balanced (t < 2)
log_contribution_to_focal_candidate	7.604	7.622	+0.018	0.070	0.25	0.80	yes
n_transactions_t	2.213	2.166	−0.046	0.118	0.39	0.70	yes
log_total_house_giving_t	11.404	11.247	−0.157	0.196	0.80	0.43	yes
dem_share_t	0.446	0.554	+0.108	0.088	1.23	0.22	yes
pac_age_in_dime	10.941	10.465	−0.476	0.750	0.63	0.53	yes
focal_candidate_cfscore	0.301	−0.039	−0.341	0.254	1.34	0.19	yes

largest is 0.82 on focal-candidate CFscore.

J.1 Inner-donut covariate balance

The de la Cuesta and Imai (2016) defense of the McCrary density rejection rests on the claim that PAC-side donor-perspective close-election RDDs use the running variable’s sign as a donor-side choice rather than a candidate-side strategic behavior, so that the appropriate validity test is covariate balance rather than density continuity at the cutoff. To tighten this defense against the methodologist concern that any precise manipulation would concentrate in the narrowest possible window around the cutoff, we re-run the six-covariate balance test on the inner-donut subsample with $|rv| < 0.5$ percentage points. The subsample contains 27,783 PAC-race observations across 9,775 unique PACs, 69 close races, and 16 cycles. Inside the donut the running variable carries negligible information, so the balance test simplifies from the `rdrobust` local-linear specification of Table 2 in the body to a linear-probability regression of each covariate on the loss indicator with race-clustered HC1 standard errors. The `loss`-coefficient and its t-statistic are the inner-donut balance statistic per covariate.

All six covariates remain balanced at the conventional $|t| < 2$ threshold inside the inner-donut window, with no individual covariate statistically distinguishable from zero (all $p > 0.18$). The maximum $|t|$ of 1.34 is on `focal_candidate_cfscore`, which car-

ries a substantially smaller effective sample (the 69-race subsample is roughly five percent of the full panel's race count) and is consistent with sampling variation rather than precise manipulation. Four of the six covariates additionally meet the stricter $|t| < 1$ threshold that the `rdrobust` test on the full panel meets (Table 2 in the body); the two that exceed $|t| = 1$ (`dem_share_t` and `focal_candidate_cfscore`) carry the largest sampling variability in this narrow window and are ideological in flavor, but neither approaches conventional statistical significance. The inner-donut balance result reinforces the body's reading that the McCrary density asymmetry is mechanical rather than evidence of precise manipulation at the threshold.

J.2 Differential PAC attrition at $t+2$

The four headline outcomes (`gave_to_winner_t2`, `gave_to_loser_t2`, `delta_log_total_giving_t2`, `delta_dem_share / delta_toward_winner`) treat dissolved PACs as zero. If PAC dissolution probability is differential at the cutoff — for example, if PACs whose pick lost are more likely to wind down or merge — all four headline estimates would be biased. We test the concern directly by constructing the binary indicator `active_t2` (equal to 1 if the PAC makes any contribution at cycle $t+2$ anywhere in DIME, and zero otherwise) and re-running `rdrobust` with the same triangular kernel, MSE-optimal bandwidth, $p=1$ local linear, and race-clustered HC1 standard errors used for the headline outcomes. The discontinuity is +0.0075 (SE 0.0043, 95-percent confidence interval [-0.0009, +0.0159], MSE-optimal bandwidth 10.31 percentage points, effective N within bandwidth approximately 472,000). The point estimate is well below the conventional 0.05 materiality threshold, the confidence interval upper bound is 1.6 percentage points, and the raw 5.5-percentage-point loss-side versus winner-side gap in unconditional means (winner-side 0.869, loss-side 0.813) collapses to essentially zero at the cutoff. The zero-treated-as-zero convention for the four headline outcomes is therefore not a source of material bias from differential PAC attrition.

J.3 Pre-trend balance for differenced outcomes

Table 2 in the body reports level-balance tests on six baseline covariates. Two of the four co-primary outcomes are themselves cycle-to-cycle differences (`delta_log_total_giving` and the partisan-composition outcome reported as `delta_dem_share` in raw form and `delta_toward_winner` in the manuscript-primary sign-aligned form), and for these the appropriate companion test is whether the pre-period version of each delta is balanced at the cutoff. A pre-period differential trend would imply that the post-period delta-RDD estimate could in principle reflect a continuation of pre-existing trajectories rather than a clean focal-race effect. We construct pre-period deltas by panel-internal lookup of each PAC’s prior-cycle giving (80.5 percent coverage at t-2; PACs without a focal race at t-2 are excluded from this check) and re-run `rdrobust` on each.

Pre-period

outcome	Estimate	SE	p	Balanced
<code>delta_log_total_giving_pre</code>	-0.007	0.037	0.62	yes
<code>delta_toward_winner_pre</code> (sign-aligned, manuscript-primary)	-0.029	0.007	9×10^{-5}	no
<code>delta_dem_share_pre</code> (raw, audit trail)	-0.017	0.008	0.010	no

The pre-period total-giving delta is balanced cleanly, which strengthens rather than weakens the headline “no expansion” interpretation: there is no pre-existing differential drift in giving levels at the cutoff and no post-focal one either. The pre-period winner-ward delta is not balanced: winner-PACs were drifting toward their candidate’s party in the pre-period by approximately 2.9 percentage points per two-cycle window more than

loser-PACs were toward their candidate's party. The raw Democratic-share-based pre-period delta carries the same imbalance under the D-direction parametrization. The post-period winnerward estimate of +0.0685 is therefore consistent with a deviation from a continuing pre-trend rather than with a pure level effect. A properly identified difference-in-differences-adjusted RDD on the composite outcome `delta_toward_winner_did = delta_toward_winner_post - delta_toward_winner_pre` (constructed by netting out each PAC's pre-period winnerward drift directly, with a sample restriction to PAC-cycle observations that have non-NA `dem_share` at all three cycles $t-2$, t , $t+2$) returns a pooled discontinuity of $|\beta| = +0.0949$ (SE 0.012, $p = 3 \times 10^{-14}$, 95 percent CI [+0.068, +0.116]; bandwidth 12.2 percentage points; effective N approximately 348,000). The pre-trend correction enlarges rather than attenuates the substantive interpretation: netting out the pre-existing drift exposes a sharper winnerward redirection at the focal-race threshold than the raw post-period estimate alone reveals.

J.4 Within-pool DiD-adjusted partisan-composition RDD

The pooled pre-trend imbalance is itself decomposable across winner-party pools. When the within-pool pre-trend balance test is run separately on D-winning and R-winning races, the D-winner pool returns a near-balanced pre-trend on the winnerward outcome (+0.010, SE 0.011, $p = 0.53$) while the R-winner pool returns a sharply imbalanced one (+0.038, SE 0.009, $p < 10^{-5}$; +0.038 on the winnerward outcome corresponds to -0.038 on the raw Democratic-share outcome). The two pre-trend estimates differ at $|z| = 1.91$ ($p = 0.06$) on the winnerward parametrization and at $|z| = 3.36$ ($p = 0.0008$) on the raw Democratic-share parametrization, so the two pools are not symmetric on this margin under either reading. To re-estimate the focal-race RDD with the pre-period drift netted out per pool, we use a composite difference-in-differences outcome that subtracts each PAC's pre-period winnerward drift from its post-period winnerward shift directly, computed as the raw composite `delta_did = (dem_share_{t+2} - dem_share_t)`

- (dem_share_t - dem_share_{t-2}) and then sign-aligned to the winnerward parametrization. The composite outcome requires non-NA PAC observations at t-2, t, and t+2; coverage is 66.2 percent of the panel (1,131,711 of 1,710,562 rows).

Pool	Raw post	β	Pre-trend β	
			(SE)	(SE)
Pooled	+0.0685 (0.008)	+0.029 (0.007)	+0.0949 (0.012)	[+0.068, +0.116]
D-winner	+0.0710 (0.012)	+0.010 (0.011)	+0.0718 (0.017)	[+0.032, +0.100]
R-winner	+0.0742 (0.012)	+0.038 (0.009)	+0.1106 (0.016)	[+0.082, +0.143]

All entries are reported with the substantive sign convention (positive value = winnerward redirection in percentage points; same sign convention used in the body and the 27-spec robustness curve). Three points warrant emphasis. First, both within-pool DiD-adjusted estimates are positive and substantively meaningful, and all three confidence intervals (pooled, D-winner, R-winner) exclude zero ($p < 10^{-13}$, $p < 0.001$, and $p < 10^{-12}$ respectively). The winnerward redirection appears in both pools symmetrically. Second, the D-winner DiD estimate (+0.072) is essentially identical to the raw post-period estimate (+0.071) because the D-winner pre-trend is itself near zero. The pre-trend correction barely alters the D-winner conclusion. Third, the R-winner DiD estimate (+0.111) is roughly 50 percent larger than the raw post-period estimate (+0.074). The R-winner pool was already trending winnerward in the pre-period (+0.038 in winnerward, equivalent to -0.038 in raw Democratic-share, over two cycles), and netting out that pre-existing trend reveals a sharper redirection at the focal-race threshold than the raw post-period estimate alone would suggest. The pre-trend imbalance, in other words, attenuates rather than inflates the within-pool R-winner power-following signal: the raw post-period within-pool magnitudes are conservative, and the case against partisan loyalty is strengthened, not

weakened, by the proper DiD adjustment. The pooled DiD-adjusted winnerward shift of +0.0949 is the load-bearing direct evidence against the partisan-loyalty alternative — visible in the pooled estimate alone, not only in the within-pool decomposition.

K Multiple-testing correction

Bonferroni and Benjamini-Hochberg adjustments for the four-outcome family applied to the four pre-registered outcomes:

Outcome	p			Passes	
	p (raw)	(Bonferroni)	p (BH)	Bonferroni	Passes BH
gave_to_winner	8.4e-101	3.4e-100	3.4e-100	yes	yes
gave_to_loser	24.4e-20	1.8e-19	8.8e-20	yes	yes
delta_log_total_giving	0.76	1.00	0.76	no	no
delta_dem_share	0.54	1.00	0.72	no	no

Two of four pre-registered co-primary outcomes are massively significant after correction. The total-giving outcome is a genuine informative null (the exit-robust margins bound any budget expansion below roughly five percent, falsifying the budget-expansion hybrid; Appendix Section Q). The raw partisan-composition outcome (`delta_dem_share`) is non-significant under family-wise correction because the pooled estimate is a cancellation artifact of two opposite-sign within-pool movements rather than a substantive null. Reported on the manuscript-primary winnerward parametrization (`delta_toward_winner`, the deterministic sign-aligned derivative documented in Section A.3), the partisan-composition outcome is itself a large positive discontinuity (point = 0.0685 on the `rdrobust` convention reported as $|\beta|$, $p \approx 8 \times 10^{-16}$, well below any family-wise threshold). The corrected pattern reads as three large effects (winner up, loser down, partisan composition winnerward) plus one informative null (total giving unchanged), which together falsify the hybrid and partisan-loyalty alterna-

tives. The winnerward parametrization is reported as a deterministic transform of the pre-registered `delta_dem_share`, not as an additional test, so the Bonferroni and BH thresholds are unchanged. The four outcomes are of course not independent tests — a committee that redirects toward the winner mechanically moves its portfolio composition winnerward — which is why the body reads them as one joint signature rather than four separate confirmations; the decomposition in Appendix Section O quantifies the mechanical component of the composition outcome directly (19.5 percent, the remainder being portfolio-wide spillover).

L Loser-recontest sensitivity (active disengagement vs mechanical absence)

The body’s loser-loyalty drop of -0.215 (Section 4.1) potentially conflates two channels: (a) **active disengagement**, where PACs cease supporting a defeated candidate even though that candidate remains on the ballot at $t+2$, and (b) **mechanical absence**, where the focal loser does not run again and `gave_to_loser_t2 = 0` is forced by construction. We assess the relative contributions through three nested estimates and an office-of-recontest decomposition.

L.1 Three estimates

Specification	β	SE	95% CI	h (pp)	N (PAC $\times$ race)
(1) Unrestricted (full panel)	-0.2155	0.0247	[-0.275, -0.178]	6.96	1,710,562
(2) Any-office recontest at t+2	-0.2327	0.0262	[-0.296, -0.193]	7.15	1,062,309
(3) Same House seat rematch at t+2	-0.2204	0.0268	[-0.285, -0.180]	7.44	939,373

Specifications (2) and (3) restrict to PAC $\times$ race observations where the focal loser appears in the DIME candidacy roster (`recipient.type = "cand"`) at cycle t+2. Spec (3) further restricts to losers who return to challenge the same House seat (state and district both match the focal race). All three estimates lie within two percentage points of one another. The same-seat rematch sub-estimate confirms active disengagement at essentially the full unrestricted magnitude: $0.220 / 0.2155 = 1.023$.

L.2 Office decomposition (close-cutoff, PAC-weighted)

Among PAC $\times$ race observations with $|rv| \leq 7$ pp (the MSE-optimal bandwidth):

Office at t+2	Share of close-cutoff obs	Share of recontesters
Same House seat (rematch)	79.6%	91.8%
Different House district	4.8%	5.6%
State office	1.6%	1.8%
Senate	0.7%	0.8%
No recontest (mechanical absence)	13.3%	—

Within the close-election sample, 80 percent of focal losers return to challenge the same House seat held by the focal-race winner; only 13 percent truly retire from electoral politics. By contrast, the population-wide retirement rate (across all rv) is 44 percent, and the contrast (44 percent versus 13 percent) is the design-level reason mechanical absence does not contaminate the cutoff estimator: retirement concentrates in lopsided races (where rv > 10 pp, 37 percent of focal losers retire) that contribute negligibly to the local-linear fit.

L.3 Implication for active-disengagement interpretation

The bracket framing — recontest-restricted estimate as a lower bound on |active disengagement| due to competence-selection bias, unrestricted estimate as an upper bound due to mechanical-absence inflation — is empirically not informative for this sample: at the close-election cutoff, 80 percent of races feature the focal loser back on the same ballot, leaving little room for selection differential to operate, and the same-seat sub-estimate is essentially identical to the unrestricted baseline. We interpret the -0.215 drop as overwhelmingly active PAC disengagement: PACs choosing the new incumbent (the focal-race winner) over their previously preferred candidate who remains on the same ballot.

Two notes on the specifications. First, spec (2) here uses the `dime_recipients` candidacy roster, which is more strict than a `contribDB`-based “any contribution to a `cand`-type recipient” check; the `contribDB`-based version yields -0.217 and reflects 96.5 percent PAC-weighted recontest at the cutoff. The roughly ten-percentage-point gap between the

two recontest measures (96.5 percent versus 86.7 percent in `dime_recipients`) is largely composed of speculative cycle activity and short-lived primary candidacies that appear as `cand`-type recipients in `contribDB` but not on the formal candidacy roster. Both measures yield the same substantive conclusion. Second, the slight increase in magnitude under restriction (-0.233 in spec (2), -0.220 in spec (3), versus -0.215 unrestricted) reflects the elimination of non-recontest races that contribute zero discontinuity by construction: when the focal loser is absent from the ballot, `gave_to_loser_t2` = 0 is forced symmetrically across the cutoff, contributing nothing to the local-linear fit.

M Within-PAC fixed-effects sensitivity

The headline estimates in Section 4 use a regression-discontinuity design that identifies a between-PAC contrast at the cutoff: the difference in average behavior between PACs whose pick barely won and PACs whose pick barely lost. The substantive language of “redirection” frames this contrast as a within-PAC counterfactual — a single PAC behaves one way when its pick wins and another way when its pick loses. The body Research Design § acknowledges this distinction explicitly. This appendix supplies the within-PAC fixed-effects sensitivity that validates the prose framing empirically.

We restrict the sample to observations within the MSE-optimal bandwidth used in the primary RDD specification ($h = 14.29$ percentage points for `gave_to_winner_t2`, $h = 12.93$ percentage points for the sign-aligned `delta_toward_winner`) and estimate a parametric local-linear LPM:

$$y_{i,t} = \alpha_i + \gamma_t + \beta \cdot \text{loss}_{i,t} + \delta \cdot \text{rv}_{i,t} + \theta \cdot (\text{loss}_{i,t} \times \text{rv}_{i,t}) + \varepsilon_{i,t},$$

where α_i is a PAC fixed effect, γ_t is a cycle fixed effect, $\text{loss}_{i,t}$ is the focal-race-outcome indicator (1 if the PAC’s pick lost), and $\text{rv}_{i,t}$ is the running variable (focal-race vote-share margin, signed positive when the PAC’s pick is the eventual winner). Standard errors are

Table A6: Within-PAC fixed-effects sensitivity at MSE-optimal bandwidth.

β_{loss} is the parametric LPM coefficient on the loss indicator with PAC and cycle fixed effects absorbed; the rdrobust loss-effect column reports the headline RDD point estimate from Table 1 re-signed as the equivalent loss=1 effect; distance is reported in units of the rdrobust standard error.

Outcome	h (pp)	N obs	PACs with both loss states	β_{loss} (within- PAC FE)	SE(β_{loss})	rdrobust loss-effect	Distance (SE units)
gave_to_winner	14.22	601,787	18,080 (30.1%)	-0.368	0.012	-0.354	0.81
delta_toward_loser	12.93	486,014	12,903 (54.4%)	+0.067	0.006	+0.069	0.21

clustered at the focal-race level. Estimation uses `fixest::feols` to accommodate the high-dimensional PAC fixed effect (approximately 16,000 to 24,000 absorbed levels per outcome after singleton removal).

For both outcomes the within-PAC coefficient is sign-consistent with the rdrobust headline and lies well within one standard error of it: 0.81 SE for the binary headline outcome, 0.21 SE for the sign-aligned partisan-composition derivative. Identifying within-PAC variation requires the same PAC to be observed under both loss = 0 and loss = 1 within the bandwidth; this holds for 30.1 percent of PACs in the binary-headline sample and 54.4 percent in the partisan-composition sample, providing a sufficient base for precise estimation. The within-PAC t-statistic on the binary headline outcome is 31.7.

Two notes on interpretation. First, the within-PAC coefficient on the binary headline is slightly larger in absolute value than the between-PAC rdrobust estimate (0.368 versus 0.354). Netting out PAC-specific baselines does not attenuate the discontinuity; the within-PAC contrast at the cutoff is at least as sharp as the between-PAC contrast. Second, the parametric form above is a sensitivity check rather than a primary specification: it does not adopt rdrobust's bias correction and is reported only at the MSE-optimal bandwidth selected for the headline outcome. The close match between the between-PAC and within-PAC estimates indicates that the strategic-redirection prose interpretation is consistent with the within-donor counterfactual the language implies.

Table A7: Within-PAC giving trajectories to the fixed focal candidates (headline bandwidth $h = 14.29$ pp).

Cell entries are the share of committee-by-race units contributing to the named target in the given relative cycle; $\tau = +1$ is the paper’s $t+2$ outcome cycle.

Side and target	$\tau = -2$	$\tau = -1$	$\tau = 0$	$\tau = +1$ (t+2)	$\tau = +2$
Loss side: P(give winner)	0.016	0.020	0.034	0.105	0.071
Loss side: P(give loser)	0.105	0.216	1.000	0.097	0.051
Win side: P(give winner)	0.177	0.319	0.997	0.473	0.250
Win side: P(give loser)	0.007	0.011	0.015	0.010	0.011

Note: Restricting to units within five percentage points of the cutoff leaves the loss-side winner-following rise essentially unchanged (0.022, 0.026, 0.033, 0.108, 0.068 across the five cycles). The $\tau = 0$ loss-side give-to-loser value of 1.000 reflects that the focal loser is, by construction, the candidate the committee supported in the focal cycle.

M.1 Within-PAC event-study trajectory

The fixed-effects sensitivity above nets out time-invariant committee heterogeneity but is still a single-period between-outcome contrast within the bandwidth. To recover the within-donor *dynamics* — the evidence the body §4.6 reports — I track each committee-by-focal-race unit’s giving to the two fixed focal candidates, the eventual winner and the eventual loser, across relative election cycles $\tau = -2, -1, 0, +1, +2$, where $\tau = 0$ is the focal cycle and $\tau = +1$ coincides with the paper’s $t+2$ outcome cycle. Outcomes are re-streamed from DIME for calendar cycles 1986–2024 and cached; as a validation, giving at $\tau = +1$ reproduces the canonical `gave_to_winner_t2` and `gave_to_loser_t2` outcomes exactly (agreement 1.0000). The sample is cutoff-local at the headline bandwidth ($h = 14.29$ pp): 235,192 loss-side and 402,338 win-side committee-by-race units, of which 18,080 committees (30.1 percent) contribute both a close win and a close loss. Estimation uses `fixest::feols` with unit fixed effects and the pre-campaign cycle $\tau = -2$ as the reference, since giving to a candidate peaks mechanically during that candidate’s own campaign ($\tau = -1$), which is not a valid pre-treatment baseline.

Table A8: Within-donor event-study estimates, loss side (unit fixed effects, reference $\tau = -2$).

The first two rows are within-unit time profiles for a single target; the third is the winner-versus-loser within-unit difference-in-differences. Coefficients are changes relative to the $\tau = -2$ baseline.

Coefficient (vs. $\tau = -2$)	$\tau = -1$	$\tau = 0$	$\tau = +1$	$\tau = +2$
P(give winner), within-unit	+0.004	+0.018	+0.089	+0.055
P(give loser), within-unit	+0.112	+0.896	-0.007	-0.054
Winner-vs-loser gap (within-unit DiD)	-0.108	-0.877	+0.096	+0.109

Note: The $\tau = +1$ within-unit rise in loss-side giving to the winner (+0.089) has $p = 8 \times 10^{-95}$. The winner-minus-loser difference-in-differences — the change in the winner-versus-loser gap relative to $\tau = -2$ — is -0.108 at $\tau = -1$, +0.096 at $\tau = +1$ ($p = 6 \times 10^{-16}$), and +0.109 at $\tau = +2$ ($p = 3 \times 10^{-25}$); combined with the -0.089 raw gap at the $\tau = -2$ baseline, the within-committee preference flips from favoring the former pick before the race to favoring the winner after it. On the win side the same difference-in-differences is positive throughout (baseline gap +0.170, rising to about +0.46 at $\tau = +1$), reflecting incumbency continuation. A naive loss-versus-win difference-in-differences on giving to the winner is confounded, because that outcome is incumbency continuation on the win side but opponent-following on the loss side: it shows a spurious +0.139 win-side pre-trend at $\tau = -2$ and a $\tau = +1$ estimate of -0.069, far below the cross-sectional level gap. The clean within-donor object is therefore the within-side winner-versus-loser comparison reported here.

N Push-versus-pull adjudication via winner committee chair status

The push-versus-pull distinction within the access view (Theory section) asks whether the post-focal redirection is initiated by the PAC (push) or by the winner’s office staff and party-leadership solicitation infrastructure (pull). Direct adjudication would require contact-level records on which side opened the post-focal relationship, which DIME does not contain. This appendix reports two indirect tests that converge against a pull reading: a decomposition by winner incumbency type (already discussed in Appendix Section B.2 and summarized below for cross-reference) and a direct heterogeneity test using external congressional committee assignment data.

N.1 Incumbency-type decomposition (summary)

For cross-reference. Among the three incumbency types of focal-race winners:

Winner incumbency type	$\hat{\beta}$ on gave_to_winner_t2	Pull prediction
Open seat (no prior staff or donor pipeline at t)	+0.46	LOWEST (infrastructure absent)
Incumbent reelected (established staff, donor pipeline)	+0.42	HIGHEST (most-developed pull capacity)
Challenger won (new staff at t+1, displaces incumbent)	+0.35	LOW (no infrastructure inherited)

The observed ordering — open seat (+0.46) > incumbent reelected (+0.42) > challenger won (+0.35) — inverts the pull-mechanism prediction (which would rank incumbent-reelected highest). Open-seat winners arrive at t+1 with the least developed donor-relations apparatus, yet they generate the strongest redirection. The pattern is more consistent with PAC-side strategic recalibration than with winner-side solicitation effort. Full discussion is in Appendix Section B.2.

N.2 Committee chair heterogeneity (Stewart-Woon House assignments)

The incumbency-type decomposition uses an aggregate proxy for winner-side solicitation capacity. A sharper test uses external data on each winner's committee position at the relevant Congress. Standing-committee chairs operate the most developed donor-relations pipelines in Congress: chairs control hearing schedules and bill markups, command the largest committee staff, and maintain the most active fundraising operations on behalf of the majority. If the post-focal redirection were materially driven by winner-side

Table A9: Subgroup rdrobust on `gave_to_winner_t2`, full sample vs by winner chair status (in-coverage subsample, Congresses 103-115).

rdrobust with triangular kernel, MSE-optimal bandwidth, race-clustered HC1 standard errors. Standard errors in parentheses.

	Winner not chair	Winner is chair
Full-committee chair		
$\hat{\beta}$ (SE)	+0.352*** (0.008)	+0.306*** (0.023)
Bandwidth h (pp)	11.4	5.3
N races (within h)	4,291	231
N obs (within h)	367,401	4,480
Any chair (full + subcmte)		
$\hat{\beta}$ (SE)	+0.349*** (0.008)	+0.523*** (0.014)
Bandwidth h (pp)	11.8	6.4
N races (within h)	4,240	282
N obs (within h)	375,524	6,576

solicitation, the chair-led winners would show systematically stronger redirection than rank-and-file members.

Data source. Stewart and Woon’s House Committee Assignment Data (Congresses 103-115; updated 2017-11-17; http://web.mit.edu/cstewart/www/data/house_assignments_103-115-3.xls) records every member-committee assignment with party-status code (1 = majority original, 2 = minority original) and rank within party status. Chairs are identified as members with party-status = 1 and rank = 1 within a standing committee in a given Congress. We construct two indicators: `is_full_cmte_chair` (chair of any standing committee, codes < 200 in Stewart-Woon’s coding) and `is_any_chair` (full or subcommittee chair). The `bonica.rid` → real ICPSR crosswalk is constructed from the DIME recipients file (DIME’s `ICPSR` field is Bonica-encoded as `ICPSR × 10000 + cycle`; `ICPSR2` is the real four-to-five-digit ICPSR used by Stewart-Woon and Voteview). Coverage spans Congresses 103-115 (focal cycles 1992-2016, `t+2` outcomes 1994-2018), corresponding to 60.8 percent of the full panel; focal cycles 1990, 2018, and 2020 are missing chair coverage.

The two subgroup rows differ in opposite directions: full-committee chairs receive marginally less redirection than non-chairs (−0.046), while any-chair winners (including subcommittee chairs) receive markedly more (+0.174). The two estimates use differ-

Table A10: Continuous loss $\times$ chair interaction (parametric LPM at MSE-optimal bandwidth, in-coverage subsample).

`gave_to_winner_t2` regressed on `loss * rv + loss * chair + rv * chair` within $|rv| \leq 11.65$ percentage points (effective $N = 385,120$; 921 races); race-clustered standard errors. The interaction coefficient `loss  $\times$  chair` is the chair-led increment to the `loss=1` effect on the binary headline.

Interaction	Coefficient	SE	95-percent CI	p
<code>loss $\times$ is_full_cmte_chair</code>	+0.007	0.094	[-0.177, +0.191]	0.94
<code>loss $\times$ is_any_chair</code>	-0.063	0.083	[-0.225, +0.099]	0.45

ent bandwidths (chair-led subgroups are bandwidth-narrow because the close-election sample contains few chair-led races), and the any-chair effect is partially attributable to ceiling-saturation on the winner-side limit: chair-led winners attract very high baseline PAC giving, and the `rdrobust` point estimate (right-limit minus left-limit) absorbs that base-rate gap rather than identifying redirection magnitude per se.

Both interaction coefficients are statistically zero, and both 95-percent confidence intervals exclude amplification of redirection by winner-side power larger than about half the headline effect (the upper bounds sit near +0.10 to +0.19 against the +0.354 headline). The full-committee chair interaction is +0.007 (the same loss effect for chair-led and non-chair-led winners); the any-chair interaction is -0.063 (chair-led winners show a marginally stronger drop, directionally opposite to pull amplification). The continuous LPM is preferred over the subgroup `rdrobust` for adjudication because the two subgroups are estimated at the same bandwidth and on the same baseline, removing the ceiling-saturation artifact that distorts the subgroup `rdrobust` on any-chair winners.

The chair test is statistically conservative on the precision side: chair-led winners are a small slice of the close-election sample (committee chairs typically hold safe seats), and the interaction confidence intervals span roughly ± 0.18 . A pull mechanism that amplified redirection by 10 percentage points among chair-led winners would lie within the confidence interval and be undetectable here. What the chair test cleanly excludes is a pull mechanism of the magnitude that “winner staff solicitation drives most of the redirection” would imply: such a mechanism would generate interaction coefficients of +0.20

or larger, well outside the estimated confidence intervals.

Combined verdict. Both the incumbency-type decomposition (N.1) and the committee chair interaction (N.2) point in the same direction: redirection magnitude is insensitive to plausible variation in winner-side solicitation capacity. The chair coverage limitation (Congresses 103-115; focal cycles 1992-2016) and the sample-size constraint inside the close-election bandwidth leave room for a small pull component the design cannot detect, but the two tests jointly rule out the pull mechanism as the dominant driver of the +0.354 winner-redirection. The framework in the Theory section identifies the push reading as the operative description of the access channel; the empirical evidence reported here supports that reading.

O Mechanical vs spillover decomposition of the winnerward partisan-composition outcome

The winnerward partisan-composition outcome `delta_toward_winner` is computed on the partisan composition of the PAC's **full House contribution portfolio** in cycle t and cycle $t+2$. The two focal-race candidates are themselves part of that portfolio. The +6.85-percentage-point pooled discontinuity (body Table 1, column 4) therefore has two analytically distinct components: a mechanical component driven by the PAC's reallocation within the focal race itself, and a spillover component driven by partisan reorientation across the PAC's other House contributions. Decomposing the headline into these two components matters for the substantive interpretation reported in §3 and §4.1: the partisan-loyalty falsification is identified on the spillover component (loyalty predicts no shift in other-race partisan composition when a single race flips), while the focal-race component is already captured by the first two co-primary outcomes. This appendix decomposes the headline into the two components and confirms that the spillover component carries the load-bearing inference.

Construction. I re-extract the DIME contribution database to recover the exact amount each PAC gave to the focal-race winner (`amt_winner_t`) and focal-race loser (`amt_loser_t`) in cycle t . The cycle- $t+2$ amounts (`amt_winner_t2`, `amt_loser_t2`) are already in the analytic panel. For each PAC-by-cycle observation, I subtract the focal-race contributions from the numerator and denominator of the PAC’s Democratic-share computation, allocating the contribution to the D side or R side according to the focal candidate’s party:

$$\text{dem_share_excl}_{i,t} = \frac{\text{dem_giving}_{i,t} - \text{focal_dem}_{i,t}}{(\text{dem_giving}_{i,t} - \text{focal_dem}_{i,t}) + (\text{rep_giving}_{i,t} - \text{focal_rep}_{i,t})},$$

where $\text{focal_dem}_{i,t}$ is the PAC’s contribution to whichever of the two focal candidates is a Democrat (zero if both are Republican), and analogously for $\text{focal_rep}_{i,t}$. The focal-excluded delta is the difference of focal-excluded shares across cycles, and the focal-excluded winnerward outcome `delta_toward_winner_excl` applies the same sign normalization as the primary outcome (positive in D-winning races, negative in R-winning races on the raw delta; sign-aligned to mean “movement toward the focal-race winner’s party in the PAC’s non-focal House portfolio”). The re-extract recovers exact contributions to both focal candidates in cycle t , including the non-dominant amounts of dual-pick PACs that the primary build’s dominant-amount rule discards: 18,454 PAC-race observations are recovered as dual-pick (where both `amt_winner_t` > 0 and `amt_loser_t` > 0).

Focal-race share of the PAC’s portfolio. The mechanical component is bounded by the share of the PAC’s total D and R giving that flowed to the focal-race candidates. Within the MSE-optimal bandwidth window ($|rv| < 13$ pp; $N = 582,759$), the mean focal share is 11.3 percent at cycle t and 2.8 percent at cycle $t+2$. The asymmetry reflects portfolio diffusion: the focal-cycle contribution is concentrated on the focal candidate by construction,

Table A11: Decomposition of the winnerward partisan-composition discontinuity into mechanical and spillover components.

Sign-normalized winnerward outcome `delta_toward_winner` and its focal-excluded counterpart `delta_toward_winner_excl`, `rdrobust` with triangular kernel and race-clustered standard errors; MSE-optimal bandwidth re-selected for each outcome. Magnitudes reported as $|\hat{\beta}|$ following the sign convention documented for the primary outcome in body §3.

Outcome	$ \hat{\beta} $	SE	p (robust)	95-percent CI	h (pp)	N effective
<code>delta_toward_winner</code> (full portfolio)	0.0693	0.0047	$< 10^{-46}$	[0.0592, 0.0778]	11.1	437,092
<code>delta_toward_winner_excl</code> (focal-race removed)	0.0558	0.0046	$< 10^{-31}$	[0.0456, 0.0638]	11.4	437,154
Mechanical (A) = full – excl	0.0135	—	—	—	—	—
Spillover (B) = excl	0.0558	0.0046	$< 10^{-31}$	[0.0456, 0.0638]	11.4	437,154

while two cycles later the PAC’s giving spreads across many other House races. Even at the higher cycle- t share, focal contributions are a minority of the PAC’s House portfolio: the partisan-composition outcome is fundamentally a portfolio measure, not a focal-race measure, regardless of how the decomposition lands.

Decomposition reading. The mechanical component (A) accounts for 19.5 percent of the full pooled discontinuity (1.35 of 6.93 percentage points), and the spillover component (B) accounts for the remaining 80.5 percent (5.58 percentage points). The spillover component is identified at $p < 10^{-31}$, with a 95-percent confidence interval that excludes zero by a factor of about ten relative to its width. The full and focal-excluded estimates use very similar MSE-optimal bandwidths (11.1 vs 11.4 percentage points) and overlapping effective samples, so the difference reflects the substantive contribution of the focal-race reallocation rather than a sample-composition artifact.

Substantive implication for the partisan-loyalty falsification. The full +6.85-percentage-point discontinuity reported in body Table 1 is consistent with two distinct stories. Under a pure-loyalty account, the PAC simply moves to the focal-race winner (which mechanically shifts the portfolio’s partisan composition by the focal-race contribution amount) but maintains its preferred-party tilt across all other House races.

Under an access-driven reorientation account, the PAC reorients across its entire portfolio toward the new winner's party. The focal-excluded decomposition distinguishes these two readings directly: the spillover component is identified on the PAC's giving to the candidates of the focal-race winner's party in non-focal House races, where mechanical contributions from the focal-race reallocation cannot operate. A +5.58-percentage-point spillover ($p < 10^{-31}$) confirms that the partisan reorientation is portfolio-wide, not just a mechanical artifact of the focal-race redirection captured by the first two co-primary outcomes. The partisan-loyalty falsification reported in body §4.1 is identified on this spillover component; the mechanical 1.35-percentage-point component is substantively redundant with the +0.354 / -0.215 discontinuities on the first two outcomes and is not load-bearing for the falsification argument.

P Donor-type heterogeneity and the individual-donor comparison

This section supports body §4.3. It documents (i) the committee-master crosswalk used to split the committee sample by Federal Election Commission committee type and its coverage, (ii) the redirection by committee type, and (iii) the parallel individual-donor design, the within-individual donation-size gradient, and the breadth-neutral robustness checks that distinguish motive from a mechanical donor-breadth artifact.

Committee-type classification. `DIME contributor_type = "C"` is the full set of registered committee filers, not a corporate subset; its only sibling value is `"I"` (individuals). To recover each committee's FEC type I map the DIME committee identifier (`bonica_cid`) to its modal FEC committee identifier (`efec_comid`) and join the FEC committee master, reading committee type (`CMTE_TP`), organization type (`ORG_TP`), and designation (`CMTE_DSGN`). Because `efec_comid` is populated only from 2012 onward while `bonica_cid` is stable across the panel, a committee's 2012-or-later FEC identifier

Table A12: Composition of the committee sample at the MSE-optimal bandwidth.

Share of committee-race observations ($ rv < 14.3$ pp) by FEC committee type.	
Committee type	Share of close-bandwidth observations
Candidate committees (CMTE_TP H/S/P)	22.5%
Party committees (CMTE_TP X/Y/Z)	19.7%
Non-connected / ideological (N/Q, no sponsor)	3.6%
Leadership PACs (CMTE_DSGN = D)	1.7%
Trade (ORG_TP = T)	1.5%
Membership / labor / cooperative (ORG_TP M/L/V)	0.8%
Corporate (strict, ORG_TP = C/W)	0.7%
Unmatched or unclassified (mostly pre-2012)	49.6%

backfills the classification of its earlier-cycle contributions. Coverage is therefore era-dependent: 86 percent of committee-race observations are classified in 2012–2020 but only 22 percent before 2002, and roughly half of all close-bandwidth observations — overwhelmingly committees that ceased filing before 2012 — remain unclassified. The CMTE_TP field (which distinguishes candidate and party committees from PACs) is well populated where present; ORG_TP (which marks a connected committee’s sponsor type as corporate, labor, trade, membership, or cooperative) is frequently blank, so a committee positively coded corporate by ORG_TP is a strict lower bound on business committees. I recover a broader business-connected set from the presence and text of the connected-organization name (CONNECTED_ORG_NM).

Party and candidate committees together account for roughly two-fifths of all close-bandwidth observations; positively coded corporate committees are under one percent, which is what makes the v4 “corporate PAC” sample label a misnomer.

Redirection by committee type. Each row of Table P2 is the body’s headline specification (rdrobust, triangular kernel, local-linear, MSE-optimal bandwidth, race-clustered standard errors) run on the indicated committee-type subset. The pooled re-run on the classified sample reproduces the body headline exactly (+0.354), validating the classification pipeline. Every committee type redirects toward the winner by between +0.25 and +0.38; non-connected ideological committees redirect the most. Excluding party and can-

Table A13: Winner redirection, loser-loyalty drop, and partisan reorientation by committee type.

gave_to_winner_t2, gave_to_loser_t2, and the sign-aligned partisan-reorientation outcome; magnitudes in the body's $|\beta|$ convention.

Committee type	→winner	→loser	reorient
Non-connected ideological	+0.376	−0.209	0.057
Party committees	+0.377	−0.207	0.042
Leadership PACs	+0.333	−0.200	0.082
Candidate committees	+0.330	−0.185	0.078
Labor	+0.315	−0.173	0.031 (n.s.)
Membership	+0.307	−0.133	0.095
Trade	+0.274	−0.154	0.120
Corporate (strict ORG_TP = C, 640 committees)	+0.246	−0.122	0.140
Business-connected (recovered, 3,150 committees)	+0.353	−0.197	0.070
PAC-only (excludes party + candidate committees)	+0.334	−0.187	0.081
Pooled (= body headline)	+0.354	−0.215	0.069

didate committees (the “PAC-only” row) leaves the redirection essentially unchanged. One genuine type difference appears on partisan reorientation: business-connected types reorient toward the winning party (≈ 0.07), whereas labor reorients negligibly (0.031, $p = 0.12$) — labor follows winners but does not switch party.

The strict-corporate row (ORG_TP = C only) is a noisy 640-committee slice and understates business giving; the recovered business-connected estimate (+0.353, 3,150 committees) is the more reliable business figure, and the body reports the recovered estimate.

Individual-donor panel. I build the individual panel from DIME contributor_type = "I" using the identical race definitions and dominant-amount pick rule as the committee panel, so that only the donor universe differs. The races are extracted from the committee analytic panel itself, making the race set byte-identical. Because the individual data are large (roughly 22 million itemized contributions in 2020 alone), I assemble the panel — aggregate per donor-race, assign the pick, and join the t+2 outcomes — in a disk-backed columnar database and collect only the matched donor-race observations. The panel covers focal cycles 1990 through 2018 (the last cycle whose t+2 outcomes are observed) and comprises 7.88 million donor-race observations across 2.97 million individual donors. For apples-to-apples comparison I re-run the committee design on the

Table A14: Individuals versus committees, and the within-individual donation-size gradient.

Committee column is the committee design re-run on the shared 1990–2018 window. Size bands classify individuals by total focal-cycle giving (small < \$200).

Outcome / donor group	→winner	reorient
Individuals (all)	+0.241	0.027
Committees (1990–2018 window)	+0.358	0.073
Individuals, small (< \$200)	+0.148	0.002 (n.s.)
Individuals, mid	+0.276	0.047
Individuals, large	+0.414	0.063

same 1990–2018 window; the committee winner redirection on that window is +0.358, reproducing the full-sample headline (+0.354) and confirming the window restriction does not drive the committee–individual gap.

Breadth-neutral robustness. The binary `gave_to_winner_t2` gradient could be a mechanical artifact: donors who give to more candidates are more likely to give to any one of them, the winner included. Three checks establish that the gradient reflects motive rather than donor breadth. First, the partisan-reorientation outcome is a share of the donor’s portfolio and is therefore invariant to how many candidates the donor supports; on it the gradient is small 0.002, mid 0.047, large 0.063 (committee 0.073) — small expressive donors do not realign at all, large donors realign at committee levels. Second, small donors are not less active at t+2 than mid-sized donors (cycle-t+2 activity rates: small 0.59–0.60, mid 0.41–0.47, large 0.55–0.66), so the weak redirection of small donors is not attrition or churn. Third, conditioning on donors active at t+2 — which removes the breadth/retention channel at the cost of post-treatment conditioning — leaves the gradient intact and sharper (small +0.246, large +0.626). The breadth-neutral partisan-reorientation share is the primary evidence; the conditional-on-active estimate corroborates it. Scripts: `code/33_pac_type_heterogeneity.R`, `code/34_individual_donor_rdd.R`, `code/35_donor_gradient_validation.R`.

Q Total-giving null: the exit-robust budget battery

The body's third co-primary outcome, the change in log total House giving, is a precise zero at the cutoff (+0.045 log-units, SE 0.138; Table 1, column 3), but its 95 percent confidence interval is wide: the upper limit, +0.312 log-units, corresponds to $\exp(0.312) - 1 \approx 37$ percent expansion, wide enough that a budget-expansion hybrid could in principle inhabit it. This appendix shows that the width is an artifact of the logarithm amplifying committee exit rather than genuine imprecision about the quantity the hybrid predicts, and that on every margin matching the hybrid's prediction the budget is fixed to within roughly five percent.

Roughly 14 percent of committees make no House contribution at all by $t+2$ (the active-at- $t+2$ rate is 0.859). For each such committee the log change is $\log(1) - \log(\text{total}_t + 1) = -\log(\text{total}_t + 1)$, a large negative number; the full-sample delta-log distribution has mean -1.44 and standard deviation 3.47 , and this heavy left tail — committees disengaging entirely — is what inflates the confidence interval. Exit is the opposite of the expansion a hybrid requires, so the wide unconditional interval does not keep the hybrid alive. Three exit-robust reformulations isolate the expansion question directly.

The log-change family (rows A–F) holds the expansion bound in the 31–40 percent range throughout, including under the two analysis-of-covariance specifications. The analysis-of-covariance result is diagnostic: controlling for the committee's baseline giving level barely moves the standard error (0.138 to 0.139), which shows the residual variance is not cross-committee level heterogeneity (which the baseline absorbs) but the within-committee exit channel (which it does not). Winsorizing at the first and ninety-ninth percentiles does not tighten the bound either, because the exit mass is roughly 14 percent of committees, not one percent.

The exit-robust family (rows E, G, H) collapses the bound to roughly five percent. Conditioning on committees that remain active at $t+2$ gives a loss effect of -1.4 percent (expansion ruled out above 3.4 percent). A within-committee fixed-effects specification

Table A15: Total-giving null across eight specifications.

Rows A–F keep the committee-exit mass in the outcome and vary the transform (delta log, inverse-hyperbolic-sine change, two analysis-of-covariance specifications controlling for the baseline giving level, and a winsorized delta log); rows E, G, H remove or bound the exit mass (conditioning on committees active at $t+2$, within-committee fixed effects, and the share-of-budget percentage change winsorized at $p1/p99$). Specifications A–F, E, H are sharp `rdrobust` (Calonico, Cattaneo and Titiunik 2014; triangular kernel, MSE-optimal bandwidth, race-clustered HC1, robust bias-corrected CI); G is a within-committee LPM (`fixest : feols`, committee and cycle fixed effects) with the running variable interacted with the loss indicator. “Expansion bound” is the upper 95 percent confidence limit — the side the budget-expansion hybrid predicts — translated to a percentage via $\exp(\cdot) - 1$ for the log/IHS rows and directly for the share row.

Spec	Formulation	Point (SE)	p	95% CI	Expansion bound	N_h
A	$\Delta \log$ total giving (baseline)	+0.045 (0.138)	0.76	[−0.227, +0.312]	36.7%	711,153
B	Δasinh total giving	+0.053 (0.146)	0.71	[−0.231, +0.340]	40.4%	717,429
C	ANCOVA $\log(t+2)$ on baseline log	+0.020 (0.139)	0.99	[−0.272, +0.274]	31.5%	655,164
D	ANCOVA $\text{asinh}(t+2)$ on baseline	+0.028 (0.147)	0.94	[−0.276, +0.298]	34.7%	665,368
F	$\Delta \log$ win-sorized ($p1/p99$)	+0.046 (0.137)	0.74	[−0.223, +0.313]	36.7%	712,848
E	$\Delta \log$ active at $t+2$	−0.014 (0.028)	0.46	[−0.075, +0.034]	3.4%	470,644
G	within-committee FE, loss coefficient	−0.003 (0.024)	0.90	[−0.050, +0.044]	4.5%	711,153
H	share-of-budget % change, win-sorized	−0.001 (0.026)	0.96	[−0.052, +0.049]	4.9%	784,215

— comparing a committee’s own close wins to its own close losses, which absorbs each committee’s exit propensity (adjusted $R^2 = 0.54$) — gives −0.3 percent (expansion ruled out above 4.5 percent). Expressing the change as a share of the committee’s prior budget,

which caps an exit at -100 percent rather than sending it to $-\infty$, gives -0.1 percent (expansion ruled out above 4.9 percent). All three are precisely estimated near-zeros.

The within-committee specification is identified off the 18,731 committees (28.9 percent of the close-bandwidth sample) that appear in close races on both the winning and losing sides — the same multi-race sub-population that identifies the within-PAC fixed-effects sensitivity (Appendix Section M) — and so inherits its sub-population caveat; the active-only specification conditions on a post-treatment survival indicator, defensible because differential attrition at the cutoff is small and roughly balanced (Section J) but post-treatment conditioning nonetheless. Neither is decisive alone. Their convergence with the share-of-budget margin — which is full-sample and conditions on nothing post-treatment — is what licenses the conclusion: three reformulations that remove the exit tail in three different ways all return a change within five percent of zero. The redirection toward the winner is financed by substitution within that near-fixed budget, the budget-constrained allocation of the model in Section 2. The bound is on the average loss effect on total giving; it does not bound heterogeneity, so a minority of committees could expand while others contract. Script: `code/31_total_giving_null.R`.

R The lobbying channel: directed-lobbying redirection on the public-firm subsample

The body reports that the redirection toward the winner appears in a second, behaviorally independent channel — directed lobbying — on the subsample where lobbying can be observed. This section gives the construction and the full battery.

Lobbying expenditures are observable, and matchable to specific legislative targets, only for committees sponsored by publicly traded firms. I build a firm-by-focal-race panel that mirrors the headline committee design exactly: the unit is a sponsoring firm (Compustat `gvkey`) crossed with a contested two-candidate House focal race; the run-

ning variable is the signed vote-share margin from the firm’s perspective (positive if the firm’s previously supported candidate won the focal race, negative if that candidate lost), with the cutoff at zero and a dominant-amount pick rule for firms that gave to both top-two candidates. Prior support is matched on the Federal Election Commission candidate identifier (`CAND_ID`), observed for 100 percent of both winners and losers in close focal races; matching on a congressional-service identifier instead would restrict the loser-backer arm to former incumbents (30 percent coverage) and induce differential selection at the cutoff. The outcome is the firm’s directed lobbying of the focal winner over the $t+2$ cycle window, summed from the firm-by-legislator-by-year lobbying panel of Kim et al. (2026), zero-filled when the firm did not lobby the target (the same zero-treated-as-zero convention used for the headline outcomes). Restricting to focal cycles 2008 through 2016 (so the $t+2$ outcome window closes by 2018, the panel’s last year) yields 116,718 firm-by-race observations across 1,074 firms and 1,669 races; the near-cutoff sample at the primary bandwidth is ample (effective $N \approx 26,500$). Because the binary “did the firm lobby the winner at all” indicator is saturated (roughly 69 percent), the primary outcome is lobbying *intensity*, the log of one plus directed lobbying dollars.

The primary directed-lobbying discontinuity is positive, significant, and stable across fixed bandwidths: +1.16 ($h = 5$), +1.24 ($h = 10$), +1.23 ($h = 14.3$), and +1.17 ($h = 20$), with the p -value improving as the effective sample grows. The two giving cross-checks reproduce the redirection inside the same donation data (+0.17 binary, +1.41 log dollars), confirming the subsample behaves like the headline; the binary cross-check, +0.171, is about half the headline +0.354, as expected for a different sample (public-firm committees, 2008–2016) and a different giving measure, and I do not read the two magnitudes as equal. The placebo — directed lobbying of the focal loser, who is out of office at $t+2$ — is a null (-0.295 , $p = 0.37$), so the discontinuity is specific to the candidate who holds the seat.

Dropping the off-year leaves the primary essentially unchanged (+1.16), so it is not a window-pooling artifact. Conditioning on firms that lobby at all isolates the intensive

Table A16: Directed-lobbying RDD on the public-firm subsample (full panel).

Sharp `rdrobust` (Calonico, Cattaneo and Titiunik 2014; triangular kernel, MSE-optimal bandwidth, race-clustered HC1 standard errors, robust bias-corrected inference), focal cycles 2008–2016. All outcomes are measured toward the focal winner at $t+2$ except the placebo (toward the out-of-office loser). The two cross-checks re-estimate the giving redirection inside the same firm-by-legislator donation data.

Outcome (toward winner, $t+2$)	Point	Robust SE	p	95% CI	h
log directed lobbying \$ (PRIMARY)	+1.245	0.445	0.003	[+0.46, +2.20]	12.3
lobbied a bill the winner sponsored (7.6%)	+0.020	0.012	0.064	[−0.001, +0.044]	15.8
any lobbying of winner (saturated, 69%)	+0.117	0.046	0.006	[+0.04, +0.22]	11.8
log lobbying, winner – loser (net)	+1.549	0.817	0.030	[+0.17, +3.37]	12.0
gave to winner (cross- check, binary)	+0.171	0.031	7×10^{-9}	[+0.12, +0.24]	12.9
log donations to winner (cross- check)	+1.410	0.253	7×10^{-9}	[+0.97, +1.96]	13.1
log lobbying of loser (PLACEBO)	−0.295	0.433	0.37	[−1.24, +0.46]	14.3

margin and if anything strengthens the discontinuity (+1.42), while the giving cross-checks reproduce and the placebo stays null on the active subsample. Directed lobbying of the winner's committees also jumps (+0.15); at base rate 0.64 this directed-and-windowed measure is no longer saturated (unlike the firm-level any-committee-connection indicator, which is saturated at 80 to 94 percent across sitting legislators and is therefore unusable as an RDD outcome), and it corroborates the intensity result.

Table A17: Robustness of the primary directed-lobbying discontinuity.

Same estimator. Block (i) restricts the outcome to the single post-election year; block (ii) restricts the firm sample to the 752 firms with any positive lobbying in the panel (conditioning out the firm-level extensive margin); row (iii) is directed lobbying of a committee on which the winner sits.

Specification	Point	Robust SE	p	h
(i) log lobbying of winner, t+2 year only [full]	+1.164	0.420	0.003	11.8
(i) log lobbying of winner, t+2 year only [active firms]	+1.332	0.484	0.003	11.6
(ii) log lobbying of winner [active firms, PRIMARY]	+1.424	0.514	0.003	12.0
(ii) lobbied a bill winner sponsored [active firms]	+0.024	0.014	0.064	15.7
(ii) gave to winner [active firms, cross-check]	+0.181	0.030	2×10^{-10}	13.5
(ii) log donations to winner [active firms, cross-check]	+1.491	0.243	2×10^{-10}	13.6
(ii) log lobbying of loser [active firms, PLACEBO]	-0.356	0.520	0.36	14.0
(iii) directed committee lobbying of winner	+0.148	0.048	0.001	12.1

The firm-side predetermined covariates — prior giving and split-giving behavior — are balanced at the cutoff. The incumbent-reelected jump is not a firm-side confound but the same mechanical winner-backing asymmetry that drives the McCrary rejection in the headline panel: firms back winning incumbents more than losing challengers, so

Table A18: Running-variable validity on the firm-by-race panel.

Firm-side predetermined covariates run as `rdrobust` outcomes (balanced at $|t| < 2$); density via `rddensity` at the cutoff.

Check	Estimate / statistic	Balanced / continuous
balance: log firm prior giving at t	-0.055 ($ t = 0.71$)	yes
balance: gave to both top-two (dual pick)	-0.007 ($ t = 0.54$)	yes
balance: open seat	-0.007 ($ t = 0.11$)	yes
balance: incumbent reelected	+0.659 ($ t = 9.19$)	no (mechanical)
density: firm-by-race panel	$T = 6.74$ ($p = 2 \times 10^{-11}$)	rejects
density: race-by-side deduplicated	$T = 0.55$ ($p = 0.585$)	continuous

incumbent-reelected races skew to the winner side of the cutoff, and the headline already shows the redirection holds within each incumbency stratum (Appendix Section B.2). The density result is the cleanest demonstration of that mechanism. The firm-replicated panel rejects density continuity ($T = 6.74$), mirroring the headline panel's rejection, because each race contributes one observation per firm-backer and winners attract more firm-backers. But when each race contributes its winner-side margin and its loser-side margin once each — removing the firm-count replication — the underlying race-level margin density is continuous ($T = 0.55$, $p = 0.585$), so the rejection is purely the firm-count replication artifact, a cleaner illustration of the de la Cuesta and Imai (2016) reading than the headline panel provides.

The subsample is 2008–2016 focal cycles on committees sponsored by publicly traded firms; the test is mechanism confirmation on the subset where both giving and lobbying are observable, not an extension of the headline universe, and the lobbying magnitude is not the giving magnitude. Scripts: `code/27_build_lobbying_bridge.R`, `code/28_lobbying_rdd.R`, `code/29_lobbying_rdd_robustness.R`.

S Replication materials

Replication materials — including the analytic panel, all numerical outputs underlying the body and appendix tables, and the canonical-script logs that produced them — are available from the authors on request.