

Television and the Majority-Party Premium in Postwar Congress

Kisoo Kim*

July 2, 2026

Abstract

Broadcast television told postwar American voters almost nothing about their own representative — yet it reshaped House elections. I document that television created accountability without selection. Exploiting the staggered rollout of television across the United States in 1948–1964, I find that each year of cumulative exposure raised majority-party House incumbents' probability of returning to the next Congress by 1.6 percentage points; minority incumbents gained nothing. The premium did not reward individual records. It was unrelated to legislative effectiveness once majority status is controlled, flat across seniority, and attached to majority status itself — following legislators into and out of the majority. Consistent with collective rather than individual attribution, it was largest precisely when Congress produced least. The mechanism is informational substitution: television replaced candidate-level with party-level knowledge, shifting the House vote onto the majority party's brand. Local journalism's contemporary displacement by national digital media implies the same substitution.

Keywords: television; mass media; U.S. Congress; incumbency advantage; party brands

*Email: kisoo@virginia.edu

1 Introduction

Between 1948 and 1964, broadcast television reached nine in ten American households — the fastest diffusion of any mass medium in U.S. history. Its news was produced nationally: three networks covered Washington and the presidency, not 435 district races, and the medium’s arrival displaced the local newspapers and radio that did cover House members (Gentzkow, 2006). The media-accountability literature is accordingly pessimistic about television as an instrument of congressional accountability — media sharpen voters’ evaluation of individual representatives only where coverage gives voters something to evaluate them with (Strömberg, 2004; Snyder and Strömberg, 2010; Prat and Strömberg, 2013) — and its prediction for the rollout is an erosion of candidate-level knowledge, not a change in how House elections are decided. A second tradition expected more: television, a medium of faces and personalities, would personalize electoral politics and reward individually visible incumbents (Ranney, 1983; Wattenberg, 1991; Prior, 2006). The era seemed to bear both out — voters’ knowledge of their representatives fell, and House results pulled steadily apart from presidential results, with split district-level verdicts climbing from a fifth of districts in 1948 to a third by 1964 (Ornstein et al., 2013). Neither tradition predicts what this article documents: television attached a large and durable electoral premium to majority-party status as such — a reward unconnected to who a representative was or what she had done. How does a medium that tells voters nothing about their representative end up organizing the House vote around the party in charge?

The answer this article develops is that television created accountability without selection. Each year of cumulative television exposure raised majority-party House incumbents’ probability of returning to the next Congress by 1.6 percentage points; minority incumbents gained nothing. Yet a convergent set of tests shows that this premium bears none of the marks of individual selection. It is unrelated to legislative effectiveness once majority status is controlled; it does not concentrate among committee chairs but is flat

across seniority quartiles; and it follows the same legislator into and out of the majority at nearly identical magnitude — it attaches to the position, not the person. Each of these facts independently severs the premium from individual records, and together they point to a single conclusion: voters were rewarding what the majority was, not what any member did. A final test makes the collective character of that reward especially vivid — the premium is *largest* when Congress produces the least legislation (the triple interaction with a low-productivity-Congress indicator is -0.021 , $p < 0.001$), the opposite of what credit-claiming for collective output would predict. The mechanism, I show, is an informational substitution: television eroded voters' candidate-specific knowledge while enriching their knowledge of the national party landscape, shifting the basis of the House vote from individual records to the party brand (Snyder and Ting, 2002). A medium that made parties, rather than persons, legible to voters organized House voting around the governing party — resolving the puzzle with which this article began.

Whether a new medium weakens or deepens voters' reliance on party depends on what *kind* of information it conveys. Two cognitively distinct uses of partisan information respond to informational shocks in opposite directions. The **candidate-inference cue** (Type A) uses party as a placeholder for individual-candidate positions; candidate-specific information substitutes for it. The **collective-responsibility cue** (Type B) uses party as an attribution device for chamber outcomes; institutional information about which party governs *enables* it. The distinction follows Snyder and Ting (2002)'s informational theory of party brands, in which the electoral value of a party's label rises with voters' ability to map the party to outcomes it institutionally produces. Broadcast television in this period was a Type B medium: structurally national news content, focused on Washington and the executive branch, oriented to the visible work of the governing party — Speakers, committee chairs, presidential vetoes — rather than to 435 individual House members. A Type B shock should reshape voter knowledge asymmetrically: less about candidates, more about parties.

Direct survey evidence from the American National Election Studies bears out this prediction. Identifying off the staggered state-level rollout, each additional year of cumulative television exposure reduced respondents' ability to recall their House candidate's name by 3.5 percentage points, reduced their ability to identify whether the candidate was an incumbent by 3.9 percentage points, and reduced correct incumbent identification by 1.0 percentage points. The Type B measures move in the opposite direction: perceived party difference rose by 1.0 percentage point per year, and caring about the presidential outcome rose by 1.0 percentage point. Voters in television-exposed regions learned less about candidates and more about parties. A direct voter-level test that interacts each knowledge measure with TV exposure in predicting party-line House voting is consistent with this account, though only partially: all three candidate-knowledge (Type A) interactions are negative, and the party-knowledge (Type B) measure most closely tied to engaged national-outcome voting is positive, while the remaining coefficients are individually indistinguishable from zero (§4.1).

The magnitude of the premium is sizable relative to the incumbency advantage of the era. The 1.6-percentage-point gain in return probability is accompanied by a rise in vote share of roughly 0.6 percentage points per year of cumulative exposure. Gelman-King-style estimates place the U.S. House advantage at roughly two to three percentage points through the 1950s (Gelman and King, 1990; Ansolabehere and Snyder, 2002), so that a single year of cumulative exposure is comparable to the era's typical advantage and two years exceed it, with the substantial, sustained advantage familiar from later decades emerging only after my window, over the late 1960s and 1970s. The scope of the estimate is majority-versus-minority rather than Democrat-versus-Republican: the 78–89 sample contains only two Republican-majority Congresses, too few for a powered cross-party test; the horse-race regression absorbs $TV \times Democrat$ into $TV \times majority$; and a Senate replication (Appendix D.4) gives a same-direction headline of essentially identical magnitude.

The identification strategy exploits the staggered FCC license freeze and post-freeze allocation that drove cross-market variation in cumulative television exposure between 1946 and 1964 (Gentzkow, 2006). An event study of the majority premium shows flat pre-trends before television arrives, television’s eventual arrival timing does not predict outcomes in the pre-television Congresses, the premium survives controls for metropolitan tide-sensitivity and for post-census redistricting, and state-level randomization inference confirms the headline effectiveness result (§3.4). This paper sits at the intersection of two literatures. It supplies the historical natural-experiment grounding for what Hopkins (2018) and Moskowitz (2021), among others, have documented as contemporary nationalization correlates of local-news decline, and it is theoretically anchored in Snyder and Ting (2002)’s brand model of party labels, with the downstream tightening of party discipline I document consistent with cartel theory (Aldrich, 1995; Cox and McCubbins, 2005).

I make three contributions. Theoretically, the Type A / Type B distinction predicts divergent signs for the same information shock and resolves apparent contradictions in the media-effects literature. Empirically, voter-level evidence of asymmetric substitution and legislator-level evidence of a majority-party premium are estimated under one identification. Substantively, I document the mechanism through which broadcast television reshaped U.S. House electoral outcomes, ruling out meritocratic selection, credit claiming for legislative output, chair-visibility amplification, and party-neutral incumbent visibility.

2 Theory

2.1 Two kinds of party cue

The starting point for my argument is a distinction that the literature on party-cue voting has tended to elide. The construct conventionally labeled “the party cue” actually refers to two cognitively distinct uses of partisan information that respond to different

kinds of informational environments and predict opposite responses to the same kind of information shock.

The first use, which I call the **candidate-inference cue** (Type A), is the standard workhorse construct of the political behavior literature. When a voter knows little about a specific candidate, she uses the candidate's party label to predict where that candidate is likely to stand. A Democratic candidate is more likely to favor union rights, social spending, and civil rights legislation; a Republican candidate is more likely to favor business interests, lower taxes, and a hawkish foreign policy. The voter who knows nothing else about Smith but knows that Smith is a Democrat can therefore form an educated guess about Smith's positions by appealing to what she knows about Democrats as a class. This kind of inference is the centerpiece of Zaller's (1992) Receive-Accept-Sample model of mass opinion, in which partisans selectively process political messages through the filter of their own party identification. It is the mechanism behind Lupia's (1994) demonstration that uninformed California voters could mimic the choices of informed voters on insurance reform initiatives by attending to the endorsements of trusted partisan organizations. It is the "low-information rationality" Popkin (1991) describes when he argues that voters use partisan shortcuts to economize on cognitive effort.

The clearest theoretical implication of the Type A cue is that information about specific candidates is a substitute for it. The voter who learns that Smith voted against her party on a key issue, or that Smith has a particular record on a salient policy area, no longer needs the party label to predict Smith's positions; she can read those positions directly. This logical implication is the basis for the long-standing intuition that "more information weakens partisan voting," though the empirical record is more complicated than that intuition suggests: Lenz (2012), for instance, finds that political information often strengthens rather than weakens the alignment between voters' partisanship and their expressed positions. The Type A / Type B distinction I develop below explains this empirical complexity by clarifying that "information" shocks have different effects depending on what

kind of information they convey.

The second use of partisan information, which I call the **collective-responsibility cue** (Type B), operates through a different cognitive process. Here the party label is not a substitute for candidate-specific information but an input to a different kind of judgment: the assignment of credit or blame for collective outcomes. Voters know that legislation requires congressional majorities and that the presidency executes laws differently depending on which party holds it. To reward or punish the parties for the state of the country — for the economy, for war, for the productivity of Congress — voters must know which party is institutionally responsible. The party label functions here as an attribution device, allowing the voter to map collective outcomes to a specific organizational target. This second tradition has its sharpest formal statement in Snyder and Ting's (2002) informational theory of party brands. In their model, the value of a party's label to its individual members is a *state-contingent informational asset*: voters use the label to forecast the party's collective behavior, and the value of doing so depends on the party's institutional position: whether it holds the majority, controls the chamber's agenda, and bears collective responsibility for what the chamber produces. The brand has more informational content, and is therefore more electorally consequential for individual members, when the party occupies a position the voter can map to collective outcomes. Closely allied accounts in the same tradition include Aldrich's (1995) treatment of parties as solutions to collective-action problems within the legislature, Cox and McCubbins's (1993; 2005) demonstration that the majority-party position confers genuine institutional power over the legislative agenda, and Fiorina's (1981) retrospective voting framework. Under all of these accounts, the party label has informational content that depends on the party's institutional position: what the label means depends on whether the party is in the majority, whether it controls the presidency, whether it bears responsibility for current policy outcomes.

The clearest theoretical implication of the Type B cue runs in the opposite direction

from the Type A prediction. Information about which party governs is not a substitute for the cue but an enabler of it: the voter who newly learns that the Democrats control Congress can newly assign Democrats credit or blame for what Congress has produced. Information shocks that target institutional knowledge — that teach voters who governs — should therefore amplify, not weaken, the electoral relevance of the party label.

2.2 The cognitive distinction matters

The distinction between Type A and Type B is not a semantic refinement; it implies that the same nominal “party cue” can be eroded or activated by the same nominal “information shock,” depending on what kind of information the shock conveys. A piece of information that increases candidate-level knowledge weakens Type A while leaving Type B unchanged. A piece of information that increases institutional knowledge strengthens Type B while leaving Type A unchanged. A medium that delivers both kinds of information would have ambiguous net effects. A medium that delivers asymmetric information — more of one kind than the other — would generate predictable asymmetric responses.

Most empirical work on media and electoral behavior has not made this distinction, and the resulting picture has appeared incoherent. One body of work finds that voters in informationally richer environments rely less on party labels — Snyder and Strömberg (2010) on the historical rollout of local U.S. newspapers being the canonical example. A second, growing body of work finds the opposite: voters in nationalized media environments engage in *stronger* party-line voting. Hopkins (2018) documents that American political behavior has nationalized since 2000, with voters in down-ballot races increasingly relying on national partisan signals rather than local incumbent-specific information. Moskowitz (2021) shows that exposure to local news causally raises voter knowledge of incumbents and weakens party-line voting in down-ballot races. Darr et al. (2018) document that newspaper closures sharpen partisan polarization in local voting. Martin and McCrain (2019) show that station-ownership consolidation shifts local news content

toward nationally homogeneous, national-party-aligned coverage. The two literatures have appeared contradictory only because they have been read as opposing claims about a single construct. They are not. The first set is consistent with the Type A theorem: local newspapers deliver primarily candidate-level information, so where local newspapers are richer, the Type A cue weakens. The second set is consistent with the Type B prediction: the media that have eroded over the contemporary period (local newspapers) deliver Type A information, while the media that have replaced them (national cable, partisan online content) deliver Type B information — and the documented strengthening of party-line voting tracks precisely the substitution of the latter for the former. My setting takes this contemporary inversion at its source. The 1948-1964 broadcast television rollout displaced local newspapers and rallies in favor of nationally produced political content; the Type A erosion and Type B amplification I predict for that transition are the same asymmetry the contemporary literature documents for its inverse, and the original natural experiment that produced it.

2.3 Why broadcast television was a Type B medium

Three structural features of broadcast television in 1948-1964 made it a much better delivery mechanism for Type B information than for Type A information.

First, the technology was national in its production. Three networks (CBS, NBC, and the early ABC) produced the bulk of news content, which was distributed identically across the markets that received their signals. The local station purchased a national news feed centered on Washington and the executive branch. Whatever was happening in a particular voter's House district — who was running, what they had said, whom they had endorsed — was almost entirely absent from the network nightly news, which by physical and economic constraints could not produce district-specific reporting for hundreds of separate races. What was present was the institutional drama of Congress: which party held the majority, what major bills were being debated, who was being investigated by

which committee, whether the President's program was advancing.

Second, broadcast television substituted for media that had been more local in orientation. Gentzkow's (2006) influential entertainment-substitution finding documents that the arrival of television displaced newspaper reading and rally attendance — precisely the media that delivered candidate-level information about House races. Local newspapers in this era covered local races in detail, including endorsements, candidate biographies, and campaign-trail reporting. Rally attendance gave voters direct exposure to candidates and their organizations. The voter whose evening shifted from reading the local political column to watching network news did not simply lose information; she traded one kind of information for another. Type A inputs declined; Type B inputs rose.

Third, the framing of network political coverage emphasized partisan conflict at the institutional level. Network news segments on Congress in this era featured the Speaker, the Majority Leader, the relevant committee chairs — the institutional symbols of party control — far more than they featured rank-and-file members. The Sam Rayburns, Joe Martins, and Lyndon Johnsons of the era were television regulars; the typical House backbencher was not. The cumulative effect was to make “which party is running things” a salient and frequently rehearsed fact, while making “what is one's House member doing” a fact that voters had to seek out elsewhere.

These three features predict the same asymmetric pattern. The Type A cue should weaken, if at all, only because the Type A inputs themselves were eroding — consistent with the Zaller-Lupia tradition. The Type B cue should strengthen because the Type B inputs were being delivered through a new and unprecedented channel that, by the mid-1950s, reached the overwhelming majority of American households for several hours a day.

2.4 Empirical implications

The argument yields three testable implications. At the voter level, television exposure should increase institutional knowledge about the parties (Type B) without correspondingly increasing candidate-level knowledge (Type A); plausibly the latter declines. At the legislator level, television exposure should raise the electoral fortunes of majority-party incumbents — through higher reelection probability, larger vote share, and reduced exit — while leaving minority-party incumbents essentially unaffected. And the premium should operate through the party label as such rather than through individual records or institutional roles within the majority: the interaction between television exposure and legislative effectiveness should be null once majority status is controlled, and the effect should not concentrate in the senior, chair-eligible tail of the majority.

The competing accounts make different predictions. Under meritocratic selection, television raises the electoral return to legislative performance, predicting a positive interaction between television and legislative effectiveness on incumbent return. Under chair amplification, television selectively amplifies the institutional visibility of senior majority members who hold committee chairs, predicting that the majority-party effect concentrates in the top quartile of seniority. Under incumbent visibility, the free exposure a new visual medium confers on sitting members advantages incumbents regardless of party (Prior, 2006), predicting a television return premium for minority as well as majority incumbents. Under pure entertainment substitution (Gentzkow, 2006), television depresses all forms of political knowledge and engagement, predicting uniformly negative effects on both Type A and Type B voter outcomes.

These implications are designed to fall together or apart. If the knowledge asymmetry fails to appear at the voter level — if television increases both kinds of knowledge symmetrically or leaves both unchanged — the proposed mechanism for the legislator-level premium is undermined. If the majority-minority asymmetry is absent at the legislator level, there is no premium to explain. If the effect concentrates in chairs or scales with

individual legislative effectiveness, the “party label as such” interpretation collapses into a record-based or role-based story.

3 Research Design

3.1 Television introduction as a natural experiment

I exploit the staggered introduction of television stations across U.S. media markets between 1946 and 1964 as a source of plausibly exogenous variation in media access. The staggered rollout was shaped by two supply-side disruptions orthogonal to local political demand. During World War II, the federal government banned construction of new broadcast facilities, halting the nascent television industry. After the war, rapid growth in station applications overwhelmed the FCC’s technical allocation system, leading to a freeze on new licenses from September 1948 to July 1952. Markets that had received licenses before the freeze experienced television years earlier than markets that waited for post-freeze allocation. This source of variation, driven by regulatory capacity constraints rather than local demand for political information, follows the identification strategy of Gentzkow (2006) and has been used in Campante and Hojman (2013) and Song (2022) on political-behavior outcomes, and in Fenton and Koenig (2025) on labor-supply outcomes.

I measure treatment at the district level as years of cumulative television exposure: the population-weighted average of county-level years since the first television station became available in the county’s designated market area (DMA), using ICPSR study 22720 (Gentzkow and Shapiro, 2008) for first-station dates and DMA-county crosswalks. Congressional districts often span multiple counties belonging to different DMAs with different introduction dates; I aggregate to the district level using 1940 census population weights. The treatment variable ranges from zero to approximately nineteen and a half years across my sample, reaching the upper end only in early-station districts that contained pre-freeze metropolitan markets.

The key identifying assumption is that, conditional on district and congress fixed effects, the timing of television introduction is uncorrelated with unobserved determinants of legislator outcomes. Early- and late-adopting districts may differ in baseline levels, but the fixed effects absorb these. I test the design directly with a pre-television placebo: in the pre-television Congresses (the 78th and 79th, 1943–1946), the eventual timing of television’s arrival, interacted with majority status, does not predict incumbent return ($p = 0.80$; Appendix A.1). An event study of the majority premium itself shows no differential pre-trend by majority status before television arrives (Figure 2), and I report further checks against a metropolitan-tide confound and post-census redistricting in Section 3.4.

3.2 Voter-level data: ANES 1952-1964

For the voter-level analysis I use the American National Election Studies Cumulative Data File 1952-1964, restricted to respondents in the thirty-nine states that overlap with my legislator panel. I construct two families of outcomes.

The Type A (candidate-level knowledge) family draws on three items asked in 1958 and 1964: whether the respondent could recall the names of House candidates running in her district (VCF0972), whether she could state whether either candidate was already in Congress (VCF0977), and whether her identification of the incumbent was correct (VCF0978). The Type B (institutional/party-level knowledge) family draws on three items: whether the respondent perceives an important difference between the parties (VCF0501, available 1952, 1960, 1964), whether she correctly identifies one of the parties as more conservative (VCF0502, available 1960 and 1964), and whether she cares which party wins the presidential election (VCF0311, available throughout). The American National Election Studies did not ask “which party controls Congress” before 1968, so the Type B measures are indirect proxies for institutional knowledge rather than direct measures.

I estimate

$$Y_{ist} = \alpha + \beta \text{years_of_tv}_{st} + X_{ist}\gamma + \delta_t + \varepsilon_{ist},$$

where Y_{ist} is one of the knowledge outcomes for respondent i in state s and election year t , X_{ist} contains demographic controls (age, education, income, race), and δ_t is a year fixed effect. Standard errors are clustered at the state level. State fixed effects cannot be added because television penetration was state-level and saturated by 1956: with state fixed effects, years of cumulative exposure becomes nearly collinear with year fixed effects. Identification therefore rests on cross-state variation in first-TV-year, with year fixed effects absorbing national time trends.

3.3 Legislator-level data: 1943-1966 House panel

For the legislator-level analysis I use a panel of all U.S. House members serving in the 78th through 89th Congresses (1943-1966). Ideology measures and the legislator roster are drawn from Voteview (Lewis et al., 2024). The Legislative Effectiveness Score follows the Volden-Wiseman convention (Volden and Wiseman, 2014) as extended to historical Congresses by Chiou and Goplerud (2024). Historical U.S. House election returns combine county-level data for 1840-1972 from ICPSR 8611 (Clubb et al., 2008) with district-level two-party vote tallies for 1946-2012 (Jacobson, 2013). The analysis sample, restricted to the thirty-seven non-South states, contains 4,036 legislator-congress observations covering 1,075 unique legislators. The 37-state count excludes the eleven Confederate states from the 48-state legislator panel; the voter-level ANES analysis above retained these states (hence the 39 panel-overlap states reported there) because the candidate-knowledge / party-knowledge mechanism I test is not subject to the one-party-South scope condition that motivates excluding the South at the legislator level.

Outcomes include the binary indicator of whether the legislator returns in the next Congress (constructed from Voteview rosters), the legislator's vote share normalized to

the share of her own party (so that Democrats and Republicans both have “vote share” measured in the same direction), and the Legislative Effectiveness Score normalized to the congress mean. I define *incumbent return* as the same legislator (`icpsr`) appearing in the next Congress; the variable encompasses reelection, primary loss, retirement, death-in-office, candidacy for higher office (Senate or Governor), and other paths to chamber persistence rather than strictly winning reelection. I construct seniority as the count of consecutive previous Congresses in which the legislator served. Majority-party status is defined per congress: a legislator is “majority” in Congress t if her party held the House majority in Congress t .

I estimate variants of

$$Y_{ict} = \alpha + \beta \text{years_of_tv}_{ct} \times \text{majority}_{it} + \gamma \text{years_of_tv}_{ct} + Z_{ct}\eta + \mu_c + \delta_t + \varepsilon_{ict},$$

where Y_{ict} is the outcome for legislator i in district c in Congress t , Z_{ct} contains district-level demographic controls (percent urban, percent nonwhite), μ_c is a district fixed effect, and δ_t is a congress fixed effect. Standard errors are clustered at the district-by-decade level, which allows within-district serial correlation across consecutive Congresses while preserving cross-decade independence; results are robust across alternative clustering schemes (state, district-only, two-way state $\times$ congress; Appendix D.3), and the randomization-inference test (1,000 state-level permutations) provides a clustering-free confirmation of the headline legislative-effectiveness result (Appendix A.2). The coefficient of interest is the interaction β .

The non-South restriction reflects the structural difference between the two-party North and the one-party South in this period. In Southern districts where general elections were uncontested, the electoral mechanism I describe cannot operate. I report the South-included results as robustness; the South attenuates the effect (spec-curve medians +0.015 full-sample vs +0.018 non-South), consistent with the scope condition.

3.4 Identification credibility

I subject the headline interaction to four checks beyond the district and congress fixed effects. First, an event study of the premium itself. Replacing cumulative exposure with congress-relative-time bins interacted with majority status (Figure 2) shows no differential pre-trend by majority status before television arrives — the pre-period interaction coefficients are jointly indistinguishable from zero (Wald $\chi^2(4) = 3.1, p = 0.54$) — while the premium emerges and grows in the congresses after arrival. This is the event-study analogue of the headline outcome, and its pre-trends are clean. Second, a pre-television placebo: in the pre-television Congresses (the 78th and 79th, 1943–1946), the eventual timing of television’s arrival interacted with majority status does not predict incumbent return ($p = 0.80$; Appendix A.1). Third, the premium is not an artifact of metropolitan, early-television districts being more responsive to national partisan tides: allowing the effect to differ by district urbanization (a percent-urban $\times$ majority interaction) and letting urban districts follow their own congress-specific trajectory (a per-congress urban slope) leaves the interaction essentially unchanged ($+0.015, p < 0.001$). Fourth, the premium survives the redistricting that followed the 1950 and 1960 censuses. Because a district’s boundaries can change across a reapportionment while its identifier does not, I re-estimate with district-by-apportionment-cycle fixed effects, so that a district is never compared to its post-redistricting self; the interaction is $+0.019 (p < 0.001)$. Consistent with the novelty of the medium, the premium is largest in the earliest apportionment cycle and attenuates as television matures; the single-cycle subwindows are individually underpowered, because the cross-district variation in exposure timing that identifies the effect compresses once television is near-universal.

For inference, I report randomization inference for the legislative-effectiveness result, whose small effective number of clusters is the natural worst case. Across 1,000 permutations of the state-level treatment timing, the permutation p -value is 0.001 against a conventional clustered 0.020, confirming that the result is not an artifact of the clustering

scheme. Randomization inference speaks to the number of effective clusters, not to differential pre-trends; the pre-trend question is addressed by the event study and the placebo above.

I do not report Goodman-Bacon (2021) or Sun and Abraham (2021) decompositions for my continuous treatment, because these estimators were developed for binary absorbing treatments and do not fit this design. When exposure is coarsened to a binary television-presence indicator, the Goodman-Bacon and Callaway–Sant’Anna routines cannot be estimated on the resulting panel, and the Sun-Abraham interaction-weighted estimator returns an insignificant point estimate (-0.05 , $p = 0.72$) — as expected, since binarizing discards the exposure-intensity variation that identifies the continuous-treatment effect (Appendix A.3).

4 Results

4.1 Voter-level: candidate knowledge declines, party knowledge rises

The American National Election Studies data provide direct micro-level evidence for the asymmetric information substitution that underlies my argument. Table 1 reports the coefficient on cumulative years of state-level television exposure for each of the six knowledge outcomes, in the preferred specification with year fixed effects, demographic controls, and state-clustered standard errors. Figure 1 displays the coefficient pattern visually.

The pattern is stark. All three Type A measures show negative coefficients, large in magnitude, and statistically significant at conventional levels. Each additional year of cumulative television exposure is associated with a 3.5 percentage point reduction in the probability that a respondent can recall her House candidate’s name, a 3.9 percentage point reduction in the probability that she correctly knows whether the candidate is incumbent, and a 1.0 percentage point reduction in the probability that she correctly

Table 1: TV exposure and voter knowledge: Type A vs Type B

Outcome	Coef.	SE	<i>p</i>	<i>N</i>	Years
<i>Type A — candidate-level knowledge</i>					
Recall House candidate name	−0.0350***	0.0092	< 0.001	2,679	1958, 1964
Knows if cand was incumbent	−0.0391***	0.0096	< 0.001	2,675	1958, 1964
Correctly IDs the incumbent	−0.0096***	0.0025	< 0.001	1,779	1958, 1964
<i>Type B — institutional/party knowledge</i>					
Perceives party difference	+0.0100**	0.0047	0.032	4,147	1952, 60, 64
Knows which party more conservative	+0.0092	0.0073	0.208	2,461	1960, 1964
Cares about presidential outcome	+0.0102**	0.0050	0.043	5,915	1952–64
<i>First stage</i>					
Watched political TV programs	+0.0324***	0.0058	< 0.001	5,643	1952–64

Coefficients on cumulative years of state TV exposure. All specifications include year fixed effects and respondent-level demographic controls (age, education, income, race). Standard errors clustered at the state level. Sample: ANES Cumulative Data File 1952–1964, restricted to 39 states overlapping the legislator panel. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.




Figure 1: Effect of cumulative TV exposure on voter knowledge. Each row shows the coefficient on years of state-level TV exposure with 95% confidence intervals from the preferred specification (year fixed effects, demographic controls, state-clustered SEs). Type A measures (candidate-level knowledge) all show statistically significant negative effects; Type B measures (party-level knowledge) all show positive effects, two of three significant. The first-stage relationship between exposure and television viewing is positive as expected.

identifies the incumbent when asked. Two of the three Type B measures show oppositely signed positive coefficients, comparable in magnitude to the correct-identification decline though smaller than the recall declines: each year of exposure raises perceived party difference by 1.0 percentage point and raises caring about the presidential outcome by 1.0 percentage point. The third Type B measure, correctly identifying the more conservative party, is positive and similarly sized but does not reach conventional significance.

The substantive magnitudes are large, and both patterns are consistent with what survey research has long established about the postwar electorate. Voters have persistently held strikingly little candidate-specific knowledge of their House members: Stokes and Miller found that in the late 1950s most voters knew nothing about either congressional candidate on their ballots (Stokes and Miller, 1962), and Mann, analyzing district surveys two decades later, showed that congressional voting rested on candidate recognition and broad evaluation rather than on detailed recall of records (Mann, 1978). My voter-level results identify a force acting on that low baseline: cumulative television exposure further depressed candidate recall and incumbent identification. The concurrent rise in perceived party difference, in turn, aligns with the growing ideological distinctiveness of the postwar parties documented by Carmines and Stimson (1989) and others. Voters were not becoming uniformly more or less informed; they were trading candidate-specific knowledge for party-level knowledge.

A first-stage check confirms that the treatment is doing the work it should: each year of cumulative exposure raises the probability that a respondent reports having watched political programs on television by 3.2 percentage points, consistent with the established first-stage relationship documented in Gentzkow (2006).

The Type A and Type B effects are estimated on different respondent samples — the relevant ANES questions were asked in different election years — so the within-respondent contrast I describe is a contrast at the level of the regression rather than at the level of the individual respondent. To strengthen the comparison, I restrict to 1964,

the only year in my period in which both Type A and Type B questions were asked, and estimate cross-state effects within that single year. The asymmetric pattern survives. Within 1964, recall declines by 2.3 percentage points per additional year of historical TV exposure, incumbency identification by 3.6 percentage points, and correct identification by 1.2 percentage points; the Type B measures are positive but smaller and not individually significant. The within-1964 magnitudes are smaller than the pooled estimates, which is expected given the much narrower range of the identifying variation in cross-section (TV exposure varies from approximately 12 to 16 years across states in 1964 alone, versus 0 to 16 years in the pooled sample), but the direction is consistent.

To test directly whether the knowledge shift translated into a shift in how voters used that knowledge, I interact each knowledge measure with cumulative TV exposure in predicting party-line House voting — a vote in which the respondent’s House choice matches her own partisan identification. The evidence is partial but points in the predicted direction. All three candidate-knowledge (Type A) interactions are negative, and the strongest, knowing whether the candidate is the incumbent, is marginally significant (-0.023): the party-line gap between informed and uninformed voters narrows under television exposure. Among the party-knowledge (Type B) measures, caring about the presidential outcome interacts positively with television ($+0.014$), so that voters engaged with national outcomes vote more strongly with their party under television. The remaining four interactions are individually indistinguishable from zero, and two of the Type B measures point the wrong way. I therefore read this test as corroborating rather than decisive: the sign pattern is uniform across the three Type A measures and appears again on the Type B measure most closely tied to engaged use of party information, but the individual coefficients are mostly imprecise, and it is the legislator-level evidence that carries the argument’s weight.

The voter-level evidence has two important caveats that I address transparently. First, state fixed effects cannot be added to the preferred specification. By 1956, every state in

my sample had at least one operational television station, so the extensive-margin variation in TV access disappears, and the intensive-margin variation in years of cumulative exposure is essentially collinear with the combination of state fixed effects and year fixed effects. Identification therefore rests on cross-state variation in first-TV-year, which may correlate with state-level differences in urbanization, education, or media-market structure that demographic controls can only partially address. Second, the within-1964 first-stage on watched-TV is null, suggesting that within-year cross-state variation in cumulative exposure may capture more than just exposure intensity. The voter-level evidence is therefore best interpreted as direct micro-evidence for the mechanism that produces the legislator-level patterns I report next, rather than as standalone causal estimates of voter-level effects. The legislator-level evidence carries the primary empirical weight in my argument.

4.2 Legislator-level: the majority-party premium

The legislator-level analysis shows the asymmetric pattern that the voter-level evidence predicts. Table 2 reports the headline interaction estimates for incumbent return, vote share, and incumbent exit, estimated on the panel of legislator-congress observations covering the 78th through 89th Congresses with district and congress fixed effects, district-by-decade clustering, and demographic controls. Figure 2 presents the event study of the majority-party premium in incumbent return, which is flat before television arrives and emerges only afterward.

The interaction between cumulative television exposure and majority-party status on incumbent return is positive and highly significant: +0.016 per year. The estimate implies that an additional year of cumulative TV exposure raises a majority-party incumbent's probability of returning to the next Congress by 1.6 percentage points relative to a minority-party incumbent's probability of returning. The minority-party main effect of TV is small and indistinguishable from zero (the separate-sample minority regression in

Table 2: TV exposure and majority-party premium

Outcome (per yr of TV exposure)	Coef.	SE	<i>p</i>	<i>N</i>
<i>Headline interactions: TV × majority_party</i>				
Incumbent return	+0.016***	0.002	< 0.001	4,036
Own-party vote share	+0.006***	0.001	< 0.001	3,610
Incumbent exit (1 – return)	–0.016***	0.002	< 0.001	4,036
<i>Subsample estimates</i>				
TV × majority — R-majority Congresses (80, 83)	–0.001	0.016	0.97	674
TV × majority — within-legislator FE	+0.019***	0.003	< 0.001	2,799
<i>Separate samples (no interaction)</i>				
TV — majority sample only, return	+0.018**	0.008	0.033	1,973
TV — minority sample only, return	–0.002	0.007	0.772	2,028
TV — majority sample, vote share	+0.007**	0.003	0.011	1,930
TV — minority sample, vote share	–0.004	0.003	0.139	1,961

All specifications include district and congress fixed effects with demographic controls (% urban, % nonwhite), except the within-legislator row, which substitutes legislator fixed effects. Standard errors clustered at the district-by-decade level. Sample: non-South, House congresses 78–89, complete LES + Voteview matching. The within-legislator FE estimate is from the switcher subsample; switcher-subsample estimates with alternative controls range +0.017 to +0.019, and the full-sample legislator-FE triple specification gives +0.022. The R-majority subsample row reports the corrected estimate following revision of the majority-party indicator for the 83rd and 84th Congresses; the small subsample size ($N = 674$ across two Republican-majority Congresses) cannot power a separate cross-party test, and the formal triple TV × majority × Republican-Congress is null ($p = 0.93$). Within-legislator FE specification restricts to the 463 House members observed in both majority and minority status. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.




Figure 2: Event study of the majority-party premium in incumbent return. Each point is the majority $\times$ relative-time interaction coefficient — the majority premium at each congress relative to television’s arrival — with 95% confidence intervals, from a district- and congress-fixed-effects specification with district-by-decade clustered standard errors; the congress before arrival is the reference. Before television arrives the premium is flat and jointly indistinguishable from zero (Wald test on the pre-period interactions, $\chi^2(4) = 3.1, p = 0.54$); after arrival it emerges and grows. The corresponding event study for the downstream legislative-effectiveness outcome, which is sensitive to event-time binning, appears in Appendix A.2.

Table 2 reports -0.002) — a null that also cuts against party-neutral visibility accounts of television incumbency (Prior, 2006). To benchmark the magnitude, Gelman-King-style estimates place the U.S. House incumbency advantage at roughly two to three percentage points through the 1950s (Gelman and King, 1990; Ansolabehere and Snyder, 2002), rising to about five to six points around 1960 before the substantial, sustained advantage of later decades emerged over the late 1960s and 1970s, beyond the end of my window (Cox and Katz, 1996). Against this period-appropriate benchmark, a single year of cumulative television exposure for a majority incumbent (1.6 points) is comparable to the era’s typical advantage, two years exceed it, and three approach even the elevated readings around 1960. The per-year point estimate itself is precisely identified; I do not extrapolate to multi-decade cumulative magnitudes because the effect attenuates over time¹ and the linear extrapolation extends beyond the within-state variation directly observed.

Whether the asymmetry operates uniformly under both Republican and Democratic majorities cannot be cleanly tested in this window. The 78–89 sample contains only two Republican-majority Congresses (the 80th and 83rd), with a combined $N = 674$, too few for a powered cross-party test. In that small subset the TV-by-majority interaction is small and statistically indistinguishable from zero, and the formal triple interaction of TV, majority status, and a Republican-Congress indicator is also null. Failing to reject equality between R-majority and D-majority Congresses is consistent with — but not strong evidence for — a uniform effect across both parties’ majorities. An independent replication in the U.S. Senate over the same 78–89 window (Appendix D.4) yields a $TV \times \text{majority}$ interaction on senator return of $+0.017$, essentially identical in magnitude to the House headline of $+0.016$. The Senate R-majority subset point estimate is positive ($+0.018$), underpowered but pointing in the same direction as the full-sample interaction and the D-majority subset. Taken together, the House and Senate R-majority point estimates are consistent with a uniform cross-party effect, with underpowering rather than mechanism

¹The premium attenuates with accumulated exposure: the quadratic interaction $TV \times \text{majority} \times \text{years}$ ² is -0.002 ($p < 0.001$).

failure as the most plausible explanation for the individual subsample results. A horse-race regression that includes both $TV \times \text{majority}$ and $TV \times \text{Democrat}$ in the same specification provides a cleaner read on whether the underlying mechanism turns on majority status or on partisan identity: majority status absorbs the interaction, with $TV \times \text{Democrat}$ statistically indistinguishable from zero. Once majority status is controlled, party identity adds nothing. The mechanism operates on majority status, not on something specific to Democrats; the R-majority subset is simply too small to test independently in either chamber alone, but the cross-chamber evidence is consistent with the uniformity interpretation.

The vote share results corroborate the binary incumbency outcome. I measure vote share as the share of the major-party vote going to the legislator's own party — so that Democrats and Republicans are measured on the same metric of “own-party vote share” rather than on a Democratic-share scale that would mechanically conflate party effects with majority effects. The TV-by-majority interaction on vote share is +0.006: each year of cumulative TV exposure raises a majority-party candidate's vote share by roughly 0.6 percentage points relative to a minority candidate's vote share. The vote share result is important because it shows the premium is not a discontinuous victory effect; it operates on the underlying margin of victory in each race, indicating that majority-party candidates win more votes in television-exposed districts, not merely that they cross the 50 percent threshold more often.

Decomposing the source of the incumbent-return effect, I find that television selectively reduces majority-party exit (the complement of return). Among majority-party incumbents, each year of TV exposure reduces the probability of exit — defined as not appearing in the next Congress, whether due to electoral defeat, retirement, death, or candidacy for higher office — by 1.8 percentage points. Among minority-party incumbents, the effect is small and insignificant. The interaction between TV and majority on exit is -0.016 , the mirror image of the return effect by construction. Television there-

fore protects sitting majority-party incumbents from departure at a substantially higher rate than it protects sitting minority-party incumbents. The effect is consistent with both reduced electoral defeat and reduced strategic exit (retirement decisions, decisions not to seek higher office); I cannot fully distinguish these subchannels in my data, but the aggregate pattern of “majority members stay longer in television districts” is robustly identified.

A within-legislator specification, restricted to the 463 House members who served in both majority and minority status during my window, provides perhaps the cleanest identification of the headline result. The 1948-1964 period contains two majority transitions in the House — the Republican gain in 1946 (taking effect in the 80th Congress) followed by the Democratic recapture in 1948 (81st), and the Republican gain in 1952 (83rd) followed by the Democratic recapture in 1954 (84th). Members who served across these transitions experienced majority and minority conditions while their personal characteristics, district, name recognition, and ideology remained constant. Within this within-legislator sample of 2,799 observations, the TV-by-majority interaction is +0.019, essentially identical in magnitude to the cross-legislator estimate. The premium is not a property of legislators who happen to belong to particular parties; it is a property of the majority-party position itself, which the same person can move into and out of as national electoral conditions change.

4.3 Mechanism: party label, not individual record

The headline asymmetry between majority and minority incumbents is consistent with several distinct mechanisms that share the same observable signature. Television could amplify the electoral return to legislative effectiveness, with majority members benefiting more because they pass more legislation that voters can credit. It could amplify the visibility of committee chairs and other senior members whose institutional power is concentrated in the majority — and which television, with its tendency to focus on visible

institutional leadership, might disproportionately broadcast. Or it could operate at the level of the party label itself, rewarding all majority members regardless of record or role. Each alternative would have been the obvious story given some plausible reading of the existing literature, and each makes a sharp, falsifiable prediction.

If television enabled voters to identify and reward more effective legislators — the story I initially entertained in an earlier version of this project, and the natural extension of Snyder and Strömberg's (2010) accountability story to the broadcast medium — then the interaction of television exposure with legislative effectiveness should be positive. Voters in television-exposed districts would discriminate among candidates partly on the basis of legislative productivity, and incumbents with strong records would benefit electorally from voters' new ability to evaluate them.

After controlling for majority-party status, the TV-by-LES interaction on incumbent return is null. The triple interaction of TV, majority status, and LES is also null. Television does not amplify the electoral premium for legislative productivity. Notably, the absence of an interaction does not reflect television eroding a previously strong reflection of records in electoral outcomes — the baseline LES-return association in this period is itself small and statistically indistinguishable from zero ($+0.003$, district FE without majority control). Television added a record-independent majority-party premium rather than displacing a record-based reward structure. The picture sharpens when I look at the productivity gradient within the majority. Among majority-party members, those in the lowest LES quartile have a point-estimate gain of approximately $+0.040$ from each additional year of TV exposure (not significant at conventional levels in this underpowered cell), while those in the highest LES quartile have a point estimate close to zero (-0.007). Among minority-party members, even those in the highest LES quartile gain nothing — though they are by construction the legislators with the strongest individual record to project to voters who might be paying attention to legislative output. Television does not reward effective legislators; if anything, it rewards ineffective majority members at least

as much as effective ones. This pattern is inconsistent with any credit-claiming story in which voters use newly acquired information to identify productive members.

A complementary test exploits variation across Congresses in aggregate legislative output. If voters are rewarding majority members for what Congress accomplishes, the majority-party premium should be larger in Congresses that pass more major legislation, where there is more for voters to credit. The opposite is true. The triple interaction of television, majority status, and a low-productivity-congress indicator is -0.021 : the majority-party premium is significantly larger in unproductive Congresses than in productive ones. When Congress is producing little, majority members benefit more from television exposure, not less. The pattern is consistent with attribution to the party as a collective unit — voters are assigning credit or blame to the majority for what it has or has not produced — but inconsistent with credit-claiming for specific accomplishments. A credit-claiming voter has nothing to credit when Congress accomplishes nothing; an attribution voter still knows whom to blame.

A second alternative is that television amplifies the institutional visibility of senior members who occupy committee chairs and other leadership positions. Network news in this era did disproportionately cover institutional leaders. If voters were learning about the work of chairs through television and using that information to reward them, the majority-party premium should concentrate in the senior tail of the majority distribution.

I test this using seniority as a proxy for committee chair status. The proxy is unusually clean for my period: in 1948-1964, House committee chairs were determined by strict seniority within the majority party, with no exceptions for performance, loyalty, or party leadership preference. Within-majority seniority quartiles therefore correspond closely to the institutional hierarchy of chairs, sub-chairs, and rank-and-file members. (Direct committee-assignment data for my period are not available; the Stewart-Woon committee dataset begins in 1993, and the U.S. Party Leadership data have a gap from 1942 to 1994.)

The TV-by-majority effect by seniority quartile is $+0.014$ in the lowest quartile, $+0.017$

in the second, $+0.022$ in the third, and $+0.013$ in the top quartile — all positive and statistically significant, with the top quartile (the chairs) no larger than the bottom quartile (the rank-and-file). A formal Wald test of equality across the four quartiles fails to reject. The triple interaction of television, majority, and a binary indicator for top-quartile seniority is -0.003 and far from significant; the same triple with a continuous standardized seniority measure is essentially zero. I cannot reject equality of the premium across the majority hierarchy; critically, the chair-eligible top quartile is no larger than the rank-and-file bottom. The chair-amplification story would have predicted a strong concentration in the top quartile; instead the top quartile sits below the third quartile in point estimate.

The within-legislator switch evidence reinforces the same conclusion. Among the 463 legislators who served under both majority and minority status during my window, I re-estimate the seniority interaction. The triple interaction of TV-by-majority with top-quartile seniority within this within-legislator sample is -0.002 : null. The within-legislator triple confirms what the cross-section quartile pattern already showed: the premium does not concentrate in chairs, and there is no detectable seniority gradient in either direction.

The cleanest test of the party-label interpretation is the within-legislator switch. The 463 legislators in my sample who served under both majority and minority conditions provide a direct comparison: the same person, observed under both regimes, with all individual characteristics (name recognition, ideology, ability, charisma, district relations) held constant. If the premium is a property of specific legislators rather than of the majority position, the within-legislator effect should be smaller than the cross-legislator effect. If the premium is a property of the position itself, the within-legislator effect should match the cross-legislator estimate.

The within-legislator TV-by-majority interaction on incumbent return is $+0.019$, essentially identical in magnitude to the cross-legislator estimate. When the same legislator's party gains the majority, television exposure increases her probability of returning to

the next Congress by approximately the same magnitude I estimate in the cross-section. When her party loses the majority, the effect disappears. The premium does not travel with the person; it travels with the position. Table 3 collects the key tests.

Table 3: Mechanism tests for the majority-party premium

Test	Estimate	SE	<i>p</i>
<i>Meritocratic selection (TV × LES)</i>			
TV × LES on incumbent return (controlling majority)	−0.0002	0.001	0.84
TV × majority × LES_top (triple)	+0.002	0.005	0.74
TV effect, majority Q1 LES (least productive)	+0.040	0.024	0.09
TV effect, majority Q4 LES (most productive)	−0.007	0.026	0.79
<i>Cross-Congress productivity gradient</i>			
TV × majority × low-productivity-Congress (triple)	−0.021***	0.005	< 0.001
<i>Chair amplification (seniority)</i>			
Q1 (junior, 1 term)	+0.014**	0.006	0.021
Q2 (2–3 terms)	+0.017***	0.006	0.004
Q3 (4–5 terms)	+0.022***	0.007	0.001
Q4 (senior, 6+ terms; chair-eligible)	+0.013**	0.007	0.045
Wald test of equality across quartiles	$\chi^2(3) = 4.58$	—	0.21
TV × majority × senior_top (triple, binary)	−0.003	0.007	0.65
TV × majority × senior_z (triple, continuous)	−0.0001	0.002	0.96
<i>Within-legislator switch</i>			
TV × majority within-legislator FE	+0.019***	0.003	< 0.001
TV × majority × senior_top (within-legislator)	−0.002	0.007	0.79

Outcome variable: incumbent return to next Congress. All specifications include district and congress fixed effects (the within-legislator rows substitute legislator FE for district FE) with demographic controls. SEs clustered at the district-by-decade level. Sample: non-South, House congresses 78–89. Seniority quartiles defined within the majority-party pool (chair-eligibility threshold ≥ 6 terms in this period).

* $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Each of these tests could have rejected the party-label interpretation; none did. Television’s effect on majority-party incumbents is a label phenomenon, not a record phenomenon and not a chair phenomenon. The voter-level evidence in Section 4.1 explains the consistency: voters in television-exposed regions are not learning more about specific candidates (Type A is declining), so they cannot be using individual records to discriminate within the majority. They are learning more about parties (Type B is rising), so they use the most informationally accessible pieces of party knowledge — including which

party holds the majority — to assign collective credit and blame.

The four tests also collectively narrow the residual interpretive space. The remaining channels within the party-label family are themselves close cousins: brand valuation in the Snyder-Ting sense, agenda-control awareness in the Cox-McCubbins sense, leadership visibility, and credit/blame attribution as in Fiorina (1981). I cannot empirically separate these; they are bundled in my aggregate “majority-party premium.” But all four sit within the Type B / collective-responsibility family that my theory predicts television should amplify, and none requires individual-record information that television does not deliver.

4.4 Downstream consequence: legislative effectiveness rises through majority retention

The empirical regularity that motivated this project — that congressional districts with longer television exposure produced more legislatively effective representatives — survives the reframing, but its interpretation changes substantially. The headline result is unchanged: the within-district effect of cumulative television exposure on the Legislative Effectiveness Score of the legislator currently representing the district is +0.042, confirmed by randomization inference across 1,000 state-level permutations. Television districts produced more effective legislators. The question is how.

A natural decomposition separates two channels. The “replacement” channel captures changes in district-level effectiveness that arise because television districts elect different legislators. The “adaptation” channel captures changes that arise because the same legislators behave differently — work harder, write more bills, manage them more effectively — under television-induced electoral monitoring. The two channels are separable in the panel by comparing district fixed effects (which capture both channels) to legislator fixed effects (which capture only adaptation, holding the legislator’s identity constant). Under district fixed effects, I estimate the +0.042 total effect. Under legislator fixed effects,

the adaptation effect is $+0.010$ with a standard error of 0.028 — small and statistically indistinguishable from zero, though imprecisely estimated. The point-estimate difference, approximately 77 percent of the total effect, is the replacement channel.

The replacement channel is real but its interior is not what the meritocratic-selection account would predict. If television were enabling voters to identify and reward effective legislators, the replacement effect should run through the systematic substitution of less effective incumbents by more effective challengers. I test this composition mechanism directly. The share of TV's district-level LES effect that runs through changes in district majority-party share accounts for approximately negative nineteen percent of the total — essentially zero in net contribution, with the wrong sign. Television does not raise district LES by shifting more districts into majority-party hands.

What does drive the replacement effect is quality selection in retention. Television districts retain higher-LES members of both parties (within the majority, $+0.073$; within the minority, $+0.034$). Combined with the majority-retention finding from Section 4.2, the majority-side selection produces the district-level LES rise through a two-step process: television districts are more likely to retain majority members, and within the majority pool, the retained members are higher-quality. Neither step requires generalized meritocratic selection at the level of the broader candidate pool. Both steps are consistent with the Type B mechanism, in which voters use majority status as the dominant filter and the residual quality differentiation — present in both parties — occurs through local information that survives television's substitution effect.

The decomposition has a built-in check. DW-NOMINATE scores, by construction, do not vary within a legislator's career, so the legislator-fixed-effects comparison is estimable only for outcomes that vary within career — like LES. The small LES adaptation estimate therefore reflects within-career behavior the design can actually detect, not a mechanical artifact of career-fixed scoring.

4.5 Downstream consequence: party discipline tightens as the brand becomes more valuable

A second downstream consequence concerns the ideological behavior of majority-party members. The interaction between television exposure and majority status on a measure of party loyalty (deviation from the majority-party median on roll-call votes) is -0.002 per year at the district-fixed-effects level. The result is robust under both common roll-call ideology measures (DW-NOMINATE career-fixed ideal points and Nokken-Poole congress-varying ideal points). Majority-party members in television-exposed districts vote closer to the party median than majority-party members in non-exposed districts. The effect is concentrated in the majority: minority-party members do not show comparable convergence.

This pattern is the natural complement of the brand-amplification mechanism. The party label functions as a collective brand whose value to any individual member depends on its consistency. Snyder and Ting's (2002) informational-brand model predicts that as the brand becomes more electorally valuable, parties gain stronger incentives to enforce brand consistency — through nomination, recruitment, and primary discipline that select more party-loyal majority candidates, and potentially through within-member adaptation as well. My data point primarily to the selection channel. The Nokken-Poole-based district-fixed-effects coefficient (-0.002) attenuates and loses significance when I condition on legislator fixed effects within the switcher subsample (-0.001); DW-NOMINATE cannot make this within-legislator comparison because DW scores are career-fixed by construction. Television made the majority-party label more valuable; the party-discipline tightening that followed operated mainly through *who served* rather than through how the same incumbents voted.

The ideology results in this paper exhibit the same selection-versus-adaptation split, with an interesting twist. At the cross-sectional district-fixed-effects level using DW-NOMINATE (career-fixed ideal points), television exposure is associated with majority

Democrats located closer to the chamber center (-0.008 for majority Democrats specifically, on absolute distance from zero) — television districts elect more centrist Democrats. At the within-legislator level using Nokken-Poole congress-varying ideal points, the same Democrats who serve under more cumulative television exposure shift modestly leftward (-0.009 per year). The two channels run in different directions: the replacement channel substitutes more centrist Democrats into majority service, while the adaptation channel finds the surviving Democrats drifting toward the (more liberal) Democratic median over their own careers. Both channels reinforce party-brand coherence — TV nationalizes the Democratic brand, and individual Democrats align more closely with that national brand over their tenure even as the marginal new Democratic incumbent is drawn from a more centrist subset of the candidate pool. The result is consistent with the Type B mechanism, with brand alignment operating through both channels — selection on the discipline margin, and within-member ideological drift toward the national party position among Democrats.

The party-discipline tightening complicates simple narratives about television's democratic effects. From one angle, it represents a kind of accountability — voters' newly active brand attribution induces majority members to behave consistently with their party's position rather than to free-ride on the party label while pursuing idiosyncratic preferences. From another, it represents a constraint on representation: members face stronger incentives to vote with their party than to vote with their districts on issues where the two diverge. The same mechanism that enables collective-responsibility attribution also restricts the dimensional flexibility of individual representation.

Table 4 collects the downstream-consequence estimates.

Table 4: Downstream consequences: legislative effectiveness and party discipline

Outcome	Coef.	SE	<i>p</i>
<i>Legislative effectiveness (LES, normalized to congress mean)</i>			
TV → district LES (district FE, total effect)	+0.042**	0.018	0.020
— randomization inference <i>p</i> (1,000 perms)	—	—	0.001
TV → within-legislator LES (legislator FE, adaptation)	+0.010	0.028	0.727
implied replacement share	77%	—	—
TV → majority-share mediation channel	~ −19%	—	—
TV → within-majority LES (selection within majority)	+0.073**	0.036	0.045
<i>Party discipline (distance from majority-party median)</i>			
TV × majority → loyalty (DW-NOMINATE, district FE)	−0.0021***	0.0006	< 0.001
TV × majority → loyalty (Nokken-Poole, district FE)	−0.0021**	0.0007	0.002
TV × majority → loyalty (Nokken-Poole, legislator FE, switchers)	−0.0013	0.0009	0.177
<i>Democrat ideology shift (selection vs adaptation)</i>			
TV → DW-NOM Dems (district FE, replacement-side)	+0.0073**	0.0031	0.021
TV → Nokken-Poole Dems (legislator FE, within-Dem adaptation)	−0.0094**	0.0039	0.015
TV → majority Democrat DW-NOM (district FE)	−0.0080**	0.0033	0.014

LES regressions: non-South legislator-congress panel, district or legislator FE as indicated, district-by-decade clustered SEs. Loyalty and ideology regressions: same panel, district + congress FE. Negative loyalty coefficient indicates convergence toward party median (less deviation). * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

5 Conclusion

Broadcast television in postwar America did not teach voters about their representatives. It taught voters who was in charge. That asymmetric information substitution — Type A knowledge replaced by Type B knowledge — translated into an electoral premium for majority-party members that did not depend on their individual records or their institutional roles. The premium was not concentrated in chairs. It was not amplified for legislatively productive members. It applied to backbenchers and committee chairs alike, to junior and senior majority members alike — and if anything more strongly when Congress produced less, as collective attribution, but not credit claiming, would predict.

The most consequential implication concerns the standard account of how information shocks interact with party-cue voting. A long tradition treats the party label as a cognitive shortcut that voters rely on when candidate-level information is scarce (Popkin, 1991), and a related tradition emphasizes the label as a collective informational asset whose value depends on the party's institutional position (Snyder and Ting, 2002; Grynaviski, 2010). A common reading of the first strand holds that information shocks weaken voters' reliance on the label by supplying the candidate-level information for which the label substitutes. My findings imply that this reading is correctly understood as a statement about a specific kind of information shock. The 1948–1964 television rollout was a powerful information shock that did not increase candidate-level information — it reduced candidate recall and incumbency identification — while increasing institutional information about national two-party politics. The cue that strengthened was not the cue the standard account addresses. There is no contradiction once Type A and Type B cues are distinguished, but the appearance of contradiction has been a recurring source of confusion in empirical work that finds nationalization or brand strengthening in periods of expanding media access (Hopkins, 2018; Moskowitz, 2021). The familiar intuition that more candidate-level information weakens party voting is not wrong; it is incomplete. Information shocks erode party heuristics when they target the candidate-level inputs

that party heuristics substitute for. They activate party heuristics when they target the institutional inputs that party heuristics enable.

This reading also clarifies where my findings sit relative to Achen and Bartels (2016)'s critique of the heuristics-rescue tradition. The candidate-inference task that the Zaller-Lupia-Popkin tradition emphasized is the task Achen and Bartels (2016) argue voters most clearly fail at, and my data are consistent with that diagnosis: the television return premium does not scale with individual productivity, and the least productive majority members benefit from television at least as much as the most productive. What survives in my setting is a thinner capacity — collective-responsibility attribution, closer in spirit to Fiorina (1981)'s retrospective voting — that operates on the coarse institutional fact of which party holds the gavel rather than on what any individual representative has done. Voters appear to perform collective attribution even where they fail at candidate-level inference.

The result is a paradox of accountability without selection. The conventional intuition holds that media exposure improves democratic accountability through individual selection: voters identify good representatives, retain them, and replace bad ones. Television did increase a kind of accountability — but it rewarded majority members who hewed to the party line over those who built independent records, and it penalized minority members regardless of their personal effectiveness or local responsiveness. It produced sharper party brands, tighter party discipline, and stronger collective electoral consequences for governing-party performance, while leaving the already-weak connection between an individual representative's record and her electoral fate unimproved. Whether this is good or bad depends on one's normative theory of representation: trustees of party governance will see health, advocates of personal-vote accountability will see loss. The trade-off itself, however, has been largely missed in empirical media-and-democracy work that treats accountability as a unitary good.

The historical setting has acquired renewed contemporary relevance. The decline of

local newspapers and the rise of nationally oriented cable and digital media closely resemble the Type A erosion and Type B amplification I document — the same substitution, at greater intensity (Hopkins, 2018). To the extent the mechanism generalizes, the present period should exhibit stronger majority-party premia, tighter party discipline, and weaker candidate-level voter information; nationalized partisan voting (Jacobson, 2015), stronger party-line voting in Congress (Lee, 2016), and declining knowledge of local representatives (Hayes and Lawless, 2021) all point in the predicted direction. Debates about media’s democratic role often present “more information” as straightforwardly good. The relevant question is not whether a medium increases information but what kind of information it increases: the medium’s structural orientation, not the quantity of political content it produces, determines which side of the accountability trade-off it tilts toward.

Four limits bound these claims. The aggregate majority-party premium bundles closely related channels — brand valuation, agenda-control awareness, leadership visibility, and credit-blame attribution — that my data cannot separate, though all live within the Type B family the theory addresses. The American National Election Studies did not ask which party controls Congress before 1968, so my Type B knowledge measures are indirect proxies for institutional knowledge. The identification rests on cross-state variation in television timing; the placebo and randomization-inference checks address the leading concerns, but unobserved state-level confounds correlated with adoption timing remain a residual concern that the two-level triangulation mitigates rather than eliminates. And the 1948–1964 broadcast setting is one information environment: the attenuation of the premium as television aged suggests that magnitudes depend on a medium’s novelty as well as its structure, and generalization to other media and other legislative systems requires care.

The mechanism I describe likely operates wherever a new medium nationalizes the information environment more than it localizes it, and wherever it transmits institutional

positioning more than candidate evaluation. Broadcast television offers a long-run window into how durably such a shift reshapes representation. The contemporary moment, in which local newspapers decline as national platforms expand, suggests the question is not closed.

References

- Achen, C. H. and Bartels, L. M. (2016). *Democracy for Realists: Why Elections Do Not Produce Responsive Government*. Princeton University Press, Princeton.
- Aldrich, J. H. (1995). *Why Parties? The Origin and Transformation of Political Parties in America*. University of Chicago Press, Chicago.
- Ansolabehere, S. and Snyder, Jr., J. M. (2002). The incumbency advantage in U.S. elections: An analysis of state and federal offices, 1942-2000. *Election Law Journal*, 1(3):315–338.
- Campante, F. R. and Hojman, D. A. (2013). Media and polarization: Evidence from the introduction of broadcast TV in the United States. *Journal of Public Economics*, 100:79–92.
- Carmines, E. G. and Stimson, J. A. (1989). *Issue Evolution: Race and the Transformation of American Politics*. Princeton University Press, Princeton, NJ.
- Chiou, F.-Y. and Goplerud, M. (2024). Effective lawmaking across congressional eras. *Journal of Politics*, 86(3):997–1012.
- Clubb, J. M., Flanigan, W. H., and Zingale, N. H. (2008). Electoral data for counties in the United States: Presidential and congressional races, 1840–1972. ICPSR Study #08611 [distributor]. County-level historical House election returns, 1840–1972.
- Cox, G. W. and Katz, J. N. (1996). Why did the incumbency advantage in U.S. House elections grow? *American Journal of Political Science*, 40(2):478–497.
- Cox, G. W. and McCubbins, M. D. (1993). *Legislative Leviathan: Party Government in the House*. University of California Press, Berkeley.
- Cox, G. W. and McCubbins, M. D. (2005). *Setting the Agenda: Responsible Party Government in the U.S. House of Representatives*. Cambridge University Press, Cambridge.

- Darr, J. P., Hitt, M. P., and Dunaway, J. L. (2018). Newspaper closures polarize voting behavior. *Journal of Communication*, 68(6):1007–1028.
- Fenton, G. and Koenig, F. (2025). Labor supply and entertainment innovations: Evidence from the U.S. TV rollout. *American Economic Journal: Applied Economics*, 17(4):1–28.
- Fiorina, M. P. (1981). *Retrospective Voting in American National Elections*. Yale University Press, New Haven, CT.
- Gelman, A. and King, G. (1990). Estimating incumbency advantage without bias. *American Journal of Political Science*, 34(4):1142–1164.
- Gentzkow, M. (2006). Television and voter turnout. *Quarterly Journal of Economics*, 121(3):931–972.
- Gentzkow, M. and Shapiro, J. M. (2008). Introduction of television to the united states media market, 1946–1960. ICPSR Study #22720 [distributor]. TV introduction years and DMA boundaries used for years_of_tv construction.
- Goodman-Bacon, A. (2021). Difference-in-differences with variation in treatment timing. *Journal of Econometrics*, 225(2):254–277.
- Grynaviski, J. D. (2010). *Partisan Bonds: Political Reputations and Legislative Accountability*. Cambridge University Press, Cambridge.
- Hayes, D. and Lawless, J. L. (2021). *News Hole: The Demise of Local Journalism and Political Engagement*. Cambridge University Press, Cambridge.
- Hopkins, D. J. (2018). *The Increasingly United States: How and Why American Political Behavior Nationalized*. University of Chicago Press, Chicago.
- Jacobson, G. C. (2013). U.S. House election returns, 1946–2012. Dataset. District-level two-party vote tallies for U.S. House general elections, compiled by Gary C. Jacobson.

- Jacobson, G. C. (2015). It's nothing personal: The decline of the incumbency advantage in US house elections. *Journal of Politics*, 77(3):861–873.
- Lee, F. E. (2016). *Insecure Majorities: Congress and the Perpetual Campaign*. University of Chicago Press, Chicago.
- Lenz, G. S. (2012). *Follow the Leader? How Voters Respond to Politicians' Policies and Performance*. University of Chicago Press, Chicago.
- Lewis, J. B., Poole, K., Rosenthal, H., Boche, A., Rudkin, A., and Sonnet, L. (2024). Vote-view: Congressional roll-call votes database. <https://voteview.com/>. Data accessed for House congresses 78–89; ideology and roster.
- Lupia, A. (1994). Shortcuts versus encyclopedias: Information and voting behavior in California insurance reform elections. *American Political Science Review*, 88(1):63–76.
- Mann, T. E. (1978). *Unsafe at Any Margin: Interpreting Congressional Elections*. American Enterprise Institute, Washington, DC.
- Martin, G. J. and McCrain, J. (2019). Local news and national politics. *American Political Science Review*, 113(2):372–384.
- Moskowitz, D. J. (2021). Local news, information, and the nationalization of U.S. elections. *American Political Science Review*, 115(1):114–129.
- Ornstein, N. J., Mann, T. E., and Malbin, M. J. (2013). *Vital Statistics on Congress*. Brookings Institution Press, Washington, DC. Table 2-16, Ticket Splitting between Presidential and House Candidates, 1900–2012.
- Popkin, S. L. (1991). *The Reasoning Voter: Communication and Persuasion in Presidential Campaigns*. University of Chicago Press, Chicago.

- Prat, A. and Strömberg, D. (2013). The political economy of mass media. In Acemoglu, D., Arellano, M., and Dekel, E., editors, *Advances in Economics and Econometrics: Tenth World Congress*, volume 2, pages 135–187. Cambridge University Press, Cambridge.
- Prior, M. (2006). The incumbent in the living room: The rise of television and the incumbency advantage in U.S. house elections. *Journal of Politics*, 68(3):657–673.
- Ranney, A. (1983). *Channels of Power: The Impact of Television on American Politics*. Basic Books, New York.
- Snyder, Jr., J. M. and Strömberg, D. (2010). Press coverage and political accountability. *Journal of Political Economy*, 118(2):355–408.
- Snyder, Jr., J. M. and Ting, M. M. (2002). An informational rationale for political parties. *American Journal of Political Science*, 46(1):90–110.
- Song, B. (2022). Television and economic voting in US presidential elections. *British Journal of Political Science*, 52(4):1622–1634.
- Stokes, D. E. and Miller, W. E. (1962). Party government and the saliency of Congress. *Public Opinion Quarterly*, 26(4):531–546.
- Strömberg, D. (2004). Mass media competition, political competition, and public policy. *Review of Economic Studies*, 71(1):265–284.
- Sun, L. and Abraham, S. (2021). Estimating dynamic treatment effects in event studies with heterogeneous treatment effects. *Journal of Econometrics*, 225(2):175–199.
- Volden, C. and Wiseman, A. E. (2014). *Legislative Effectiveness in the United States Congress: The Lawmakers*. Cambridge University Press, Cambridge.
- Wattenberg, M. P. (1991). *The Rise of Candidate-Centered Politics: Presidential Elections of the 1980s*. Harvard University Press, Cambridge, MA.

Zaller, J. R. (1992). *The Nature and Origins of Mass Opinion*. Cambridge University Press, Cambridge.

A Appendix A: Identification

A.1 Pre-television placebo

The pre-television placebo asks whether the eventual timing of television’s arrival predicts electoral outcomes before television existed. In the pre-television Congresses (the 78th and 79th, 1943–1946; $N = 680$), the interaction of the district’s eventual TV-introduction year with majority status does not predict incumbent return ($+0.008$, $p = 0.804$). If TV-treated districts had been on a different political trajectory before TV arrived, I would expect a non-zero coefficient; I find none. Implementation: script `17_remaining_gaps.R`.

A.2 Randomization inference

For the headline LES result, I report randomization inference based on 1,000 permutations of the state-level treatment timing. The observed test statistic falls at the 99.9th percentile of the permutation distribution, yielding $p = 0.001$ — consistent with the conventional clustered p -value of 0.020 and reassuring against the concern that the small effective number of clusters understates uncertainty. Randomization inference addresses the effective number of clusters, not differential pre-trends; the pre-trend evidence for the headline majority-premium outcome is the event study in Figure 2 of the main text (implementation: script `07_randomization_inference.R` for the permutation test, script `25_event_study_headline.R` for the majority-premium event study).

Figure A.1 shows the corresponding event study for the downstream legislative-effectiveness (LES) outcome. Unlike the majority-premium event study in the main text, its pre-trend verdict is sensitive to how event time is binned: under the reference binning the district-fixed-effects pre-trends are jointly insignificant ($F = 1.48$, $p = 0.19$) and the legislator-fixed-effects pre-trends pass comfortably, while an alternative binning yields a marginal district-FE pre-trend ($F = 2.84$, $p = 0.04$). Because the LES result

is a downstream consequence rather than the paper’s headline, and because the headline majority-premium event study has clean pre-trends, I treat the LES event study as illustrative and lean the identification argument on the majority-premium event study, the pre-television placebo, and randomization inference. Implementation: script `04_event_study_verification.R`.




Figure A1: Event study of cumulative television exposure on the Legislative Effectiveness Score (downstream outcome). District- and congress-fixed-effects specification; reference period is the congress before television arrives. Pre-trend tests are reported in the text and are sensitive to event-time binning.

A.3 Goodman-Bacon and Sun-Abraham diagnostics

I attempted Goodman-Bacon (2021) decompositions and Sun and Abraham (2021) interaction-weighted estimators, both designed for binary absorbing treatments rather than the continuous exposure I study. Coarsening exposure to a binary television-presence indicator (any exposure versus none) produces degenerate treatment-timing groups under the compressed postwar rollout: the Goodman-Bacon and Callaway-Sant’Anna routines fail to return an estimate, and the Sun-Abraham interaction-weighted

estimator yields an insignificant aggregate ATT (-0.05 , $p = 0.72$; the binary two-way fixed-effects analogue is -0.08 , $p = 0.45$). This is expected: a binary indicator discards the exposure-intensity variation that identifies my continuous-treatment effect. Details: `script 06_staggered_did.R`.

B Appendix B: ANES Type A vs Type B — full specifications

B.1 Robustness across specifications

For each of the six knowledge outcomes, I report estimates under four specifications: (a) bivariate regression on years of TV exposure; (b) plus year fixed effects; (c) plus year fixed effects and demographic controls (preferred); (d) plus state fixed effects (where identification permits).

Specification (d) drops most outcomes because TV penetration was state-level and saturated by 1956; with state fixed effects, years of exposure becomes nearly collinear with year fixed effects. Of the six outcomes, only “cares about presidential outcome” and “perceives party difference” retain identifying variation under state FE; both become statistically indistinguishable from zero, with standard errors an order of magnitude larger than in the preferred specification (cares-about-outcome $+0.007$, SE 0.049 ; party difference -0.044 , SE 0.045), consistent with state fixed effects absorbing the cross-state identifying variation.

B.2 Within-1964 cross-section

The 1964 ANES is the only year in my window with both Type A and Type B knowledge questions. Restricting to 1964 alone, I estimate cross-state effects of cumulative TV exposure on each outcome:

Outcome	Coef.	SE	<i>p</i>	<i>N</i>
Type A: recall House cand name	−0.0225	0.0126	0.074	1,350
Type A: knows if cand is incumbent	−0.0360	0.0128	0.005	1,370
Type A: correctly IDs incumbent	−0.0122	0.0032	< 0.001	947
Type B: perceives party difference	+0.0094	0.0085	0.270	1,382
Type B: knows cons party	+0.0034	0.0053	0.515	1,371
Type B: cares about pres outcome	+0.0079	0.0073	0.278	1,479
First stage: watched political TV	−0.0001	0.0058	0.993	1,381

The within-1964 first-stage on watched-TV is null because by 1964 nearly all respondents in TV-exposed states were watching political TV programs at saturation rates — the cross-state variation in cumulative exposure is no longer driving variation in current viewing intensity. The Type A and Type B effects retain their asymmetric pattern.

C Appendix C: Mechanism — additional detail

C.1 Marginal effects table (4 quartiles by 2 majority states)

The seniority-by-majority interaction in the chair-amplification test implies a 4×2 matrix of marginal TV effects. Computed from the full triple-interaction specification:

	Q1 (junior)	Q2	Q3	Q4 (senior)
Minority	+0.010 (0.007)	+0.002 (0.006)	+0.006 (0.007)	+0.003 (0.007)
Majority	+0.014 (0.006)	+0.017 (0.006)	+0.022 (0.007)	+0.013 (0.007)

Cell entries: marginal TV effect on incumbent return per year of cumulative exposure (standard error in parentheses). Implementation: `script 19_majority_subgroup.R`.

C.2 Open-seat dynamics

A potential concern with the within-legislator switch test is that majority members who lose the majority position may be on a different selection process than incumbents who run continuously. I address this concern with separate analysis of open seats — districts where the previous incumbent did not run for reelection.

Among 1,046 open-seat observations (528 following a majority-party departure, 518 following a minority-party departure), the effect of TV on whether the new member is from the majority party is statistically null in pooled, post-majority-departure, and post-minority-departure subsamples. However, I find a strongly significant effect on party-switching at open seats: TV exposure reduces the probability of an open-seat party switch by 3.4 percentage points per year ($p = 0.004$). Television therefore preserves incumbent-party retention at open seats, consistent with the broader pattern of party-label amplification.

This open-seat analysis was also designed to assess whether the small district-level negative effect of TV on majority share (estimated at -0.016 per year) reflects systematic flipping of majority-held seats to minority parties at moments of incumbent departure.

It does not. The pooled and conditional open-seat results show null effects on majority share, and I conclude the small district-level negative coefficient most likely reflects a measurement artifact arising from the interaction of incumbent persistence with national majority transitions during 1948–1964 (when the same persistent incumbent’s “majority status” variable changes mechanically as national majority changes between the 80th, 81st, 83rd, and 84th Congresses). I report this analysis in detail for transparency but exclude it from the main text. Implementation: `script 20_open_seat_retirement.R`.

C.3 Vote margin by competitiveness

Among majority-party incumbents, the TV effect on own-party vote share is larger in competitive districts (40–65% incumbent vote share) than in safe districts: +0.005 per year ($p < 0.001$, $N = 2,474$) versus +0.004 per year ($p = 0.092$, $N = 1,136$). The point estimate is larger in competitive than in safe districts, though the gap is small and not formally tested; the safe-seat estimate is marginal. The pattern is consistent with the effect operating where it matters electorally — in races that are decided rather than locked in. Implementation: `script 14_vote_margin.R`.

C.4 Within-majority quality selection

Among majority-party legislators specifically, the within-majority effect of TV exposure on individual LES is +0.073 ($p = 0.045$, $N = 2,024$); the within-minority counterpart is +0.034 ($p = 0.006$, $N = 2,012$). Television districts therefore retain higher-LES members of both parties; the majority-side selection is what combines with majority retention to produce the district-level LES rise reported in the main text. I interpret the quality margin as residual local information that survives television’s substitution effect — voters who do learn about specific candidates use that information at the margin to retain effective members. Implementation: `script 13_majority_mediation.R`.

D Appendix D: Robustness

D.1 Specification curve

I construct a specification curve for the headline $TV \times$ majority interaction on incumbent return across 48 specifications, varying: (a) fixed-effects structure (district + congress; legislator + congress); (b) clustering scheme ($CD \times$ decade, state, district, two-way state $\times$ congress); (c) controls (none; demographic — % urban + % nonwhite; demographic + seniority); (d) sample (non-South 37 states; full 48 states).

The interaction is positive in all 48 specifications (median estimate +0.0168, range [+0.012, +0.021]). It is significant at the 10% level in 48 of 48 (100%), at the 5% level in 44 of 48 (91.7%), and at the 1% level in 40 of 48 (83.3%). The four specifications that miss the 5% threshold are all combinations of district + congress fixed effects with two-way state-by-congress clustering, where the test statistic clusters between $p = 0.05$ and $p = 0.10$. Under legislator + congress fixed effects every specification is significant at the 5% level. Implementation: script 21_spec_curve_majority_premium.R.

I separately report a specification curve for the downstream LES result (script 02_e2_les.R), which yields a positive median estimate with 25% of specifications significant at the 5% level, consistent with modest per-spec power. A curve for the uncontrolled $TV \times$ LES interaction (script 03_e3_mechanism.R) is positive in all specifications and significant in half — and that interaction is fully absorbed once $TV \times$ majority enters the model (Table 3), consistent with majority status rather than productivity carrying the premium.

D.2 By-congress heterogeneity

I estimate the $TV \times$ majority interaction separately for each congress in my window. By-congress estimates exhibit substantial volatility consistent with the small per-congress sample sizes (each congress contributes roughly 330 non-South legislator-congress obser-

vations); the pooled triple $TV \times \text{majority} \times \text{Republican-Congress}$ is statistically indistinguishable from zero ($p = 0.93$). The 80th and 83rd Congresses (the only Republican-majority Congresses in this window) are too few in number, and the per-congress samples too small, to support a powered separate test of cross-party generalizability (see Section 4.2 of the main text and the R-majority subsample row in Table 2).

D.3 Alternative clustering

I report the headline $TV \times \text{majority}$ interaction under four clustering schemes: $CD \times \text{decade}$ (preferred), state, district, and two-way (state and congress). Under the primary specification (non-South, district + congress fixed effects, demographic controls), the point estimate is +0.0156 in all four schemes; the conventional p -value is below 0.001 under $CD \times \text{decade}$, state, and district clustering, and is 0.048 under two-way clustering — significant at the 5% level in all four. The randomization-inference test on the headline legislative-effectiveness result (Appendix A.2) provides a complementary, clustering-free benchmark for the design.

D.4 Senate replication

I replicate the headline $TV \times \text{majority}$ interaction in the U.S. Senate over the same 78–89 Congress window. The Senate analysis sample contains 942 senator-Congress observations covering 242 senators in 37 non-South states; after singleton removal under state + congress fixed effects, the regression sample is 788 observations. State-level cumulative TV exposure is the population-weighted aggregate of the district-level measure used in the House analysis. Standard errors are clustered at the state level (statewide elections render district-level clustering moot).

The headline result (+0.017, $p = 0.003$) is essentially identical in magnitude to the House headline (+0.016, $p < 0.001$). The Senate R-majority subset point estimate (+0.018,

Table A1: Senate replication of the majority-party premium

Test	Estimate	SE	<i>p</i>	<i>N</i>
<i>Headline interaction</i>				
TV × majority on incumbent return (state FE)	+0.017***	0.005	0.003	788
<i>Cross-party generalizability</i>				
R-majority subset (Cong 80, 83)	+0.018	0.029	0.53	129
D-majority subset (Cong 78–89, ex 80, 83)	+0.015**	0.006	0.021	659
Triple TV × maj × Republican-Congress	+0.009	0.027	0.75	788
<i>Mechanism tests</i>				
Horse race: TV × majority (with TV × Democrat)	+0.024*	0.014	0.085	788
Horse race: TV × Democrat (with TV × majority)	−0.009	0.015	0.58	788
TV × majority on return (controlling LES)	+0.018**	0.005	0.002	788
TV × LES on return (controlling majority)	−0.002	0.002	0.39	788
Triple TV × majority × LES_top (binary)	−0.013	0.010	0.18	788
<i>Within-legislator</i>				
TV × majority (legislator FE, full sample)	+0.008	0.008	0.28	739
TV × majority (legislator FE, switcher subsample)	+0.012	0.008	0.13	547
<i>Party discipline</i>				
TV × maj on dist from party median (DW, state FE)	−0.0038	0.0023	0.12	787
TV × maj on dist from party median (NP, state FE)	−0.0048*	0.0025	0.066	786
TV × maj on dist from party median (NP, legislator FE)	−0.0001	0.0028	0.98	737

Senate sample: 78–89 Congress, 37 non-South states. State + congress fixed effects; SE clustered at state level. Sample size variation reflects observations dropped by fixed-effect singletons and missing covariates. “Switcher subsample” restricts to the 114 senators observed under both majority and minority status (vs. 463 House switchers). Compare House Tables 2, 3, 4. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$. Implementation: script 22_senate_replication.R.

$p = 0.53$, $N = 129$) is positive and points in the same direction as the full-sample interaction, although underpowered. Combined with the House R-majority subset's near-zero estimate (-0.001 , $p = 0.97$, $N = 674$), the cross-chamber evidence is consistent with a uniform effect across both parties' majorities — with the cross-party test simply underpowered in the available 78–89 window. The triple $TV \times \text{majority} \times \text{Republican-Congress}$ is null in both chambers (House $p = 0.93$; Senate $p = 0.75$).

The horse race shows the same pattern as in the House: $TV \times \text{majority}$ survives controlling for $TV \times \text{Democrat}$ ($+0.024$, $p = 0.085$); $TV \times \text{Democrat}$ is null and slightly negative (-0.009 , $p = 0.58$). The $TV \times \text{LES}$ interaction is null after controlling for majority (-0.002 , $p = 0.39$), and the triple $TV \times \text{majority} \times \text{LES}$ is null ($p = 0.18$) — the same merit-based-selection rejection as in the House.

The party-discipline pattern follows the selection-versus-adaptation split documented for the House. Cross-section party loyalty point estimates are same-signed and larger than the House's under both DW-NOMINATE (-0.0038 , $p = 0.12$) and Nokken-Poole (-0.0048 , $p = 0.066$) measures of distance from the party median, though only marginally significant given the smaller Senate sample. Under legislator fixed effects on Nokken-Poole, the within-legislator effect is statistically indistinguishable from zero (-0.0001 , $p = 0.98$), matching the House finding that the discipline tightening operates through who serves rather than how the same incumbents vote.

The within-legislator FE estimate on incumbent return is statistically insignificant ($+0.008$, $p = 0.28$). Two Senate-specific limits explain the lower power compared to the House. First, six-year Senate terms mean that approximately two-thirds of senator-Congress observations are mid-term, and the binary outcome is mechanically equal to one for those observations (overall sample mean is 0.81). Second, only 114 senators are observed in both majority and minority status during the window (vs. 463 House members), and the switcher-subsample legislator-FE estimate ($+0.012$, $p = 0.13$) is positive but underpowered. The headline state-FE result, however, is essentially identical

in magnitude to the House estimate, and the cross-section party-loyalty estimates are same-signed though only marginally significant.

The Senate replication therefore provides three independent forms of cross-chamber confirmation of the mechanism: (i) the headline majority-premium magnitude is essentially identical; (ii) the merit-based-selection alternative is rejected the same way; (iii) the party-discipline channel shows the same selection-not-adaptation signature, with the cross-section estimates marginal in the smaller Senate sample. The R-majority point estimate further alleviates the cross-party-uniformity concern that the House subsample alone could not resolve. Implementation: script `22_senate_replication.R`.

D.5 Metropolitan-tide confound and redistricting

Two threats specific to the district-panel design deserve separate treatment.

Metropolitan tide-sensitivity. Early-television districts were disproportionately pre-freeze metropolitan markets, and urban districts may respond more strongly to national partisan tides for reasons unrelated to television. Because the within-window majority variation is driven by the wave elections of 1946–1954, such differential tide-sensitivity could in principle generate the $TV \times \text{majority}$ interaction spuriously. I test this by allowing the majority premium to vary with district urbanization (a percent-urban $\times$ majority interaction) and by letting urban districts follow their own congress-specific trajectory (a per-congress urban slope). The interaction is unmoved:

Specification	TV $\times$ majority	SE	p	N
Baseline (Table 2)	+0.0156	0.0024	< 0.001	4,031
+ percent-urban $\times$ majority	+0.0156	0.0025	< 0.001	4,031
+ per-congress urban slope	+0.0149	0.0024	< 0.001	4,031
+ both, and percent-nonwhite $\times$ majority	+0.0149	0.0025	< 0.001	4,031

Redistricting. The panel spans the reapportionments that followed the 1950 and 1960 censuses; a district's identifier can persist across a reapportionment even when its geography changes, so a naïve district fixed effect could compare non-comparable districts. I neutralize this by interacting the district fixed effect with the apportionment cycle (1940-census districts, congresses 78–82; 1950-census districts, 83–87; 1960-census and court-redrawn districts, 88–89), so that a district is never compared to its post-redistricting self. The interaction is +0.019 ($p < 0.001$, $N = 4,030$), if anything larger than the pooled estimate. Restricting instead to a single apportionment cycle attenuates the point estimate as television matures — +0.013 ($p = 0.38$) in the 1940-census cycle (congresses 78–82) and +0.008 ($p = 0.21$) in the 1950-census cycle (83–87) — but these single-cycle windows are individually underpowered: each spans only four to five congresses, and the cross-district variation in exposure timing that identifies the effect shrinks once television is near-universal by the mid-1950s. The larger early-cycle point estimate is consistent with the novelty attenuation documented in the main text, not with a redistricting artifact. Implementation: `script 24_headline_robustness.R`.

E Appendix E: Variable construction details

E.1 Television exposure measure

Years of TV exposure for district d in congress t is constructed as:

$$\text{years_of_tv}_{dt} = \sum_{c \in d} w_c \cdot \max(0, \text{year}_t - \text{first_tv}_c)$$

where c indexes the counties in district d , w_c is the 1940 census population weight, and first_tv_c is the first year of commercial television station availability in county c 's designated market area. Source: ICPSR 22720 (Gentzkow and Shapiro, 2008) for TV introduction years and DMA boundaries; 1940 census for population weights.

E.2 Legislative effectiveness scores

I use Chiou and Goplerud's (2024) extension of the Volden-Wiseman LES framework (Volden and Wiseman, 2014) to historical congresses. The score weights legislative accomplishments by their substantive significance, following the original Volden-Wiseman convention. I normalize to a congress-specific mean of 1.

E.3 Majority-party indicator

The majority-party indicator equals 1 if the legislator's party held a majority of House seats in the relevant Congress and 0 otherwise. Majority composition: 78–79 (D), 80 (R), 81–82 (D), 83 (R), 84–89 (D).

E.4 ANES Type A and Type B variables

Variable	ANES item	Asked in	Coding
recall House cand name	VCF0972	1958, 1964	1 if Yes recalled, 0 if No/DK
knows if cand was incumbent	VCF0977	1958, 1964	1 if Yes, 0 if No/DK
correctly IDs incumbent	VCF0978	1958, 1964	1 if correct, 0 if incorrect
perceives party difference	VCF0501	1952, 60, 64	1 if Yes, 0 if No/DK
knows cons party	VCF0502	1960, 64	1 if picks one party, 0 if DK/no diff
cares about pres outcome	VCF0311	1952–64	1 if cares a good deal, 0 otherwise
watched political TV	VCF0724	1952–64	1 if watched, 0 if not

Implementation: `script 18_anes_typea_typeb.R`.

E.5 Seniority quartiles

Seniority is the number of consecutive previous Congresses in which the legislator served. Quartile cuts are computed within the majority-party pool to ensure comparability across cells. For the 1948–1964 sample, the cuts are: $Q1 \leq 1$ term, $Q2 \leq 3$ terms, $Q3 \leq 5$ terms, $Q4 > 5$ terms. The 6+ terms threshold corresponds approximately to chair-eligibility in the seniority system that prevailed in this period. Implementation: `script 19_majority_subgroup.R`.